

ABSTRACT

Title of dissertation:

ALGORITHMS AND DATA STRUCTURES
FOR FASTER NEAREST-NEIGHBOR
CLASSIFICATION

Alejandro Flores-Velazco,
Doctor of Philosophy, 2022

Dissertation directed by:

Professor David M. Mount
Computer Science

Given a set P of n labeled points in a metric space $(\mathcal{X}, \mathbf{d})$, the *nearest-neighbor rule* classifies an unlabeled query point $q \in \mathcal{X}$ with the class of q 's closest point in P . Despite the advent of more sophisticated techniques, nearest-neighbor classification is still fundamental for many machine-learning applications. Over the years, this has motivated numerous research aiming to reduce its high dependency on the size and dimensionality of the data. This dissertation presents various approaches to reduce the dependency of the nearest-neighbor rule from n to some smaller parameter k , that describes the intrinsic complexity of the class boundaries of P . This is of particular significance as it is usually assumed that $k \ll n$ on real-world training sets.

One natural way to achieve this dependency reduction is to reduce the training set itself, selecting a subset $R \subseteq P$ to be used by the nearest-neighbor rule to answer incoming queries, instead of using P . Essentially, reducing its dependency from n , the size of P , to the size of R . We propose different techniques to select R , all of which select subsets whose sizes are proportional to k , and have varying degrees of correct classification guarantees.

Another alternative involves building data structures designed to answer these classification queries, bypassing the preprocessing step of reducing P . We propose the Chromatic AVD to answer ε -approximate nearest-neighbor classification queries, and whose query times and storage requirements dependent on k_ε , which describes the intrinsic complexity of the ε -approximate class boundaries of P .

ALGORITHMS AND DATA STRUCTURES FOR
FASTER NEAREST-NEIGHBOR CLASSIFICATION

by

Alejandro Flores-Velazco

Dissertation submitted to the Faculty of the Graduate School of the
University of Maryland, College Park in partial fulfillment
of the requirements for the degree of
Doctor of Philosophy
2022

Advisory Committee:

Professor David M. Mount: *Advisor, Chair*

Professor Yancy Diaz-Mercado: *Dean's Representative*

Professor MohammadTaghi HajiAghayi

Professor John P. Dickerson

Professor Hanan Samet

Professor Samir Khuller: *External Member*

© Copyright by
Alejandro Flores-Velazco
2022

To Nelly.

Table of Contents

Dedication	ii
1 Introduction	1
1.1 Training Set Reduction	2
1.2 Nearest-Neighbor Search vs Classification	4
2 Literature Review	6
2.1 Classification Boundaries	6
2.2 Boundary Preservation	8
2.2.1 Clarkson's algorithm	8
2.2.2 Eppstein's algorithm	8
2.2.3 Bremner <i>et al.</i> 's algorithm	9
2.3 Nearest-Neighbor Condensation	10
2.3.1 Criteria	10
2.3.2 Hardness Results	11
2.3.3 Algorithms	11
2.4 Nearest-Neighbor Search	13
2.4.1 Exact Search	13
2.4.2 Approximate Search	13
2.4.3 Chromatic Search	14
3 Boundary Preserving Algorithms	15
3.1 Introduction	15
3.2 Eppstein's algorithm	16
3.2.1 The Initialization Phase	16
3.2.2 The Search Phase	17
3.2.3 The Inversion Method	17
3.3 A Simpler Initialization	17
3.3.1 Correctness Proof	19
3.3.2 Completeness Proof	20
4 Condensation Algorithms	24
4.1 Introduction	24

4.2	Lower Bounds	25
4.3	Consistent Subsets	25
4.3.1	CNN	26
4.3.2	FCNN	27
4.3.3	SFCNN	33
4.4	Selective Subsets	34
4.4.1	NET	34
4.4.2	MSS	35
4.4.3	RSS	37
4.4.4	VSS	38
4.5	Experimental Comparison	40
5	ε -coresets	43
5.1	Introduction	43
5.2	Coreset Characterization	45
5.2.1	Approximation-Sensitive Condensation	46
5.2.2	Guarantees on Classification Accuracy	47
5.3	Coreset Computation	51
5.3.1	Hardness Results	51
5.3.2	An Algorithm for α -Selective Subsets	53
5.3.2.1	Size in Doubling spaces	53
5.3.2.2	Size in Euclidean space	56
5.3.2.3	Subquadratic Implementation	58
5.3.3	An Algorithm for α -Consistent Subsets	60
5.4	Experimental Comparison	62
5.4.1	Algorithm Comparison	63
5.4.2	Subquadratic Approach	63
6	Chromatic Approximate Voronoi Diagram	65
6.1	Introduction	65
6.2	Preliminary Ideas and Intuition	68
6.3	Chromatic AVD Construction	70
6.3.1	The Build Step	70
6.3.2	The Reduce Step	74
6.4	Tree-size Analysis	76
6.4.1	Initial Size Bounds	76
6.4.2	Spatial Amortization	78
7	Conclusions	81
7.1	Training Set Reduction	81
7.1.1	Boundary Preservation	81
7.1.2	Condensation Algorithms	82
7.1.3	ε -coresets	82
7.2	Chromatic Nearest-Neighbor Search	83

¹ Chapter 1: Introduction

² In the context of non-parametric classification in *machine learning*, we are
³ given a training set P which consists of n points in a metric space $(\mathcal{X}, \mathbf{d})$, where
⁴ $P \subseteq \mathcal{X}$ and metric $\mathbf{d} : \mathcal{X}^2 \rightarrow \mathbb{R}^+$ defines the distance between any two points in \mathcal{X} .
⁵ Additionally, this training set P is partitioned into a finite set of *classes*, meaning
⁶ that each point $p \in P$ is assigned a *label* $l(p)$ that indicates the class to which p
⁷ belongs to. Finally, given an *unlabeled* query point $q \in \mathcal{X}$, the goal of a *classifier* is
⁸ to predict q 's label using the training set P .

⁹ The *nearest-neighbor rule* (or *nearest-neighbor classifier*) is among the best-
¹⁰ known classification techniques [1]. It classifies any query point $q \in \mathcal{X}$ with the
¹¹ label of its closest point in P according to the distance function \mathbf{d} . This idea can
¹² be generalized to predict q 's class using its first k nearest-neighbors in P , instead of
¹³ only one. Such generalized approach is known to as the k nearest-neighbor classifier,
¹⁴ or simply k -NN. Throughout this book, we focus exclusively on the 1-NN classifier.

¹⁵ Despite being conceptually simple, numerous results show that the nearest-
¹⁶ neighbor rule exhibits good classification accuracy both experimentally and theoreti-
¹⁷ cally [2–4]. In fact, its probability of error is bounded by twice the Bayes probability
¹⁸ of error, which is the best achievable error by any theoretical classifier. However, the
¹⁹ nearest-neighbor rule is often criticized for its high space and time complexities, as
²⁰ a straightforward implementation of this approach implies storing all the points of
²¹ P to answer such queries. This would therefore make the nearest-neighbor rule be
²² highly dependent on the size n of the training set P .

²³ This drawback raises an important question of whether it is possible to reduce
²⁴ the dependency of the nearest-neighbor classifier, from n to some smaller parameter.
²⁵ In this thesis, we proof this is indeed possible, proposing new approaches for nearest-
²⁶ neighbor classification that are dependent on some parameter k , where commonly
²⁷ $k \ll n$ on real training sets. This parameter is defined as the number of *border*
²⁸ *points* of P , and depicts the inherent complexity of the boundaries between points
²⁹ of different classes in the training set.

³⁰ Some of the approaches presented in this thesis deal with computing a subset
³¹ of P whose size is dependent on k , which can then be used by the nearest-neighbor
³² rule to answer classification queries, instead of using the entire training set. These
³³ are known as training set reduction techniques, and are explored in Chapters 3 to 5.
³⁴ Additionally, Chapter 6 explores a completely different approach, proposing a tailor-
³⁵ made data structure that can directly answer nearest-neighbor classification queries,
³⁶ thus bypassing the preprocessing step of having to reduce the training set. In
³⁷ terms of the underlying metric space, Chapters 3, 4 and 6 assume P lies in \mathbb{R}^d for

38 constant dimension d , and using ℓ_2 as the distance metric, while Chapter 5 deals with
39 training sets in general metric spaces. Finally, in terms of the model of computation,
40 Chapters 3 and 4 assume that nearest-neighbor queries are computed exactly, while
41 Chapters 5 and 6 allow multiplicative approximation errors when answering such
42 classification queries.

43 Preliminaries

44 The set of *border points* of the training set P are those that define the “bound-
45 aries” between points of different classes, and whose omission from the training set
46 would imply the misclassification of some query point in \mathcal{X} . Formally, two points
47 $p, \hat{p} \in P$ are border points of P if they belong to different classes, and there exist
48 some point $q \in \mathcal{X}$ such that q is equidistant to both p and \hat{p} , and no other point of
49 P is closer to q than these two points.

50 For any given point $p \in P$, define an *enemy* of p to be any point of P in a
51 different class than p . The *nearest-enemy* of p , denoted $\text{ne}(p)$, is the closest such
52 point according to metric \mathbf{d} . Finally, define the *nearest-enemy distance* of p to be
53 $\mathbf{d}_{\text{ne}}(p) = \mathbf{d}(p, \text{ne}(p))$, and the *nearest-enemy ball* of p to be the metric ball centered at
54 p with radius $\mathbf{d}_{\text{ne}}(p)$. Let κ be the number of distinct nearest-enemy points of P , that
55 is, the cardinality of the set $\{\text{ne}(p) \mid p \in P\}$. Similarly, denote the nearest-neighbor
56 of p as $\text{nn}(p)$, and the nearest-neighbor distance as $\mathbf{d}_{\text{nn}}(p) = \mathbf{d}(p, \text{nn}(p))$.

57 In fact, in this thesis we prove that every nearest-enemy of P is also a border
58 point, implying that $\kappa \leq k$. Thus, we will use these two parameters to analyze our
59 proposed approaches, showing their dependency on the complexity of the boundaries
60 between points of different classes in P .

61 Through a suitable uniform scaling, we may assume that the *diameter* of P
62 (*i.e.*, the maximum distance between any two points in the training set) is 1. The
63 *spread* of P , denoted as Δ , is defined to be the ratio between the largest and smallest
64 distances in P . Define the *margin* of P , denoted γ , to be the smallest nearest-enemy
65 distance in P . Clearly, this implies that $1/\gamma \leq \Delta$.

66 Additionally, a metric space $(\mathcal{X}, \mathbf{d})$ is said to be *doubling* [5] if there exist some
67 bounded value λ such that any metric ball of radius r can be covered with at most
68 λ metric balls of radius $r/2$. Its *doubling dimension* is the base-2 logarithm of λ ,
69 denoted as $\text{ddim}(\mathcal{X}) = \log \lambda$. Many natural metric spaces of interest are doubling,
70 including d -dimensional Euclidean space whose doubling dimension is $\Theta(d)$.

71 1.1 Training Set Reduction

72 The idea behind training set reduction techniques is to select a subset of the
73 original training set P , and then use this reduced set to answer any nearest-neighbor
74 classification queries. Evidently, the goal is to select a subset as small as possible,
75 subject to maintaining the classification accuracy of the nearest-neighbor classifier.

76 By definition, if instead of applying the nearest-neighbor rule with the entire
77 training set P we use the set of border points of P , the complexity of answering

78 classification queries is now dependent on k instead of n , while still obtaining the
 79 same classification for any query point in \mathcal{X} . In other words, this approach would
 80 maintain the same boundaries between points of different classes, before and after
 81 reducing the training set. For this reason, algorithms for finding the set of border
 82 points of P are known as *boundary preservation* algorithms.

83 For training sets $P \subset \mathbb{R}^2$ in the Euclidean plane, Bremner *et al.* [6] proposed
 84 an output-sensitive algorithm for finding the set of border points of P in $\mathcal{O}(n \log k)$
 85 worst-case time. For almost three decades, the best result for training sets in \mathbb{R}^d ,
 86 assuming constant d , was Clarkson's [7] algorithm that runs in $\mathcal{O}(\min(n^3, kn^2 \log n))$
 87 worst-case time. Recently, Eppstein [8] proposed a significantly faster algorithm for
 88 the d -dimensional Euclidean case, which runs in $\mathcal{O}(n^2 + nk^2)$ worst-case time. In
 89 Chapter 3, we propose an improvement over Eppstein's algorithm [8] to compute the
 90 set of border points of any training set $P \subset \mathbb{R}^d$, where dimension d is assumed to be
 91 constant. Moreover, our algorithm reduces the complexity of computing the set of
 92 border points of P to $\mathcal{O}(nk^2)$ worst-case time, where k is the size of such subset.




Figure 1.1: Illustrative example of the subsets selected by different condensation algorithms from an initial training set P in \mathbb{R}^2 of 10^4 points. It includes existing algorithms (b)-(e), along with new algorithms introduced by our work (f)-(h).

93 Another approach for training set reduction is called condensation. Formally,
 94 the problem of *nearest-neighbor condensation* consists of finding a consistent subset
 95 of P , where a subset $R \subseteq P$ is said to be *consistent* [9] if and only if for every point
 96 $p \in P$ its nearest-neighbor in R is of the same class as p . Intuitively, R is consistent
 97 if and only if all points of P are correctly classified using the nearest-neighbor rule
 98 over R . Compared to the boundary preservation techniques, using consistent subsets
 99 only guarantees the correct classification of the points of P , while using the set of
 100 border points extends this guarantee for every point of \mathcal{X} .

101 While finding subsets of P that are consistent can be done efficiently, computing
102 minimum cardinality consistent subsets is much harder. In fact, it is known to be an
103 NP-hard problem [10–12], even to approximate within practical factors. Therefore,
104 most research has focused on proposing practical heuristics to find either consistent
105 subsets (for comprehensive surveys, see *e.g.*, [13–15]). Some of the best known heuristic
106 algorithms for this problem are known as CNN [9], FCNN [16], and MSS [17], the last
107 two being considered state-of-the-art for computing consistent subsets in $\mathcal{O}(n^2)$ time.
108 However, while these heuristics have been extensively studied experimentally [18],
109 theoretical results are scarce, with no upper-bounds known for the size of the subsets
110 selected by these algorithms. In Chapter 4, we present the first theoretical results on
111 upper-bounding the subset sizes of condensation algorithms. In particular, we show
112 that FCNN and MSS can not be bounded in terms of k , and propose new quadratic-
113 time algorithms called SFCNN, RSS and VSS, that can be effectively upper-bounded
114 in terms of k (see Figure 1.1 for illustrative examples on their selected subsets).

115 So far, these approaches for training set reduction have made a key assumption:
116 that after reduction, the nearest-neighbor rule will answer classification queries (*i.e.*,
117 compute nearest-neighbors on the reduced training set) exactly. However, in practice,
118 nearest-neighbors are frequently computed approximately rather than exactly. In
119 this context, given an approximation parameter $\varepsilon \in [0, 1]$, a query point $q \in \mathcal{X}$ can
120 be assigned the class of any point of P whose distance to q is at most $1+\varepsilon$ times the
121 distance from q to its true nearest-neighbor. Sadly, the classification guarantees given
122 by both boundary preservation and condensation algorithms, rely on the assumption
123 that nearest-neighbors are computed exactly, and break when approximate query
124 answers are allowed. In Chapter 5, we propose a framework for training set reduction
125 that is sensitive to these approximations, along with a characterization of ε -coresets
126 for the problem of nearest-neighbor classification.

127 1.2 Nearest-Neighbor Search vs Classification

128 Due to its conceptual closeness, the problem of nearest-neighbor classification
129 is extremely related to the problem of nearest-neighbor search. Evidently, one can
130 reduce the problem of classifying a query point $q \in \mathcal{X}$ via the nearest-neighbor rule, to
131 simply compute q 's closest point in the training set, and assign q to the class of such
132 point. Unsurprisingly, this approach is standard for nearest-neighbor classification,
133 and under such problem reduction, there are only two alternatives to improve the
134 complexity of answering these classification queries. Either by reducing the training
135 set (as explained in the previous section), or by improving the techniques to answer
136 nearest-neighbor queries, which is a known and extensive line of research [19–22].

137 However, there exist an alternative approach towards achieving more efficient
138 nearest-neighbor classification. This would imply avoiding the reduction to the search
139 problem, and instead having algorithms and data structures to directly compute
140 the class of the nearest-neighbor of any given query point; *i.e.*, without computing
141 the nearest-neighbor itself. In computational geometry, this is usually known as the
142 *chromatic nearest-neighbor search* problem.

143 In Chapter 6, we propose a tailor-made data structure for approximate nearest-
144 neighbor classification (or approximate chromatic nearest-neighbor search), which we
145 call the *Chromatic AVD*. Given a training set P in d -dimensional Euclidean space
146 (assuming constant d and the ℓ_2 metric) and an approximation parameter $\varepsilon \in [0, \frac{1}{2}]$,
147 we construct a quadtree-based partitioning of space to answer any classification
148 query approximately. That is, for any query point $q \in \mathbb{R}^d$ this data structure returns
149 the class of any of q 's valid ε -approximate nearest-neighbors in P . Moreover, the
150 Chromatic AVD is designed as a simplification of state-of-the-art AVDs [20] for
151 approximate nearest-neighbor search. Thus, reducing its dependency from n to a
152 parameter k_ε that describes the complexity of the approximate class boundaries.

₁₅₃ Chapter 2: Literature Review

₁₅₄ The problem of classification is of high relevance in the field of machine learning.
₁₅₅ It has motivated numerous approaches that are able to predict the class of a given
₁₅₆ query point, by constructing and using some model “trained” from a given training
₁₅₇ set. However, and despite the advent of more sophisticated techniques (*e.g.*, support-
₁₅₈ vector machines [23] and deep neural networks [24]), nearest-neighbor classification
₁₅₉ is still widely used in practice [25–27], proving its value in constructing resilient
₁₆₀ defense strategies against adversarial [28] and poisoning [29] attacks, as well as in
₁₆₁ achieving interpretable and reliable classification models [30, 31].

₁₆₂ Its importance has motivated diverse research in many fields, from statistics to
₁₆₃ computational geometry. In this book, we approach the nearest-neighbor classification
₁₆₄ problem from the perspective of computational geometry, leveraging algorithms, data
₁₆₅ structures, and analysis techniques from this field. This chapter delves into the most
₁₆₆ important concepts related to this classification technique, as well as the previous
₁₆₇ work on ways to improve its efficiency.

₁₆₈ 2.1 Classification Boundaries

₁₆₉ The concept of *boundaries* in classification can be vague, as it depends on the
₁₇₀ particularities of the classification model being studied. The terms *class* boundaries,
₁₇₁ *classification* boundaries, and *decision* boundaries often refer to the same concept.
₁₇₂ Broadly speaking, they refer to the separation between two regions in space, such
₁₇₃ that the model of interest classifies points lying on either side of the boundary (*i.e.*,
₁₇₄ on each region) with different classes.

₁₇₅ In order to understand the class boundaries of the nearest-neighbor rule, we
₁₇₆ first need to introduce a well-known concept in computational geometry called
₁₇₇ *Voronoi Diagrams*. The Voronoi diagram of a point set $P \subseteq \mathcal{X}$ is a partition of the
₁₇₈ underlying space \mathcal{X} into different regions called *cells*. Each of these Voronoi cells has
₁₇₉ an associated point $p \in P$ (often call the *site* of the cell) such that for every point q
₁₈₀ inside of the cell, p is q ’s closest point in P according to the distance metric d .

₁₈₁ Throughout this book, all the figures depicting Voronoi diagrams will assume
₁₈₂ that the underlying metric space is Euclidean, and the distance function is the ℓ_2
₁₈₃ metric. Here, the Voronoi cell of each point $p \in P$ is a convex region of \mathbb{R}^d , and can
₁₈₄ be described as the intersection of $n - 1$ closed halfspaces, each being the halfspace
₁₈₅ containing p that is bounded by the bisector between p and another point of P .

₁₈₆ Two sites $p, \hat{p} \in P$ are said to be *Delaunay neighbors*, if and only if there exists
₁₈₇ a point q in the underlying space that is part of the cells of both p and \hat{p} . That is, if q

is equidistant to both p and \hat{p} , and no other points of P are closer to q . This defines the edges of the *Delaunay triangulation* of P , which is characterized as the dual graph of the Voronoi Diagram. Moreover, we can define the boundaries between the Voronoi cells of P as the union of all such points q in the space that are equidistant to at least two points/sites of P .

When considering the classes of these points, we can introduce a few useful concept. First, any edge of the Delaunay triangulation of P that connects two points of different classes is called a *bichromatic edge*. Additionally, we can define the class boundaries of the nearest-neighbor rule similarly to the definition of the boundaries between the Voronoi cells of P . We say the *class boundaries* of P are the union of points q in the underlying space that are equidistant to at least two points of P from different classes. Intuitively, these boundaries separate different *class regions*, where each such region is defined as the union of neighboring Voronoi cells corresponding to points of the same class. See Figure 2.1 for an illustrative example on these concepts.




Figure 2.1: Example of a training set $P \in \mathbb{R}^2$ with points of three classes: *red*, *blue* and *yellow*. (a) shows the different class regions defined by the union of adjacent Voronoi cells of the same class, (b) highlights the border points of P , while (c) highlights the nearest-enemy points of P . Both (a) and (b) draw the class boundaries with solid black lines.

We say that the class boundaries of P are defined by a subset of the training set called the set of *border points* of P . Basically, these border points are the endpoints of all the bichromatic edges of the Delaunay triangulation of P . Formally, two points $p, \hat{p} \in P$ are border points of P if they belong to different classes, and there exist some point q in the underlying space such that q is equidistant to both p and \hat{p} , and no other point of P is closer to q than these two points. See Figure 2.1a and 2.1b for an example of a training set $P \subset \mathbb{R}^2$ and its set of border points. Throughout, we denote k to be the total number of border points in the training set. Unsurprisingly, k is key in understanding the complexity of the class boundaries of P , making it an ideal candidate to analyze the complexity of the classification problem.

Another concept that is useful to characterize the class boundaries of P is that

213 of nearest-enemy points. Just as defined in Chapter 1, we say the *nearest-enemy*¹ of
 214 a point $p \in P$ (denoted by $\text{ne}(p)$) is p 's closest point in P of different class. Then,
 215 we let κ be the number of distinct nearest-enemy points of P , that is, the cardinality
 216 of the set $\{\text{ne}(p) \mid p \in P\}$. See an example of this concept in Figure 2.1c. While not
 217 immediately obvious, we are able to (see Section 4.4) that every nearest-enemy point
 218 of P is also a border point of P . Therefore, the set of nearest-enemy points is a subset
 219 of those points defining the class boundaries, and can then be used to characterize
 220 its complexity.

221 2.2 Boundary Preservation

222 In this section, we review known algorithms for finding the set of border points
 223 of the training set P , also referred to as *boundary preservation algorithms*. Note that
 224 for each of these algorithms, P is assumed to lie in d -dimensional Euclidean space
 225 (*i.e.*, $P \subset \mathbb{R}^d$), and the distance function is assumed to be the ℓ_2 metric.

226 2.2.1 Clarkson's algorithm

227 In 1994, Clarkson [7] proposed the first algorithm to compute the set of border
 228 points of a training set P in \mathbb{R}^d . His algorithm runs in $\mathcal{O}(\min(n^3, kn^2 \log n))$ time, by
 229 leveraging linear programming formulations along with output-sensitive techniques
 230 to compute extreme points.

231 Starting with an empty selection, the algorithm incrementally adds newly found
 232 border points to its current selection. To achieve this, it performs a “non-borderness”
 233 test on every point p of the training set. This test can only decide with certainty if
 234 p is *not* a border point, but not otherwise. However, if the test’s result is uncertain,
 235 Clarkson proposes a method to find some other point $\hat{p} \in P$ that is certainly a border
 236 point. Further details on this algorithm are left for the interested reader, and can be
 237 found in Section 5 of Clarkson’s paper.

238 2.2.2 Eppstein’s algorithm

239 In early 2022, Eppstein [8] proposed a new algorithm to find the set of border
 240 points of a given training set P in \mathbb{R}^d that runs in $\mathcal{O}(n^2 + nk^2)$ worst-case time.
 241 The algorithm’s simplicity is striking, specially when considering the small progress
 242 achieved for this problem in the last three decades. Intuitively, the algorithm works
 243 somewhat as an implicit graph traversal, where the border points of P are nodes in
 244 this graph, connected by certain implicit edges. Formally, the algorithm begins by
 245 selecting some initial set of border points of P , one point from every class region.
 246 From here, it uses a series of subroutines (which we group together and denote as the
 247 “*inversion method*”) to find the remaining border points of P . To better understand
 248 this algorithm, we study its two phases: the initialization and search phases.

¹This concept was first introduced by Dasarathy [32] in 1995 under the name of *nearest unlike neighbors* (or NUN), and was later renamed by Wilson & Martinez [33] in 1997 as nearest-enemies.

249 The initialization phase consists of finding a subset of points of P , such that
250 all these must be border points, and there is at least one point for every class region
251 of P . Eppstein’s approach to find these points is to compute the Minimum Spanning
252 Tree (MST) of P , identify the bichromatic edges of this MST, and then select the
253 endpoints of all such edges. Evidently, this phase takes a total of $\mathcal{O}(n^2)$ time.

254 The search phase completes the algorithm, and consists of finding all the
255 remaining border point of the training set. This phase iterates over all the currently
256 selected points, and for each of these points p , it performs the inversion method on p ,
257 identifying a subset of border points that are “visible” by p . These are then selected
258 by the algorithm, and the search continues. Basically, the inversion method unveils
259 the edges of some implicit graph that are incident on p , allowing the algorithm to
260 fully traverse this implicit graph. Once the inversion method has been applied on
261 every selected point, the algorithm terminates with the guarantee of having selected
262 all the border points of P . Each call of the inversion method takes $\mathcal{O}(nk)$ time,
263 making the total runtime of this phase being $\mathcal{O}(nk^2)$.

264 It is important to note that Eppstein’s algorithm works under the assumption
265 that this implicit graph can have multiple connected components. Therefore, the
266 initialization phase looks to select at least one point on every connected component,
267 to guarantee that all the border points will eventually be discovered by the algorithm.
268 In Chapter 3, we further improve Eppstein’s algorithm by showing that in fact, this
269 implicit graph has a single connected component. Hence, rendering the initialization
270 phase unnecessary, and reducing the total runtime to $\mathcal{O}(nk^2)$.

271 2.2.3 Bremner *et al.*’s algorithm

272 This is a special case algorithm intended only for the planar case; *i.e.*, then the
273 training set lies in \mathbb{R}^2 . It was proposed by Bremner *et al.* [6] in 2003, as the result of
274 a research workshop, and has a runtime of $\mathcal{O}(n \log k)$.

275 Their algorithm consists of two levels. First, the higher level algorithm, which
276 consists of repeatedly guessing the value of k , and running a lower level algorithm
277 under this assumption. This lower level algorithm works as follows: given some value
278 m , if $m \geq k$ it finds the set of all border points of P , and otherwise it fails. Moreover,
279 it does this within $\mathcal{O}((m^2 + n) \log m)$ time. Therefore, by repeatedly using values of
280 $m = 2^{2^i}$, with $i = 0, 1, 2, \dots, \lceil \log \log k \rceil$, and stopping when $m \geq k$ or $m \geq \sqrt{n}$, the
281 total runtime of the higher level algorithm equals $\mathcal{O}(n \log k)$.

282 Interestingly, the lower level algorithm works in a similar way to Eppstein’s
283 algorithm. First, it finds an initial bichromatic edge of the Delaunay triangulation
284 of P (*i.e.*, identifying two border points), and then follows iteratively by applying a
285 series of “pivot operations” on the border points found so far, to unveil new border
286 points of P . If at some point, the algorithm finds more than m border points, it fails
287 as the assumption that $m \geq k$ is false. Otherwise, the algorithm terminates with the
288 guarantee of having found all the border points of P .

289 The remaining details of this algorithm are left to the reader, and can be found
290 in the original paper [6]. However, we describe the pivot operation in higher detail, as
291 it would be leveraged throughout this book as a useful tool in some of our algorithms

292 and analyses. Basically, any pivot operation receives three parameters: a point p ,
293 a vector \vec{v} , and a set of points $S \subset \mathbb{R}^d$. Intuitively, the pivot operation grows an
294 empty ball with p on its surface and its center on the direction of \vec{v} , until the ball
295 hits a point of $\hat{p} \in S$. Then, the operation returns point \hat{p} . Despite its simplicity,
296 this operation is a useful tool to find border points, as seen in the remaining of this
297 book.

298 2.3 Nearest-Neighbor Condensation

299 Essentially, the problem of *nearest-neighbor condensation* consists of selecting
300 some subset of the original training set P , subject to the classification accuracy of
301 the nearest-neighbor rule not being “greatly reduced” when replacing the original
302 training set by this subset of points. One way of achieving this would be to simply
303 select the set of border points of P via a boundary preserving algorithm. However,
304 this approach is too strict, which has lead to the introduction of more relaxed criteria
305 for the purpose of condensation.

306 In this section, we explore the main criteria used for condensation, hardness
307 results, as well as the most relevant algorithms proposed for this problem.

308 2.3.1 Criteria

309 The central definition to understand the problem of condensation is that of
310 consistent subsets, proposed by Hart [9] in 1968. We say a subset $R \subseteq P$ is *consistent*
311 if and only if for every point $p \in P$ its nearest-neighbor in R is of the same class as
312 p . Intuitively, this means that any subset R is consistent if and only if all points of
313 P are correctly classified using the nearest-neighbor rule “trained” on R . That is,
314 with a nearest-neighbor rule that answers queries by searching among the points in
315 R , instead of searching among the points of the original training set P .

316 Another criterion frequently used for condensation is known as *selectiveness*,
317 and was proposed by Ritter *et al.* [34] in 1975. A subset $R \subseteq P$ is said to be *selective*
318 if and only if for all points $p \in P$ its nearest-neighbor in R is closer to p than its
319 nearest-enemy in P . Clearly selectiveness implies consistency, as the nearest-enemy
320 distance in R of any point of P would at least be its nearest-enemy distance in P . In
321 fact, selective subsets were introduced as a stricter version of consistency, hoping this
322 would allow the existence of polynomial-time algorithms to find minimum cardinality
323 selective subsets. As we will see in the following subsections, this would later be
324 discovered to be impossible unless $P=NP$.

325 Recall that none of these two criteria, namely consistency and selectiveness,
326 imply that every query point in \mathcal{X} would be correctly classified after condensation.
327 Instead, they simply guarantee the correct classification of the points in P . The only
328 techniques that guarantee correct classification of every query point are boundary
329 preservation algorithms, as explored in Section 2.2. It is easy to see that any subset
330 that preserves the class boundaries of P (*i.e.*, the set of border points of P) must be
331 selective, and that any selective subset of P must also be consistent. Therefore, any

332 algorithm that finds subsets of P holding any of these properties is considered to be
333 a condensation algorithm.

334 2.3.2 Hardness Results

335 Clearly the stricter notion of condensation, involving the full preservation of the
336 class boundaries of P , can be achieved in polynomial-time as described in Section 2.2.
337 The relaxation to the consistent and selective criteria imply that subsets holding
338 these properties can be computed even faster. These subsets always exist, as P itself
339 is both consistent and selective. Evidently, the more interesting research question
340 deals with finding ideally small subsets of P under these two condensation criteria.

341 Therefore, understanding the complexity of finding such subsets while minimiz-
342 ing their sizes becomes of major relevance. It took close to 30 years since consistency
343 was originally defined in 1968 for the first hardness results to appear, and it took
344 another 30 years to close some of the remaining gaps into understanding the hardness
345 of approximation for these problems.

346 Denote MIN-CS and MIN-SS to be the problems of finding minimum cardinality
347 consistent and selective subsets of P , respectively, we know that these two problems
348 are NP-hard to solve, and their decision versions are NP-complete [10–12]. Originally,
349 Wilfong [10] proved this complexity for when P lies in Euclidean space, and restricting
350 the result for MIN-CS on training sets with at least 3 classes. These results were
351 later generalized by Zukhba [11] and Khodamoradi *et al.* [12] for the cases when P
352 has at least two classes, and also when it lies in some general metric space (\mathcal{X}, d) .

353 Hardness of approximation results have also been proposed for this problem.
354 The first result dates back to 2000, when Nock & Sebban [35] found that unless $\text{NP} \subseteq$
355 $\text{DTIME}(n^{\log \log n})$, the MIN-SS problem is NP-hard to approximate in polynomial-
356 time within a factor of $(1 - o(1)) \ln n$. Later in 2014, Gottlieb *et al.* [36] found that
357 the MIN-CS problem is NP-hard to approximate in polynomial-time within a factor
358 of $2^{(\text{ddim}(\mathcal{X}) \log 1/\gamma)^{1-o(1)}}$, where $\text{ddim}(\mathcal{X})$ is defined as the doubling dimension of the
359 space \mathcal{X} , and γ is the margin or minimum nearest-enemy distance of P .

360 More recently, Chitnis [37] presented the first results on the parameterized
361 complexity of the decision version of MIN-CS. This version of the problem decides
362 whether there exists a consistent subset of P of size $\leq m$, assuming $P \subset \mathbb{Z}^d$ and ℓ_p
363 being the distance metric. Then, Chitnis proves two main results. First, that this
364 problem is W[1]-hard parameterized by $m + d$, meaning that unless FPT=W[1], there
365 is no $f(m, d) \cdot n^{O(1)}$ time algorithm for any computable function f . Additionally, that
366 under the Exponential Time Hypothesis (ETH) there is no $d \geq 2$ and computable
367 function f such that this problem can be solved in $f(m, d) \cdot n^{o(m^{1-1/d})}$ time.

368 2.3.3 Algorithms

369 Due to the importance of this problem, numerous algorithmic approaches have
370 been proposed to compute such subsets. These include some optimal algorithms,
371 as well as some case-specific, and approximation algorithms. However, due to the
372 complexity of computing even close to optimal solutions for the condensation problem,

373 most research has focused on proposing practical heuristics to find either consistent
374 or selective subsets of P . For comprehensive surveys on these heuristics, see [13–15].

375 Here we review some of the most relevant algorithms proposed for this problem,
376 in chronological order. In 1968, along with the definition of consistency, Hart [9]
377 proposed the CNN (*Condensed Nearest-Neighbor*) algorithm to compute consistent
378 subsets of P . Even though it has been widely used in the literature, CNN suffers
379 from several drawbacks: its running time is cubic in the worst-case, and the resulting
380 subset is order-dependent, meaning that the points selected are determined by the
381 order in which they are considered by the algorithm. Additionally, CNN tends to
382 select points far from the class boundaries, which is usually undesirable as these
383 points are unlikely to contribute to such boundaries. To combat this, Gates [38]
384 proposed the RNN (*Reduced Nearest-Neighbor*) algorithm in 1972, which consists
385 of first running CNN and then proceed with a postprocessing technique to further
386 reduce the subset. While this resolves the issue of selecting points far from the class
387 boundaries, this approach still runs in worst-case cubic time and is order-dependent.

388 Also in 1975, Ritter *et al.* [34] proposed the SNN (*Selective Nearest-Neighbor*)
389 algorithm, along with the definition of selective subsets. The authors hoped it would
390 be possible to compute minimum cardinality selective subsets in polynomial time,
391 and propose this algorithm to find them. Unsurprisingly, it was later proved that its
392 runtime is worst-case exponential [10], even though Wilson & Martinez [33] claimed
393 that the algorithm’s average runtime is quadratic.

394 The new century saw a big leap in heuristic approaches for condensation. In
395 particular, two algorithms stand out: MSS (*Modified Selective Subset*) proposed by
396 Barandela *et al.* [17] in 2005, and FCNN (*Fast CNN*) proposed by Angiulli [16] in
397 2007. Both algorithms compute consistent and selective subsets, respectively. But
398 more importantly, both algorithms run in quadratic worst-case time, and are order-
399 independent. These characteristics, together with their ease of implementation, and
400 specially, their exceptional performance in practice, have established these algorithms
401 as state-of-the-art for condensation.

402 In 2014, Gottlieb *et al.* [36] proposed an approximation algorithm called NET,
403 along with the almost matching hardness lower-bounds described in Section 2.3.2.
404 The algorithm is fairly simple, consisting on just computing a γ -net of P where γ is
405 P ’s margin. This clearly results in a consistent subset, which can be proven to be a
406 tight approximation of the MIN-CS problem. However, in practice, γ tends to be
407 small in real training sets, thus making the subsets selected by the NET algorithm
408 of much higher cardinality than those selected by state-of-the-art heuristics.

409 Recently, in 2019 Biniaz *et al.* [39] proposed a subexponential-time algorithm
410 for finding minimum cardinality consistent subsets of point sets $P \subset \mathbb{R}^2$ in the
411 Euclidean plane, along with other case-specific algorithms for special instances of the
412 problem in \mathbb{R}^2 . Their main algorithm runs in $n^{\mathcal{O}(\sqrt{m})}$ where m is the size of minimum
413 cardinality consistent of P , which after the hardness results presented by Chitnis [37]
414 in 2022, it has been proven to be asymptotically tight for the planar case.

415 It is important to note that, while the heuristics described in this section (*i.e.*,
416 CNN, RNN, FCNN, and MSS) have been extensively studied experimentally [18],
417 theoretical results are scarce. Before our work described in Chapter 4, little was

418 known about any theoretical guarantees on the size of their selected subsets. There
419 was a clear gap between practical heuristics without any theoretical guarantee, and
420 approximation algorithms with poor performance in practice.

421 2.4 Nearest-Neighbor Search

422 Given a set $P \subseteq \mathcal{X}$ of n points in a metric space (\mathcal{X}, d) and a query point $q \in \mathcal{X}$,
423 the goal of the nearest-neighbor search problem is to compute q 's closest point in P
424 according to the distance function d . The most common assumption for this problem
425 is that P is given initially to be preprocessed into a data structure, while query points
426 are later received as a stream to be answered using the preprocessed data structure.

427 In this section, we discuss different techniques to compute nearest-neighbors,
428 either exactly or approximately. Additionally, we consider the related problem of
429 chromatic nearest-neighbor search, which assuming that P is a training set (*i.e.*, that
430 each point in P is labeled), its goal is to compute the class of each query point's
431 nearest-neighbor, and not necessarily its nearest-neighbor itself.

432 2.4.1 Exact Search

433 There are two main approaches to compute nearest-neighbors exactly. The
434 first, via exhaustive search over the entire point set P , which evidently implies linear
435 query times and storage. Otherwise, the approaches usually apply point location
436 algorithms over space partitioning data structures. For the low dimensional case
437 of $d \leq 2$, it is possible to obtain $\mathcal{O}(\log n)$ query times and linear storage. Sadly,
438 when the dimension is $d > 2$, the “*curse of dimensionality*” starts to kick in, and
439 the overhead between query times and storage requirements grows increasingly
440 fast. On either extreme, this means that we have solutions with logarithmic query
441 time but roughly $\mathcal{O}(n^{d/2})$ storage [40], or solutions with linear storage but barely
442 sublinear query times [41]. Therefore, the applicability of exact nearest-neighbor
443 search becomes severely limited for high-dimensional data.

444 2.4.2 Approximate Search

445 In practice, nearest-neighbors are often computed approximately rather than
446 exactly, as this relaxation allows for further improvements on query time and storage
447 requirements. Formally, given an approximation parameter $\varepsilon \in [0, 1]$, the problem of
448 ε -approximate nearest-neighbor searching (or ε -ANN) deals with computing a point
449 $p \in P$ whose distance from the query point q is within a factor of $1 + \varepsilon$ of q 's distance
450 to its nearest-neighbor in P . We say that any such point p is a valid ε -approximate
451 nearest-neighbor of q , and can be retrieved as a valid answer to q 's query.

452 This problem has been studied extensively, both theoretically [19–22, 42, 43]
453 and from a practical standpoint [44, 45]. When the number of dimensions is low, most
454 techniques involve some sort of space partitioning; *e.g.*, kd-Trees [46], Quadtrees [47–
455 49], and Approximate Voronoi Diagrams [19, 20, 43]. Otherwise, alternatives for high-
456 dimensional data include hashing techniques like Locality-Sensitive Hashing [21, 50],

457 and proximity graph techniques like Hierarchical Navigable Small Worlds graphs [22].
458 However, throughout this dissertation we focus on low-dimensional data approaches,
459 as many of our results have exponential dependency on the number of dimensions.

460 2.4.3 Chromatic Search

461 Evidently, given some query point q , to classify q with the nearest-neighbor
462 rule involves retrieving the class of q 's nearest-neighbor in the training set. Note that
463 this is slightly different from retrieving the nearest-neighbor itself. Potentially, this
464 difference opens the possibility of improvements over the standard search problem. In
465 the context of computational geometry, this is sometimes referred to as the problem
466 of *chromatic nearest-neighbor search*, where classes are intuitively described as colors.

467 Evidently, to answer *exact* chromatic queries there is one clear approach: to
468 compute the set of border points of the training set, and use any of the techniques
469 described in Section 2.4.1 to answer exact search queries over this subset. Despite
470 having increased preprocessing times, this approach would lead to query times and
471 storage requirements dependent on k instead of n . However, it is unclear if it is
472 possible to achieve similar results without having to recur to boundary preservation
473 algorithms and the reduction to the standard search problem.

474 However, this is indeed possible in the approximate setting. Some work has
475 been done along these lines by Mount *et al.* [51], showing that when query points
476 are far from the class boundaries and surrounded by points of the same class, query
477 times can be significantly reduced. However, these query times are still dependent on
478 n . Moreover, the data structure itself dates back to 2000, and is based on standard
479 search techniques of the time. Evidently, many improvements have been achieved in
480 the last two decades on approximate nearest-neighbor searching, making this data
481 structure clearly outdated.

482 In Chapter 6, we propose new data structure for this problem called *Chromatic*
483 AVD, thought as a simplification of state-of-the-art AVDs [20] from 2017. This new
484 approach has query times and storage dependencies on a parameter k_ε that describes
485 the complexity of the approximate class boundaries (similarly to how k describes
486 the complexity of the exact class boundaries).

488 3.1 Introduction

489 As previously discussed, a common approach to reduce the dependency on n of
 490 the nearest-neighbor rule is to reduce the training set itself. However, most training
 491 set reduction techniques offer limited guarantees on the effect that this reduction
 492 makes on the accuracy of the classifier. Only a handful of works [6–8] have proposed
 493 algorithms that guarantee the same classification of every query point, before and
 494 after the reduction took place. These are called *boundary preserving* algorithms, and
 495 are the focus of this chapter.

496 The results presented on this chapter can also be found here [52].




Figure 3.1: On the left, a training set $P \in \mathbb{R}^2$ with points of three classes: *red*, *blue* and *yellow*. On the right, a subset of these points corresponding to the set of border points of P . The solid black lines highlight the *boundaries* of P between points of different classes. By definition, the class boundaries remain the same in both cases.

497 While other problems in the realm of training set reduction are NP-hard [10–12]
 498 to solve exactly (*e.g.*, those of finding minimum cardinality *consistent* subsets and
 499 *selective* subsets), the problem of preserving the class boundaries of the nearest-
 500 neighbor rule is tractable. As discussed in Sections 2.1 and 2.2, this problem is
 501 equivalent to that of finding the set of border points of P .

502 The set of *border points* of the training set P are those that define the boundaries

503 between points of different classes, and whose omission from the training set would
 504 imply the misclassification of some query points in \mathbb{R}^d . Formally, as seen in Section 2.1,
 505 two points $p, \hat{p} \in P$ are border points of P if they belong to different classes, and
 506 there exist some point $q \in \mathbb{R}^d$ such that q is equidistant to both p and \hat{p} , and no
 507 other point of P is closer to q than these two points (*i.e.*, the empty ball property of
 508 Voronoi Diagrams). See Figure 3.1 for an example of a training set P in \mathbb{R}^2 and its
 509 set of border points. Throughout, we let k denote the total number of border points
 510 in the training set. By definition, if instead of applying the nearest-neighbor rule
 511 with the entire training set P we use the set of border points of P , its dependency is
 512 reduced from n to k , while still obtaining the same classification for any query point
 513 in \mathbb{R}^d . This becomes particularly relevant for applications where $k \ll n$.

514 For training sets $P \subset \mathbb{R}^2$ in 2-dimensional Euclidean space, Bremner *et al.* [6]
 515 proposed an output-sensitive algorithm for finding the set of border points of P in
 516 $\mathcal{O}(n \log k)$ worst-case time. However, how to generalize this algorithm for higher
 517 dimensions remained unclear. Until very recently, the best result for the higher
 518 dimensional case was that of Clarkson [7]. He proposed an algorithm to find the
 519 set of border points of $P \subset \mathbb{R}^d$, with bounded d , that runs in $\mathcal{O}(\min(n^3, kn^2 \log n))$
 520 worst-case time. For almost three decades, this remained the best result for training
 521 sets in \mathbb{R}^d . Recently, Eppstein [8] proposed a significantly faster algorithm for the
 522 d -dimensional Euclidean case, which runs in $\mathcal{O}(n^2 + nk^2)$ worst-case time.

523 In this chapter, we propose an improvement over Eppstein’s algorithm [8] to
 524 compute the set of border points of any training set $P \subset \mathbb{R}^d$, where dimension
 525 d is assumed to be constant. While the original algorithm computes such set in
 526 $\mathcal{O}(n^2 + nk^2)$ time, where k is the number of border points of P , our new algorithm
 527 computes the same set in $\mathcal{O}(nk^2)$ time.

528 3.2 Eppstein’s algorithm

529 This algorithm is strikingly simple, yet full of interesting ideas (see a formal
 530 description in Algorithm 1), and it works as follows: it begins by selecting an initial
 531 set of border points of P , one point from every class region. From here, the algorithm
 532 uses a series of subroutines which we will group together and denote as the “*inversion*
 533 *method*”, to find the remaining border points of P . Thus, the algorithm can be
 534 naturally split into two phases: the initialization of R with some border points, and
 535 the search process for the remaining border points of P .

536 3.2.1 The Initialization Phase

537 The initialization phase (lines 1–2 of Algorithm 1) involves finding a subset of
 538 border points such that at least one point for every class region is selected. Eppstein
 539 observes that this can be achieved by computing the Minimum Spanning Tree (MST)
 540 of P , identifying the edges of the MST that connect points of different classes
 541 (denoted as *bichromatic* edges), and selecting the endpoints of all such edges. This
 542 phase takes $\mathcal{O}(n^2)$ time, but we will prove that it is not necessary.

Algorithm 1: Eppstein's algorithm [8] to find the set of border points of P .

Input: Initial training set P
Output: The set of border points of P

- 1 Let M be the MST of P
- 2 Initialize R with the end points of every bichromatic edge of M
- 3 **foreach** $p \in R$ **do**
 - 4 Let c be p 's class and P_c be the points of P that belong to class c
 - 5 Let S_p be the inverted points of $P \setminus P_c$ around p
 - 6 Find all extreme points of S_p and their corresponding original points E_p
 - 7 $R \leftarrow R \cup E_p$
- 8 **return** R

543 3.2.2 The Search Phase

544 The search phase (lines 3–6 of Algorithm 1) is in charge of finding every
545 remaining border point of P . This phase iterates over all selected points, and for
546 each such point p , it performs what we call the inversion method. This method
547 identifies a subset of border points of P , which are added to R . Once the algorithm
548 has done the inversion method on every point of R , it terminates with the guarantee
549 of having selected every border point of P .

550 3.2.3 The Inversion Method

551 Given any point $p \in P$, the inversion method on p is described in lines 4–6 of
552 Algorithm 1. Let c be p 's class, and P_c be the points of P that belong to class c , the
553 inversion method on p consists of: (i) inverting all points of $P \setminus P_c$ around a ball
554 centered at p (call the set of these inverted points as S_p and include p itself in the
555 set), (ii) computing the set of extreme points of S_p , and finally (iii) returning the set
556 E_p of those points of P that correspond to the extreme points of S_p before inversion.
557 For a detailed description and proof of correctness of this method, we refer the reader
558 to Eppstein's paper [8]. However, for the purposes of this chapter we only need a
559 property presented in Lemma 3 of [8]: the points in E_p reported by the inversion
560 method are the Delaunay neighbors of p with respect to the set $(P \setminus P_c) \cup \{p\}$.

561 Every call of the inversion method takes $\mathcal{O}(nk)$ time by leveraging well-known
562 output-sensitive algorithms for computing extreme points. Given that this method
563 is called exclusively on every border point of the training set, this yields a total of
564 $\mathcal{O}(nk^2)$ time to complete the search phase of the algorithm. Overall, this implies that
565 Eppstein's algorithm computes the entire set of border points of P in $\mathcal{O}(n^2 + nk^2)$
566 worst-case time.

567 3.3 A Simpler Initialization

568 We propose a simple modification to Eppstein's algorithm, which avoids the
569 step of computing the MST of the training set P , along with the subsequent selection

570 of bichromatic edges to produce the initial subset of border points.

571 Instead, we simply start the search process with any arbitrary point of P . The
572 rest of the algorithm remains virtually unchanged (see Algorithm 2 for a formal
573 description). We show that this new approach is not only correct, meaning that
574 it only finds border points of P , but also complete, as all border points of P are
575 eventually found by our algorithm. Additionally, by avoiding the main bottleneck of
576 the original algorithm, our new algorithm computes the same result in $\mathcal{O}(nk^2)$ time,
577 eliminating the $\mathcal{O}(n^2)$ term.

Algorithm 2: New algorithm to find the set of border points of P .

Input: Initial training set P
Output: The set of border points of P

```
1 Let  $s$  be any “seed” point from  $P$ 
2  $R \leftarrow \emptyset$ 
3 foreach  $p \in R \cup \{s\}$  do
4     Let  $c$  be  $p$ ’s class and  $P_c$  be the points of  $P$  that belong to class  $c$ 
5     Let  $S_p$  be the inverted points of  $P \setminus P_c$  around  $p$ 
6     Find all extreme points of  $S_p$  and their corresponding original points  $E_p$ 
7      $R \leftarrow R \cup E_p$ 
8 return  $R$ 
```

578 Before proceeding, it is useful to explore why Eppstein’s algorithm computes
579 the MST of the training set P . First, note that the original algorithm only applies
580 the inversion method on border points of P . In fact, Eppstein’s correctness proof
581 relies on it: Lemma 6 in [8] proves that all points in E_p are border points by assuming
582 that point p is also a border point. From the description of our algorithm, note that
583 we initially apply the inversion method on a “seed” point s , which might not be a
584 border point. Therefore, we need to generalize Lemma 6 in [8] for the case where p is
585 not a border point of P . Additionally, using the points from all bichromatic pairs of
586 the MST of P guarantees that Eppstein’s algorithm starts the search phase with at
587 least one point from every boundary of P . Eppstein’s completeness proof shows that
588 the search phase can then “move along” any given boundary and eventually select
589 all its defining points. We show that the search process is far more powerful, and
590 can even “jump” between nearby boundaries, thus rendering the MST computation
591 unnecessary.

592 The following description outlines the necessary steps to prove both the correct-
593 ness and completeness of our new algorithm, which are unfolded in the rest of this
594 section. (i) By applying the inversion method to any point of P , not necessarily a
595 border point, we must prove that all reported points are border points of P . This is
596 established in Lemma 3.1, generalizing the statement of Lemma 6 of [8] for non-border
597 points. (ii) For any class boundary of P , once the algorithm selects a point from
598 this boundary, we must prove that it will eventually select every other point defining
599 the same boundary. This is originally proved in Lemma 10 [8], however, we provide
600 simpler proofs in Lemmas 3.2 and 3.3. (iii) Given two disconnected boundaries

601 separated by a class region, we must prove that if our algorithm selects a defining
 602 point from one of the boundaries, it will eventually select all defining points from
 603 both boundaries. This is established in Lemma 3.4.

604 All together, these lemmas are used to prove the main result: the correctness,
 605 completeness, and worst-case time complexity of Algorithm 2, as stated in
 606 Theorems 3.5 and 3.6.




Figure 3.2: Example showing the inversion method from any point $p \in P$. On the left, training set P . The middle figure shows every non red point of P , except for p itself, along with a point q selected from the inversion method on p . On the right, we see evidence that q is a border point of P .

607 3.3.1 Correctness Proof

608 **Lemma 3.1.** *Let $p \in P$ be any point of the training set. Then every point selected
 609 using the inversion method on p must be a border point of P .*

610 *Proof.* Let E_p be the points of P corresponding (before inversion) to the extreme
 611 points of S_p . According to Lemma 3 [8], every point in E_p is a neighbor of point
 612 p with respect to the Voronoi Diagram of set $(P \setminus P_c) \cup \{p\}$. This implies that for
 613 every point $q \in E_p$ other than p , there exists a ball such that both p and q are on its
 614 surface and no points of $P \setminus P_c$ lie inside (see Figures 3.2a and 3.2b). We can now
 615 leverage similar techniques to the ones described in [6], to find a “witness” point to
 616 the hypothesis that q must be a border point of P .

617 Recall that the empty ball we just described, as illustrated in Figure 3.2b, is
 618 empty from points of $P \setminus P_c$. However, there might be points of P_c inside. And
 619 moreover, we know that at least one point of P_c , point p , lies on its surface. Now,
 620 let r be the center of this ball, we grow an empty ball, this time with respect to
 621 the entire training set P , such that its center lies on the line \overline{qr} and point q is on
 622 its surface (see Figure 3.2c). This ball will grow until it hits another point \hat{p} of P ,
 623 which we are guaranteed it will be of the same class as point p , and thus, of different

624 class as point q . Finally, we have just found an empty ball with respect to P , which
625 has points q and \hat{p} on its surface, and were the class of both points differ. Therefore,
626 this implies that q is a border point of P . \square




Figure 3.3: By definition, any two adjacent walls w_1 and w_2 of the Voronoi Diagram of P hold the empty ball property with the points that define them. When these walls are part of the class boundaries of P , the points that define them belong to at least two classes.

627 3.3.2 Completeness Proof

628 Before continuing, we need to formally define a few concepts. First, we define
629 a *wall* of P as any $(d - 1)$ -dimensional face of the Voronoi Diagram of P . By known
630 properties of these structures, every wall w is *defined* by two distinct points $p, q \in P$
631 such that any point on w has p and q as its two equidistant nearest-neighbors in the
632 training set. We say two walls are *adjacent* if their intersection is not empty. That
633 is, if there exists a point in \mathbb{R}^d with all the defining points of these two walls as its
634 equidistant nearest-neighbors in P .

635 Additionally, we define a *class boundary* (or just *boundary*) of P as the union
636 of adjacent walls, where each of these walls is defined by two points of different
637 classes. Similarly, we define a *class region* of P as the union of adjacent Voronoi cells
638 whose defining points belong to the same class. Based on these definitions, note that
639 class boundaries are the ones that separate different class regions of P . Figure 3.4
640 illustrates a training set in \mathbb{R}^2 with points of three classes, whose Voronoi Diagram
641 describes five class regions and two class boundaries.

642 **Lemma 3.2.** *Let w_1 and w_2 be two adjacent walls in a class boundary of P . If the
643 algorithm selects one of the points defining one of these walls, it eventually selects
644 the remaining points defining both walls.*

645 *Proof.* Let \mathcal{W} be the set of points defining both walls w_1 and w_2 (see Figure 3.3). By
646 definition, these two walls of the Voronoi Diagram of P are adjacent if there exists

647 an empty ball with all the points of \mathcal{W} on its surface. Knowing these two walls are
 648 part of the class boundaries of P , the set \mathcal{W} must contain at least three points, and
 649 at least two classes.

650 Let p_1 be the first point of \mathcal{W} to be selected by the algorithm. When doing the
 651 inversion method on point p_1 , the algorithm will select all points of \mathcal{W} of different
 652 class than p_1 , of which we know there is at least one. Let p_2 be one such point.
 653 Finally, when doing the inversion method on point p_2 , the algorithm will select
 654 the remaining points of \mathcal{W} of the same class as p_1 . Therefore, all points of \mathcal{W} will
 655 eventually be selected by the algorithm. \square




Figure 3.4: A training set with five class regions (one *blue*, two *red*, and two *yellow* regions), along with two disconnected class boundaries that separate all these regions. On the left, a boundary separating the blue region and the leftmost yellow and red regions. On the right, a boundary that separates the same blue region and the remaining red and yellow regions.

656 **Lemma 3.3.** *Let \mathcal{A} be a class boundary of P , and assume that the algorithm selects
 657 one of the defining points of \mathcal{A} . Then, the algorithm will eventually select all defining
 658 points of \mathcal{A} .*

659 This comes as a direct consequence of Lemma 3.2 and the definition of a class
 660 boundary of the training set P . It remains to show what happens with boundaries
 661 that are disconnected.

662 **Lemma 3.4.** *Let \mathcal{A} and \mathcal{B} be two disconnected boundaries of P , such that there exists
 663 a path in space from \mathcal{A} to \mathcal{B} that is completely contained within one color region.
 664 Without loss of generality, say that every point that defines \mathcal{A} has been selected by the
 665 algorithm. Then, every point that defines \mathcal{B} must also be selected by the algorithm.*

666 *Proof.* Given these two disconnected boundaries \mathcal{A} and \mathcal{B} , we assume there exists
 667 some path \mathcal{P} in \mathbb{R}^d going from a wall of \mathcal{A} to a wall of \mathcal{B} , such that this path
 668 passes exclusively through a single class region (see Figure 3.5a). Without loss of

generality, say this is a *red* class region. Formally, for every point r along \mathcal{P} we know r 's nearest-neighbor in P is red. Additionally, we assume that every border point defining \mathcal{A} is selected by the algorithm. Hence, the proof consists of showing that there exists a sequence of border points $\langle p_1, \hat{p}_1, p_2, \hat{p}_2, \dots, p_m, \hat{p}_m \rangle$ such that (i) p_1 and \hat{p}_m are defining points of \mathcal{A} and \mathcal{B} , respectively, (ii) \hat{p}_i is retrieved by the inversion method on p_i , for every $i \in [1, m]$, and finally (iii) points p_i and \hat{p}_{i-1} are both defining the same boundary, for every $i \in [2, m]$. See Figure 3.5 for a visual description.




Figure 3.5: On the right, two disconnected boundaries \mathcal{A} and \mathcal{B} enclosing a red class region. Thus, there is a path \mathcal{P} completely contained inside such region and connecting both boundaries. Other boundaries can also be enclosing the same region and be near path \mathcal{P} . On the left, we proof that there exists a sequence of points that can be retrieved by calls to the inversion method, such that if points of \mathcal{A} are selected by the algorithm, eventually points of \mathcal{B} will also be selected.

By definition, for every point r along path \mathcal{P} we know r 's nearest-neighbor is a red point. Now, let's delete every red point from consideration, including the ones defining boundaries \mathcal{A} and \mathcal{B} (see Figure 3.5b). This immediately implies that r 's nearest-neighbor just became a non-red border point of P . The fact that r 's new nearest-neighbor is a border point is easy to proof, using similar arguments as the ones laid down in Lemma 3.1. Additionally, these border points could be defining other boundaries apart from \mathcal{A} and \mathcal{B} , as seen in Figure 3.5b.

Let's start moving along the path \mathcal{P} , starting from the end-point of the path that lies on a wall of boundary \mathcal{A} . Then, find all r_i points along the path, where each r_i has two equidistant nearest-neighbors among the remaining non-red points, and both points define two distinct boundaries of P . We say there are m of these points along the path, and denote r_i 's two equidistant nearest-neighbors as $q_{i,1}$ and $q_{i,2}$ for $i \in [1, m]$. Clearly, $q_{i,1}$ and $q_{i-1,2}$ are border points defining the same boundary, for

690 all $i \in [2, m]$. See Figure 3.5b, where the three black points along the path are the r_i
 691 points, and the *yellow* and *blue* points on the surface of the balls centered at each r_i
 692 are the corresponding $q_{i,1}$ and $q_{i,2}$ points.

693 For now, let's fix the analysis on one such r_i point, and consider the ball
 694 centered at r_i with both $q_{i,1}$ and $q_{i,2}$ on its surface. There must exist some other
 695 point $q_{i,3}$ lying inside of r_i 's ball, such that $q_{i,3}$ is one of the deleted red points
 696 defining the same boundary as $q_{i,1}$. It is now easy to see that there exist an empty
 697 ball, with respect to the set $P \setminus P_{\text{red}} \cup \{q_{i,3}\}$, with both $q_{i,3}$ and $q_{i,2}$ on its boundary.
 698 This implies that $q_{i,2}$ is retrieved by the inversion method on $q_{i,3}$. Therefore, let's
 699 add $p_i \leftarrow q_{i,3}$ and $\hat{p}_i \leftarrow q_{i,2}$ to the sequence of points that we are looking for. Repeat
 700 this for every r_i with $i \in [1, m]$ to identify all points in the sequence.

701 Finally, we have the sequence of border points $\langle p_1, \hat{p}_1, p_2, \hat{p}_2, \dots, p_m, \hat{p}_m \rangle$ such
 702 that for any $i \in [1, m]$ assuming that the algorithm selects the points defining the
 703 same boundary as p_i , it will also select \hat{p}_i , and leveraging Lemma 3.3 it will eventually
 704 select all other points defining the same boundary as \hat{p}_i . Given that p_1 and \hat{p}_m are
 705 defining border points of boundaries \mathcal{A} and \mathcal{B} , respectively, and by the assumption
 706 that all points defining \mathcal{A} are selected by the algorithm, we know that eventually, all
 707 points defining \mathcal{B} will be selected too. \square

708 **Theorem 3.5.** *The algorithm selects every border point of P in $\mathcal{O}(nk^2)$ time.*

709 *Proof.* Proving the worst-case time complexity of our algorithm follows directly from
 710 the time complexity of the search phase of Eppstein's algorithm [8]. However, the
 711 correctness and completeness of our algorithm follows from Lemmas 3.1 to 3.4.

712 First, we know by Lemmas 3.1-3.3 that Algorithm 2 will select the defining
 713 border points of at least one class boundary of P . Denote this boundary as \mathcal{A} and
 714 consider any other boundary \mathcal{B} of P . Evidently, we can draw a path \mathcal{P} from \mathcal{A} to
 715 \mathcal{B} , which would generally pass through several class regions. Then, let's split \mathcal{P}
 716 into several subpaths $\mathcal{P}_1, \mathcal{P}_2, \dots, \mathcal{P}_m$ such that each subpath is completely contained
 717 within a single class region. From this, we can directly apply Lemma 3.4 on each of
 718 the intermediate boundaries that "cut" \mathcal{P} into these subpaths. Finally, this implies
 719 that our algorithm will eventually select every defining point of boundary \mathcal{B} , and
 720 similarly, it will do the same with all other boundaries of P . \square

721 **Theorem 3.6.** *Leveraging Chan's algorithm [53] for finding extreme points, the
 722 algorithm selects every border point of P in randomized expected time $\mathcal{O}(nk \log k)$
 723 for $d = 3$, and in*

$$\mathcal{O}\left(k(nk)^{1-\frac{1}{\lfloor d/2 \rfloor + 1}} (\log n)^{\mathcal{O}(1)}\right)$$

724 *time for all constant dimensions $d > 3$.*

725 Just as with Eppstein's original algorithm, we can use Chan's randomized
 726 algorithm [53] for finding extreme points of point sets in \mathbb{R}^d in order to reduce the
 727 expected time complexity of our improved algorithm. The remaining of the proof is
 728 the same as for Theorem 3.5.

729 Chapter 4: Condensation Algorithms

730 4.1 Introduction

731 As previously discussed, preserving the class boundaries of a given training
732 set P is a desirable objective towards the goal of more efficient nearest-neighbor
733 classification. However, the high complexity of computing the entire set of border
734 points of P has motivated the study of alternative approaches for training set
735 reduction, which include the problems of computing either *consistent* or *selective*
736 subsets of P . Thus, even though computing minimum cardinality subsets that are
737 either consistent or selective is known to be NP-hard [10–12], computing non-minimal
738 subsets holding these properties can be done efficiently. Algorithms that compute
739 such subsets are frequently referred to as nearest-neighbor condensation heuristics,
740 as they “heuristically” select points that are close to the class boundaries induced
741 by P . Therefore, such selected subsets induce new class boundaries that resemble
742 the original boundaries, albeit not exactly the same.

743 However, one significant shortcoming in research on practical nearest-neighbor
744 condensation algorithms is the lack of theoretical results on the sizes of their selected
745 subsets, where typically, their performance has been established experimentally. This
746 chapter presents the first theoretical guarantees on the performance of both new and
747 existing condensation algorithms. These results have been published in [54–56].

748 The bounds presented in this chapter are established with respect to the size
749 of two well-known and structured subsets of points: (*i*) the set of all nearest-enemy
750 points of P of size κ , and (*ii*) the set of border points of P of size k . Additionally,
751 some bounds depend on the minimum nearest-enemy distance of P denoted as γ .

Algorithm	Subset size
CNN	$\mathcal{O}(\kappa \log 1/\gamma)$
FCNN	Unbounded w.r.t. k and κ
SFCNN •	$\mathcal{O}(\kappa \log 1/\gamma)$
NET	$\mathcal{O}(\kappa \log 1/\gamma)$
MSS	Unbounded w.r.t. k and κ
RSS •	$\mathcal{O}(\kappa)$
VSS •	$\leq k$

Table 4.1: Contributions on the worst-case analysis on the sizes of subsets selected by different condensation algorithms. Those marked with • are new algorithms.

752 4.2 Lower Bounds

753 If we hope to leverage both k and κ to analyze the worst-case performance of
754 different nearest-neighbor condensation heuristics, we first need to establish realistic
755 expectations on what upper-bounds can be achieved. Therefore, before analyzing
756 each of these algorithms individually, we introduce two lower-bounds on the size of
757 consistent subsets, one in terms of κ and the other in terms of k .

758 First, lets recall the definition of *consistency*. A subset $R \subseteq P$ is said to be
759 *consistent* [9] if and only if for every $p \in P$ its nearest-neighbor in R is of the same
760 class as p . Intuitively, R is consistent if and only if all points of P are correctly
761 classified with the nearest-neighbor rule *w.r.t.* R .

762 **Theorem 4.1.** *There exists a training set $P \subset \mathbb{R}^d$ with κ number of nearest-enemy
763 points, for which any consistent subset has $\Omega(\kappa c^{d-1})$ points, for some constant c .*

764 *Proof.* Lets construct a training set P in d -dimensional Euclidean space, which
765 contains points of two classes: *red* and *blue*. Consider the following arrangement of
766 points: create a red point p , and take *every* point at distance 1 from p as a blue
767 point. Simply, point p plus the points on the surface of the unit ball centered at p .

768 Take any consistent subset of P and consider some point \hat{p} in this subset, along
769 with the bisector between p and \hat{p} . The intersection between this bisector and the unit
770 ball centered at p describes a cap of the ball of height $1/2$. Any point located inside
771 this cap is closer to \hat{p} than p , and hence, correctly classified. Clearly, by definition of
772 consistency, all points in the ball must be covered by at least one cap. By a simple
773 packing argument, we know such a covering needs $\Omega(c^{d-1})$ points, for some constant
774 c . The training set constructed so far has only two nearest-enemy points; *i.e.*, the
775 red point p and the one blue point closest to p (assuming general position). Then,
776 we can repeat this arrangement $\kappa/2$ times using sufficiently separated center points.
777 This gives us a new training set P with κ nearest-enemy points in total, for which
778 any consistent subset has size $\Omega(\kappa c^{d-1})$. \square

779 **Theorem 4.2.** *There exists a training set $P \subset \mathbb{R}^d$ with k number of border points,
780 for which any consistent subset has at least k points.*

781 Basically, Theorem 4.1 shows that the best upper-bound that can proved on
782 the size of the subsets selected by condensation algorithm, in terms of κ , is $\mathcal{O}(kc^{d-1})$.
783 Similarly, Theorem 4.2 proves that in terms of k , the best upper-bound is k itself.
784 Clearly, these results provide a steady floor on which to compare the results presented
785 in the coming sections of this chapter.

786 4.3 Consistent Subsets

787 In this section, we explore several algorithms that compute consistent subsets
788 of P , and present a formal worst-case analysis on the size of the subsets that these
789 algorithms select. Namely, we study the CNN, FCNN, and SFCNN algorithms.

790 4.3.1 CNN

791 The CNN algorithm (or *Condensed Nearest-Neighbor*) was the first algorithm
 792 to be proposed for computing consistent subsets [9]. In fact, it was introduced right
 793 along the definition of consistency in Hart's seminal paper. For more than 30 years,
 794 and until the introduction of the FCNN algorithm, CNN was considered to be the
 795 state-of-the-art among condensation heuristics.

Algorithm 3: CNN

Input: Initial training set P
Output: Consistent subset $R \subseteq P$

```

1  $R \leftarrow \emptyset$ 
2 while true do
3    $R' \leftarrow R$ 
4   foreach  $p_i \in P$ , where  $i = 1 \dots n$  do
5     Let  $\text{nn}(p_i, R)$  be the nearest-neighbor of  $p_i$  w.r.t.  $R$ 
6     if  $l(p_i) \neq l(\text{nn}(p_i, R))$  then
7        $R \leftarrow R \cup \{p_i\}$ 
8   Exit the while loop if  $R = R'$ 
9 return  $R$ 

```

796 This algorithm is fairly simple (see Algorithm 3 for a formal description).
 797 Beginning with an empty R set, it repeatedly scans all the points of P . For each
 798 point $p \in P$, the algorithm checks if its nearest-neighbor within the current R set is
 799 of the same class as p . That is, the algorithm checks if p would be correctly classified
 800 using R . If not, p is added to R . Otherwise, the algorithm continues checking the
 801 next point of P being scanned. Note that P can be scanned multiple times, as
 802 adding a point to R can induce the misclassification of a previously checked point.

803 Evidently, this is a cubic-time algorithm. More specifically, a straightforward
 804 implementation of CNN yields a $\mathcal{O}(n^2m)$ worst-case time complexity, where m
 805 denotes the size of the selected subset. Another characteristic about this algorithm
 806 is that it is *order-dependent*, meaning that the resulting subset depends on the order
 807 in which the points of P were scanned by the algorithm. Unsurprisingly, this is an
 808 undesirable property for practitioners.

809 **Theorem 4.3.** *The CNN algorithm selects $\mathcal{O}\left(\kappa \log \frac{1}{\gamma}\right)$ points.*

810 *Proof.* This follows by a charging argument on every nearest-enemy point in P .
 811 Therefore, consider one such point $p \in \{\text{ne}(r) \mid r \in P\}$ and some value $\sigma \in [\gamma, 1]$,
 812 and define $R_{p,\sigma}$ to be the subset of points selected by CNN whose nearest-enemy is
 813 p , and whose distance to p is between σ and 2σ . That is, $R_{p,\sigma} = \{r \in R \mid \text{ne}(r) =$
 814 $p \wedge d(r, p) \in [\sigma, 2\sigma]\}$. All these subsets define a partitioning of R when considering
 815 all nearest-enemy points of P , and values of $\sigma = \gamma 2^i$ for $i = \{0, 1, 2, \dots, \lceil \log \frac{1}{\gamma} \rceil\}$.

816 Now, let's consider any two points $a, b \in R_{p,\sigma}$ in these subsets. We want to
 817 prove that $d(a, b) \geq \sigma$. Evidently, if these two points belong to different classes, then

818 by definition we have that $d(a, b) \geq \sigma$. Thus, lets consider the case when both points
 819 belong to the same class. Lets assume *w.l.o.g.* that point a was added to R before
 820 point b . Note that when the algorithm selected point b , the following must be true
 821 $d(a, b) \geq d_{\text{ne}}(b, R)$, as otherwise point b would have been correctly classified using R .
 822 Moreover, by our partitioning of R , we also know that $d_{\text{ne}}(b, R) \geq d_{\text{ne}}(b) \geq \sigma$.

823 Thus, we have proved that $d(a, b) \geq \sigma$. By a simple packing argument based
 824 on d -dimensional Euclidean balls, we have that $|R'_{p, \sigma}| \leq 5^d$. Altogether, by counting
 825 over all the $R_{p, \sigma}$ sets for every nearest-enemy in the training set and values of σ , the
 826 size of the subset R selected by CNN is upper-bounded by $|R| \leq \kappa \lceil \log 1/\gamma \rceil 5^{d+1}$.
 827 Assuming d to be constant, this completes the proof. \square

828 Moreover, this analysis is tight. A simple example showcasing this behavior can
 829 be described as follows: consider a training set $P \subset \mathbb{R}$ in 1-dimensional Euclidean
 830 space, and let $\gamma \ll 1$ be a small value. Place a single *red* point in the origin, and
 831 make all the points in the segment $[\gamma, 1]$ *blue* points. Now consider running CNN on
 832 this training set, scanning the points starting with the single red point and following
 833 with the blue points in decreasing order on their coordinate (*i.e.*, in decreasing order
 834 on their distance to the origin, or also, their nearest-enemy distance). It is easy to
 835 see that this would make CNN select $\mathcal{O}(\log 1/\gamma)$ points. This arrangement can be
 836 repeated to increase the value of κ , and can also be discretized to make n finite and
 837 independent from k , κ , and γ .

838 The result of Theorem 4.3 confirms the empirical behavior observed for CNN,
 839 where we see many selected points located far from the boundaries of P (see Fig-
 840 ure 1.1b). Evidently, the many drawbacks of the CNN algorithm (*e.g.*, its empirical
 841 behavior, order-dependence, and cubic worst-case time complexity) has motivated
 842 numerous research on improved approaches towards computing consistent subsets.

843 4.3.2 FCNN

844 The FCNN algorithm (or *Fast CNN*) became the state-of-the-art algorithm
 845 for computing consistent subsets [16] improving over many of CNN's drawbacks.
 846 When introduced, it represented a big leap forward on practical methods for nearest-
 847 neighbor condensation, being among the first worst-case quadratic time algorithms
 848 for this problem, with an implementation that runs in $\mathcal{O}(nm)$ time, where m is the
 849 size of its selected subset. Additionally, its selected subset is order-independent,
 850 which is a highly desirable property among practitioners. More importantly, in
 851 practice, this algorithm significantly outperforms other condensation techniques.

852 This is an iterative algorithm that incrementally builds a consistent subset R
 853 of P (see Algorithm 4). First, it begins by selecting the set of centroids of each class,
 854 and then continues with an iterative process until the subset becomes consistent.
 855 During each iteration, the algorithm identifies all points of P that are misclassified
 856 with the current subset R (*i.e.*, whose nearest-neighbor is of different class), and
 857 adds some of these points to the subset. In particular, for every point p already in
 858 the subset, FCNN selects one *representative* among all the points not yet selected,
 859 whose nearest-neighbor is p , and that belong to a different class than p . That is, the




Figure 4.1: Example of the *voren* function for any point $p \in R$. Light-colored points belong to $P \setminus R$, and points in $\text{voren}(p, R, P)$ are outlined (defined as the enemies of p inside its voronoi cell w.r.t. R).

representative is selected from the set $\text{voren}(p, R, P)$ which is defined as:

$$\text{voren}(p, R, P) = \{q \in P \mid \text{nn}(q, R) = p \wedge l(q) \neq l(p)\}$$

See Figure 4.1 for an illustrative example. Usually, the representative chosen is the one closest to p , although different approaches can be used. Finally, during each iteration, the representative of each point in R is added to the subset (all in a batch operation). This is repeated until no misclassified points are left; *i.e.*, until no point of R has a representative to choose.

Algorithm 4: FCNN

Input: Initial training set P
Output: Consistent subset $R \subseteq P$

```

1  $R \leftarrow \emptyset$ 
2  $S \leftarrow \text{centroids}(P)$ 
3 while  $S \neq \emptyset$  do
4    $R \leftarrow R \cup S$ 
5    $S \leftarrow \emptyset$ 
6   foreach  $p \in R$  do
7      $S \leftarrow S \cup \{\text{rep}(p, \text{voren}(p, R, P))\}$ 
8 return  $R$ 

```

Unfortunately, to the best of our knowledge, no upper-bounds for the selection of FCNN were known before our work. The remaining of this section presents the first results along this line, showing that the selection size of this algorithm cannot be upper-bounded in terms of either κ or k , as stated in Theorems 4.4 and 4.5 respectively.

871 **Theorem 4.4.** For any $0 < \xi < 1$, there exists a training set $P \subset \mathbb{R}^d$ in Euclidean
 872 space, with constant number of classes, for which FCNN selects $\Omega(k + 1/\xi)$ points.

873 *Proof.* Consider the arrangement in Figure 4.2b (left), consisting of points of four
 874 classes. The centroids of the blue, yellow, and red classes are the only points
 875 labeled as such. By placing a sufficient number of black points far at the top of
 876 this arrangement, we avoid their centroid to be any of the three black points in the
 877 arrangement. Beginning with the centroids, the first iteration of FCNN would have
 878 added the points outlined in Figure 4.2b (right). Now each of these points have one
 879 black point inside their Voronoi cells, and therefore, these black points will be the
 880 representatives added in the second iteration. This small example, with $k = 5$, shows
 881 how to force FCNN to add all the border points plus two non-border points. Out
 882 of these two non-border black points, one is the centroid added in the initial step.
 883 The remaining non-border black point, however, was added by the algorithm during
 884 the iterative process. This scheme can be extended to larger values of k , without
 885 increasing the number of classes.




Figure 4.2: Example of a training set $P \subset \mathbb{R}^2$ for which FCNN selects more than k points.

886 The previous is the first building block of the entire training set, shown in
 887 Figure 4.2a. To this “middle” arrangement, we append “side” arrangements of
 888 points, as the one illustrated in Figure 4.2c, which will have similar behavior to the
 889 middle arrangement. This particular side arrangement will be appended to the right
 890 of the middle one, such that the distance between the red points is greater than

891 the distance from the yellow to the red point. Every time we append a new side
 892 arrangement, its blue and red labels are swapped. The arrangements appended to
 893 the left side are simply a horizontal flip of the right arrangement. Now, the behavior
 894 of FCNN in such a setup is illustrated with the arrows in Figure 4.2a. The extreme
 895 point of the previous arrangement adds the yellow point at the center of the current
 896 arrangement, which then adds the red point next to the blue point, as is closer than
 897 the other red point. Next, this red point adds the blue point, and the yellow point
 898 adds the remaining red point. Finally, the Voronoi cells of these points will look as
 899 shown in Figure 4.2c (right), and in the next iteration, the tree black points will be
 900 added.

901 After adding side arrangements as needed (same number of the left and right),
 902 it is easy to show that the centroids are still the tree points in the middle arrangement
 903 and the black point at the top (by adding a sufficient number of black points in the
 904 top cluster). This implies that FCNN will be forced to select more than k points. \square


905 **Theorem 4.5.** *For any $0 < \xi < 1$, there exists a training set $P \subset \mathbb{R}^d$ in Euclidean
 906 space, with constant number of classes, for which FCNN selects $\Omega(\kappa/\xi)$ points.*

907 *Proof.* Without loss of generality, let $\xi = 1/2^t$ for some value $t > 3$, we construct a
 908 training set $P \subset \mathbb{R}^3$ with constant number of classes, and number of nearest-enemy
 909 points κ equal to $\mathcal{O}(1/\xi)$, for which FCNN is forced to select $\mathcal{O}(1/\xi^2)$ points. The
 910 key downside of the algorithm occurs when points are added to the subset in the
 911 same iteration. In general, during any given iteration, the representatives of two
 912 neighboring points in FCNN can be arbitrarily close to each other. This flaw can be
 913 exploited to force the algorithm to add $\mathcal{O}(1/\xi)$ such points.


Intuitively, our constructed training set P consists of several layers of points
 arranged parallel to the xy -plane, and stacked on top of each other around the z -axis
 (see Figure 4.3). Each layer is a disk-like arrangement, formed by a center point and
 points at distance 1 from this center. Thus, define the backbone points of P as the
 center points $c_i = 2i\vec{v}_z$ for $i \geq 0$. We now describe the different arrangements of
 points as follows (see Figure 4.3):

$$\begin{aligned}
 \mathcal{B} &= c_0 \cup \{y_j = c_0 + \vec{v}_x R_z(j\pi/4) \mid j \in [0, \dots, 8]\} \\
 \mathcal{M}_i &= \{c_{2i}, c_{2i+1}, m_i = (c_{2i} + c_{2i+1})/2\} \\
 &\cup \{r_{ij} = c_{2i} + \vec{v}_x R_z(j\pi/2^{1+i}) \mid j \in [0, \dots, 2^{2+i}]\} \\
 &\cup \{b_{ij} = c_{2i+1} + \vec{v}_x R_z(j\pi/2^{1+i}) \mid j \in [0, \dots, 2^{2+i}]\} \\
 &\cup \{w_{ij} = c_{2i+1} + \vec{v}_x R_z((2j+1)\pi/2^{2+i} - \xi^2) \mid j \in [0, \dots, 2^{2+i}]\} \\
 \mathcal{R}_i &= \{c_{2i}, c_{2i+1}\} \\
 &\cup \{r_{ij} = c_{2i} + \vec{v}_x R_z(j2\pi/\xi) \mid j \in [0, \dots, \xi]\} \\
 &\cup \{b_{ij} = c_{2i+1} + \vec{v}_x R_z(j2\pi/\xi) \mid j \in [0, \dots, \xi]\}
 \end{aligned}$$


914 These points belong to one of 11 classes named $\{1, \dots, 8, \text{red, blue, white}\}$.
 915 Then define the labeling function l as follows: $l(c_i)$ is red when i is even and blue
 916 when i is odd, $l(m_i)$ is white, $l(y_j)$ is the j -th class, $l(r_{ij})$ is red, $l(b_{ij})$ is blue, and
 917 $l(w_{ij})$ is white.




(a) Entire arrangement of points, by stacking the different arrangements along the z -axis. The arrows illustrate the selection process by FCNN on a multiplicative arrangement \mathcal{M}_i .



(b) Base arrangement \mathcal{B} . Each point in the circumference belongs to a unique class (here colored in yellow and numbered 1 to 8 for clarity).



(c) A multiplier arrangement \mathcal{M}_i . This forces FCNN to duplicate the number of representatives around the circumference selected on an iteration.



(d) A repetitive arrangement \mathcal{R}_i . This maintains the number of representatives selected by FCNN on each iteration of the algorithm.

Figure 4.3: Example of a training set $P \subset \mathbb{R}^3$ for which FCNN selects $\Omega(\kappa/\xi)$ points.

918 Base arrangement \mathcal{B} : Consists of one single layer of points, with one red center
 919 point c_0 and 8 points y_j in the circumference of the unit disk (parallel to the xy -plane),
 920 each labeled with a unique class j (see Figure 4.3b). The goal of this arrangement is
 921 that each of these points is the centroid of its corresponding class. The centroids of
 922 the blue and white classes can be fixed to be far enough, so we won't consider them
 923 for now. Hence, the first iteration of FCNN will add all the points of \mathcal{B} . In the next
 924 iteration, each of these points will select a representative in the arrangement above.
 925 Clearly, the size of \mathcal{B} is 9, and it contributes with 8 nearest-enemy points in total.

926 Multiplier arrangement \mathcal{M}_i : Our final goal is to have $\mathcal{O}(1/\xi)$ arbitrarily close
 927 points selecting representatives on a single iteration; currently, we only have 9 (the
 928 base arrangement). While this could be simply achieved with $\mathcal{O}(1/\xi)$ points in \mathcal{B}
 929 each with a unique class, we want to use a constant number of classes. Instead, we
 930 use each multiplier arrangement to double the number of representatives selected.
 931 \mathcal{M}_i consists of (1) a layer with a blue center c_{2i} and 2^{2+i} red points r_{ij} around
 932 the unit disk's circumference, (2) a layer with a red center c_{2i+1} and 2^{3+i} blue b_{ij}
 933 and white w_{ij} points around the unit disk's circumference, and (3) a middle white
 934 center point m_i between the red and blue center points (see Figure 4.3c). Suppose at
 935 iteration $3i - 1$ all the points r_{ij} and c_{2i} of the first layer are added as representatives
 936 of the previous arrangement, which is given for \mathcal{M}_1 from the selection of \mathcal{B} . Then,
 937 during iteration $3i$ each r_{ij} adds the point b_{ij} right above, while c_{2i} adds point m_i
 938 (see the red arrows in Figure 4.3a). Finally, during iteration $3i + 1$, m_i adds c_{2i+1} ,
 939 and each b_{ij} adds point w_{ij} as its the closest point inside the voronoi cell of b_{ij} (see
 940 the blue arrows in Figure 4.3a). Now, with all the points of this layer added, each
 941 continues to select points in the following arrangement (either \mathcal{M}_{i+1} or \mathcal{R}_{i+1}). The
 942 size of each \mathcal{M}_i is $3(1 + 2^{2+i}) = \mathcal{O}(2^{3+i})$, and contributes with $3 + 2(2^{2+i}) = \mathcal{O}(2^{3+i})$
 943 to the total number of nearest-enemy points. In order to select $1/\xi = 2^t$ points in a
 944 single iteration, we need to stack \mathcal{M}_i 's for $i \in [1, \dots, t - 3]$.

945 Repetitive arrangement \mathcal{R}_i : Once the algorithm reaches the last multiplier
 946 layer \mathcal{M}_{t-3} , it will select $1/\xi$ points during the following iteration. The repetitive
 947 arrangement allows us to continue adding these many points on every iteration, while
 948 only increasing the number of nearest-enemy points by a constant. This arrangement
 949 consists of (1) a first layer with a blue center c_{2i} surrounded by $1/\xi$ red points r_{ij}
 950 around the unit disk circumference, and (2) a second layer with red center c_{2i+1} and
 951 blue points b_{ij} in the circumference (see Figure 4.3d). Once the first layer is added all
 952 in a single iteration, during the following iteration c_{2i} adds c_{2i+1} , and each r_{ij} adds
 953 b_{ij} . The size of each \mathcal{R}_i is $2(1 + 1/\xi) = \mathcal{O}(1/\xi)$, and it contributes with 4 points to
 954 the total number of nearest-enemy points. Now, we stack $\mathcal{O}(1/\xi)$ such arrangements
 955 \mathcal{R}_i for $i \in [t - 2, \dots, 1/\xi]$, such that we obtain the desired ratio between selected
 956 points and number of nearest-enemy points of the training set.

957 The training set is then defined as $P = \mathcal{B} \bigcup_{i=1}^{t-3} \mathcal{M}_i \bigcup_{i=t-2}^{1/\xi} \mathcal{R}_i \cup \mathcal{F}$, where \mathcal{F}
 958 is a set of points designed to fix the centroids of P . These extra points are located
 959 far enough from the remaining points of P , and are carefully placed such that
 960 the centroids of P are all the points of \mathcal{B} , plus a blue and white point from \mathcal{F} .
 961 Additionally, all the points of \mathcal{F} should be closer to it's corresponding class centroid
 962 than to any enemy centroid, and they should increase the number of nearest-enemy
 963 points by a constant. This can be done with $\mathcal{O}(n)$ extra points.

964 All together, by adding up the corresponding terms, the ratio between the size
 965 of FCNN and κ (the number of nearest-enemy points of P) is $\mathcal{O}(1/\xi)$. Therefore,
 966 there exists a training set in 3-dimensional Euclidean space for which FCNN selects
 967 $\mathcal{O}(\kappa/\xi)$ for any $\xi < 1/8$. \square

968 These results show that adding points in batch on every iteration of FCNN
 969 prevents the algorithm to have any guarantee on the size of its selected subset, in

970 terms of either k or κ . However, this design choice is not key for any of the features
 971 of the algorithm, and can therefore be avoided.

972 **4.3.3 SFCNN**

973 The SFCNN algorithm (or *Single* FCNN) is a modified version of FCNN such
 974 that only *one* single representative is added to the subset R on each iteration of
 975 the algorithm. Surprisingly, such a simple change in the selection process allows
 976 us to successfully analyze the size of SFCNN in terms of κ , and even prove that it
 977 computes a tight approximation of the minimum cardinality consistent subset of
 978 P on general metrics. These results can be seen as corollaries of the more general
 979 Theorems 5.16 and 5.17, which are presented later in Chapter 5.

Algorithm 5: SFCNN

Input: Initial training set P
Output: Consistent subset $R \subseteq P$

```

1  $R \leftarrow \emptyset$ 
2  $S \leftarrow \text{centroids}(P)$ 
3 while  $S \neq \emptyset$  do
4    $R \leftarrow R \cup \{\text{Choose one point of } S\}$ 
5    $S \leftarrow \emptyset$ 
6   foreach  $p \in R$  do
7      $S \leftarrow S \cup \{\text{rep}(p, \text{voren}(p, R, P))\}$ 
8 return  $R$ 
```

980 Following similar arguments to the ones described for the upper-bound proved
 981 on CNN (see Theorem 4.3), the following theorem proofs a comparable upper-bound
 982 for SFCNN in terms of κ and γ .

983 **Theorem 4.6.** *The SFCNN algorithm selects $\mathcal{O}(\kappa \log \frac{1}{\gamma})$ points.*

984 *Proof.* Similarly to the upper-bound proof of CNN, this result follows by a charging
 985 argument on each nearest-enemy point in the training set. Consider one such point
 986 $p \in \{\text{ne}(r) \mid r \in P\}$ and a value $\sigma \in [\gamma, 1]$. We define $R_{p,\sigma}$ to be the subset of points
 987 selected by SFCNN whose nearest-enemy is p , and whose distance to p is between
 988 σ and 2σ . That is, $R_{p,\sigma} = \{r \in R \mid \text{ne}(r) = p \wedge d(r, p) \in [\sigma, 2\sigma]\}$. Clearly, these
 989 subsets define a partitioning of R when considering all nearest-enemy points of P ,
 990 and values of $\sigma = \gamma 2^i$ for $i = \{0, 1, 2, \dots, \lceil \log \frac{1}{\gamma} \rceil\}$.

991 Consider any two points $a, b \in R_{p,\sigma}$ in these subsets. Assume *w.l.o.g.* that
 992 point a was selected by the algorithm before point b (*i.e.*, in a prior iteration). We
 993 show that $d(a, b) \geq \sigma$. By contradiction, assume that $d(a, b) < \sigma$, which immediately
 994 implies that a and b belong to the same class. Moreover, recalling that b 's nearest-
 995 enemy in P is p , at distance $d(b, p) \geq \sigma$, this implies that b is closer to a than to
 996 any enemy in R . Therefore, by the definition of the *voren* function, b could never be
 997 selected by SFCNN, which is a contradiction.

998 This proves that $\mathbf{d}(a, b) \geq \sigma$. Now, just as with CNN, using a simple packing
 999 argument based on d -dimensional Euclidean balls, we have that $|R'_{p,\sigma}| \leq 5^d$. Alto-
 1000 together, by counting over all the $R_{p,\sigma}$ sets for every nearest-enemy in the training set
 1001 and values of σ , the size of R is upper-bounded by $|R| \leq \kappa \lceil \log 1/\gamma \rceil 5^{d+1}$. Assuming
 1002 d to be constant, this completes the proof. \square

1003 4.4 Selective Subsets

1004 Another common criterion used for nearest-neighbor condensation is known as
 1005 *selectiveness* [34]. A subset $R \subseteq P$ is said to be *selective* if and only if for every point
 1006 $p \in P$ its nearest-neighbor in R is closer to p than its nearest-enemy in P . Clearly
 1007 selectiveness implies consistency, as the nearest-enemy distance in R of any point of
 1008 P is at least its nearest-enemy distance in P .

1009 Similarly to the previous section, this section studies several algorithms that
 1010 compute selective subsets of P , along with a formal worst-case analysis on the sizes of
 1011 their selected subsets. Namely, these are the NET, MSS, RSS, and VSS algorithms.

1012 4.4.1 NET

1013 The NET algorithm [36] was proposed as an approximation algorithm for the
 1014 problem of finding minimum cardinality consistent subsets. Their paper also presents
 1015 almost matching hardness lower-bounds for this problem. While this algorithm is
 1016 proposed for computing consistent subsets of P , it can easily be shown that its
 1017 selection is actually selective. Hence, the NET algorithm is described and analyzed
 1018 in this section, and not in section 4.3.

1019 The algorithm is fairly simple, and works by computing a γ -net of P where
 1020 γ is the minimum nearest-enemy distance in P . Evidently, this subset must be
 1021 selective. Moreover, the authors of the paper prove that this algorithm computes a
 1022 tight approximation of the minimum cardinality consistent subsets, being the first
 1023 algorithm to show such guarantees. However, in practice, γ tends to be small, which
 1024 results in subsets of much higher cardinality than needed. To overcome this issue,
 1025 the authors proposed a post-processing pruning technique to further reduce the
 1026 selected subset, which is formally described in Algorithm 6. But even with this extra
 1027 pruning, NET is often outperformed on typical training sets (*w.r.t.* both runtime
 1028 and selection size) by the more practical heuristics studied in this chapter.

1029 **Theorem 4.7.** *The NET+PRUNE algorithm selects $\mathcal{O}(\kappa \log \frac{1}{\gamma})$ points.*

1030 *Proof.* Proving this result follows basically the same arguments described on the
 1031 proofs of Theorems 4.3 and 4.6. Similarly, we must define the sets $R_{p,\sigma}$ as the points
 1032 selected by the NET+PRUNE algorithm whose nearest-enemy is some point $p \in P$,
 1033 and its nearest-enemy distance is between $[\sigma, 2\sigma)$. The main difference lies on the
 1034 lower-bound on the distance between any two points $a, b \in R_{p,\sigma}$ in such subsets. In
 1035 this case, we can only guarantee that $\mathbf{d}(a, b) \geq \sigma/2 - \gamma$. Even though this slightly
 1036 complicates the packing argument, we can still argue that there exists a constant

Algorithm 6: NET+PRUNE

Input: Initial training set P
Output: Selective subset $R \subseteq P$

- 1 $R \leftarrow$ Compute a γ -net of P
- 2 **foreach** $i \in \{1, 0, \dots, \lfloor \log \gamma \rfloor\}$ **do**
- 3 **foreach** $p \in R$ with $d_{\text{ne}}(p, R) \geq 2^{i+1}$ **do**
- 4 $R \leftarrow R \setminus \{q \in R \mid q \neq p \wedge d(p, q) < 2^i - \gamma\}$
- 5 **return** R

1037 c such that $|R_{p,\sigma}| \leq c^d$. Therefore, this implies that NET+PRUNE selects a subset
1038 with at most $\kappa \lceil \log \frac{1}{\gamma} \rceil c^d$ points, which completes the proof. \square

1039 **4.4.2 MSS**

1040 The MSS algorithm (or *Modified Selective Subset*) has been considered a state-
1041 of-the-art algorithm for computing selective subsets [17], due to its good performance
1042 in practice and its $\mathcal{O}(n^2)$ worst-case runtime.

1043 The selection process of the algorithm can be simply described as follows: for
1044 every $p \in P$, MSS selects the point with smallest nearest-enemy distance contained
1045 inside the nearest-enemy ball of p . Clearly, this approach computes a selective subset
1046 of P , which by definition, is order-independent. Unfortunately, the selection criteria
1047 of MSS can be too strict, requiring one particular point to be added for each point
1048 $p \in P$. Note that any point inside the nearest-enemy ball of p suffices for achieving
1049 selectiveness. In practice, this can lead to much larger subsets than needed. This
1050 intuition is formalized in the following theorem, where we show how MSS can select
1051 a subset of unbounded size as a function of either κ or k .

1052 **Theorem 4.8.** *There exists a training set $P \subset \mathbb{R}^d$ in Euclidean space, with constant
1053 number of classes, nearest-enemy points, and border points, such that MSS selects
1054 $\Omega(1/\xi)$ points, for any $0 < \xi < 1$.*


1055 *Proof.* Recall that for each point in P , the MSS algorithm selects the point inside
1056 its nearest-enemy ball with smallest nearest-enemy distance. Given a parameter
1057 $0 < \xi < 1$, we construct a training set in 1-dimensional Euclidean space, as illustrated
1058 in Figure 4.4a. Create two points r_1 and r_2 , and assign them to the class of *red*
1059 points. *w.l.o.g.* the distance between these two points is 1. Let \vec{u} be the unit vector
1060 from r_1 to r_2 , create additional points $b_i = r_1 + \frac{i\xi}{4}\vec{u}$ for $i = \{1, 2, \dots, 3/\xi\}$. Assign
1061 all b_i points to the class of *blue* points. The set of all these points constitute the
1062 training set P . It is easy to prove that P has only four nearest-enemy points and
1063 four border points, corresponding to r_1, r_2, b_1 and $b_{3/\xi}$.

1064 Let's discuss which points are added by MSS for each point in P (see Fig-
1065 ure 4.4b). For points r_1 and r_2 , the only points inside their nearest-enemy balls
1066 are themselves, so both r_1 and r_2 belong to the subset selected by MSS. For points
1067 b_i with $i \leq 2/\xi$, the point with smallest nearest-enemy distance contained inside


Algorithm 7: MSS

Input: Initial training set P
Output: Selective subset $R \subseteq P$

- 1 Let $\{p_i\}_{i=1}^n$ be the points of P sorted increasingly w.r.t. their nearest-enemy distance $d_{ne}(p_i)$
- 2 $R \leftarrow \emptyset$
- 3 $S \leftarrow P$
- 4 **foreach** $p_i \in P$, where $i = 1 \dots n$ **do**
- 5 $add \leftarrow false$
- 6 **foreach** $p_j \in P$, where $j = i \dots n$ **do**
- 7 **if** $p_j \in S \wedge d(p_j, p_i) < d_{ne}(p_j)$ **then**
- 8 $S \leftarrow S \setminus \{p_j\}$
- 9 $add \leftarrow true$
- 10 **if** add **then**
- 11 $R \leftarrow R \cup \{p_i\}$
- 12 **return** R



(a) Initial training set of collinear points, where both the number of nearest-enemy points and the number of border points equal to 4. That is, $\kappa = k = 4$.



(b) Subset of points computed by MSS from the original training set (fully colored points belong to the subset, while faded points do not). The size of the subset is $\Omega(1/\xi)$.

Figure 4.4: Unbounded example for MSS with respect to κ and k .

their nearest-enemy ball is b_1 , which is also added to the subset. Now, consider the points b_i with $2/\xi < i < 5/2\xi$. Let $j = i - 2/\xi$, it is easy to prove that the point with smallest nearest-enemy distance inside the nearest-enemy ball of b_i is b_{2j+1} (see Figure 4.4b). Therefore, this implies that the number of points selected by MSS equals $5/2\xi - 2/\xi = 1/2\xi = \Omega(1/\xi)$. \square

1073 4.4.3 RSS

1074 We propose the RSS algorithm (or *Relaxed Selective Subset*) with the idea
 1075 of relaxing the selection process of MSS, while still computing a selective subset.
 1076 For any given point $p \in P$ in the training set, while MSS requires to add the point
 1077 with smallest nearest-enemy distance inside the nearest-enemy ball of p , in RSS
 1078 any point inside the nearest-enemy ball p suffices. The idea is rather simple (see
 1079 Algorithm 8). Points of P are examined in increasing order with respect to their
 1080 nearest-enemy distance, and we add any point whose nearest-enemy ball contains
 1081 no point previously added by the algorithm. This tends to select points close to
 1082 the decision boundaries of P (see Figure 1.1g), as points far from the boundary are
 1083 examined later in the selection process, and are more likely to already contain points
 1084 inside their nearest-enemy ball.

Algorithm 8: RSS

Input: Initial training set P
Output: Selective subset $R \subseteq P$

- 1 $R \leftarrow \emptyset$
- 2 Let $\{p_i\}_{i=1}^n$ be the points of P sorted increasingly w.r.t. their nearest-enemy
 distance $d_{ne}(p_i)$
- 3 **foreach** $p_i \in P$, where $i = 1 \dots n$ **do**
- 4 **if** $d_{nn}(p_i, R) \geq d_{ne}(p_i)$ **then**
- 5 $R \leftarrow R \cup \{p_i\}$
- 6 **return** R

1085 Just like MSS, RSS is order-independent and computes a selective subset of
 1086 P in $\mathcal{O}(n^2)$ worst-case time. Its selectiveness is evident, as every point in P is
 1087 either added to the subset, or has a point in the subset inside its nearest-enemy ball.
 1088 Its order-independence follows from the initial sorting step of the algorithm. Now,
 1089 the time complexity of RSS can be analyzed as follows. The initial step requires
 1090 $\mathcal{O}(n^2)$ time for computing the nearest-enemy distances of each point in P , plus
 1091 additional $\mathcal{O}(n \log n)$ time for sorting the points according to such distances. The
 1092 main loop iterates through each point in P , and searches their nearest-neighbor in
 1093 the current subset, incurring into additional $\mathcal{O}(n^2)$ time using a simple linear search.
 1094 All together, the worst-case time complexity of the algorithm is quadratic.

1095 **Theorem 4.9.** *The RSS algorithm selects $\mathcal{O}(\kappa)$ points.*

1096 *Proof.* The proof follows by a charging argument on each nearest-enemy point of
 1097 P . Consider a nearest-enemy point $p \in P$, and let R_p be the set of points selected
 1098 by RSS such that p is their nearest-enemy. Let $a, b \in R_p$ be two such points, and
 1099 w.l.o.g. say that $d_{ne}(a) \leq d_{ne}(b)$. By construction of the algorithm, we also know
 1100 that $d(a, b) \geq d_{ne}(b)$. Now, consider the triangle $\triangle pab$. Clearly, \overline{pa} is the largest side
 1101 of the triangle, making the angle $\angle apb \geq \pi/3$. This means that the angle between
 1102 any two points in R_p with respect to p is at least $\pi/3$.

1103 By a standard packing argument, this implies that $|R_p| = \mathcal{O}((3/\pi)^{d-1})$. Finally,
 1104 we obtain that the number of points selected by RSS is $\sum_p |R_p| = \kappa \mathcal{O}((3/\pi)^{d-1})$. \square

1105 Additionally, we prove that the size of the subset selected by RSS can be
 1106 bounded as a function of κ and the dimensionality of P . Moreover, we show that
 1107 RSS computes an approximation of both the consistent and selective subsets of
 1108 minimum cardinality. These results come as corollaries of Theorems 5.10 and 5.11,
 1109 which will be described in later sections. However, different parameters from κ can be
 1110 used to bound the selection size of condensation algorithms: consider k , the number
 1111 of border points in the training set P . From the example illustrated in Figure 4.5,
 1112 we know that RSS can select more points than k (see Figure 4.5b). Repeating such
 1113 arrangement forces RSS to select $\Omega(k + 1/\xi)$ points.




Figure 4.5: Example where RSS selects $k + 1$ points.


1114 4.4.4 VSS

1115 We now propose the VSS algorithm (or *Voronoi Selective Subset*). This new
 1116 algorithm comes as a modification of RSS, based on an observation from the proof
 1117 of the following lemma. As stated in Chapter 2, we are able to prove that every
 1118 nearest-enemy is also a border point of P , formalized here.

1119 **Lemma 4.10.** *Any nearest-enemy point of a point in P is also a border point of P .*

1120 *Proof.* Take any point $p \in P$. Consider the *empty* ball of maximum radius, tangent
 1121 to point $\text{ne}(p)$, and with center in the line segment between p and $\text{ne}(p)$. Being
 1122 maximal, this ball is tangent to another point $p^* \in P$ (see Figure 4.6a). Clearly, p^*
 1123 is inside the nearest-enemy ball of p , which implies that p and p^* belong to the same

1124 class, making p^* and $\text{ne}(p)$ enemies. By the empty ball property, this means that
 1125 both p^* and $\text{ne}(p)$ are border points of P . \square



(a) The largest empty ball tangent to $\text{ne}(p)$ and center in $p \text{ ne}(p)$, is also tangent to some point p^* , making p^* and $\text{ne}(p)$ border points.
 (b) How to compute the radius of a ball with center in the line segment between p and $\text{ne}(p)$, and tangent to both $\text{ne}(p)$ and p' .

Figure 4.6: Relation between nearest-enemy points and border points.

1126 Therefore, from Lemma 4.10 we know that in d -dimensional Euclidean space,
 1127 the number of nearest-enemy points of P is at most the number of border points of
 1128 P . That is, $\kappa \leq k$. Moreover, this result can be easily extended to ℓ_p metric spaces
 1129 with $p \geq 2$. While this result implies an easy extension of the upper-bounds found
 1130 for RSS, CNN, SFCNN, and NET+PRUNE, now in terms of k instead of κ , it is
 1131 unclear if the other factors in those upper-bounds can be improved.

1132 Coming back to VSS, this proof opens an alternative idea for condensation.
 1133 In order to prove Lemma 4.10, we show that there exist at least one border point
 1134 inside the nearest-enemy ball of any point $p \in P$. Therefore, by selecting such border
 1135 points, we can guarantee that the resulting subset is selective and its size is $\leq k$.

1136 We call this algorithm VSS (see Algorithm 9 for a formal description). Essentially,
 1137 we show this algorithm computes a selective subset of P of size at most k . By
 1138 construction, for any point in $p \in P$ the algorithm selects one border point inside
 1139 the nearest-enemy ball of p , which implies that the resulting subset is selective, and
 1140 contains no more than k points.

1141 **Theorem 4.11.** *The VSS algorithm selects at most k points.*

1142 Finally, we describe an efficient implementation of VSS, showing this algorithm
 1143 can be implemented to run in quadratic time. Recall that for every point $p \in P$, the
 1144 algorithm selects a border point inside its nearest-enemy ball. *w.l.o.g.* implement
 1145 VSS to compute the point p^* that minimizes the radius of an empty ball tangent

Algorithm 9: VSS

Input: Initial training set P
Output: Selective subset $R \subseteq P$

- 1 $R \leftarrow \emptyset$
- 2 Let $\{p_i\}_{i=1}^n$ be the points of P sorted increasingly w.r.t. their nearest-enemy distance $d_{ne}(p_i)$
- 3 **foreach** $p_i \in P$, where $i = 1 \dots n$ **do**
- 4 **if** $d_{ne}(p_i, R) \geq d_{ne}(p_i)$ **then**
- 5 Find a border point that lies inside the nearest-enemy ball of p_i and add it to R
- 6 **return** R

1146 to both $ne(p)$ and p^* , and center in the line segment between p and $ne(p)$. For any
1147 given point p' inside the nearest-enemy ball of p , denote $r(p, p')$ to be the radius
1148 of the ball tangent to p' and $ne(p)$ and center in the line segment between p and
1149 $ne(p)$. As illustrated in Figure 4.6b, let vectors $\vec{u} = \frac{p - ne(p)}{\|p - ne(p)\|}$ and $\vec{v} = p' - ne(p)$,
1150 the radius of this ball can be derived from the formula $r(p, p') = \|\vec{v} + r(p, p')\vec{u}\|$ as
1151 $r(p, p') = \vec{v} \cdot \vec{v} / 2\vec{u} \cdot \vec{v}$. As p^* is defined as the point that minimizes such radius, a
1152 simple scan over the points of P suffices to identify the corresponding p^* for any
1153 point $p \in P$. Therefore, this implies that VSS can be computed in $\mathcal{O}(n^2)$ worst-case
1154 time.

1155 4.5 Experimental Comparison

1156 Historically, the importance of some condensation algorithms rely on their
1157 performance in practice, despite the lack of theoretical guarantees. Therefore, a
1158 natural question is how the algorithms proposed in this chapter compare to existing
1159 ones when evaluated in real-world training sets.

1160 Thus, to get a clearer impression of the relevance of these results in practice, we
1161 performed experimental trials on several training sets, both synthetically generated
1162 and widely used benchmarks. First, we consider 21 training sets from the UCI
1163 *Machine Learning Repository*¹ which are commonly used in the literature to evaluate
1164 condensation algorithms [18]. These consist of a number of points ranging from 150
1165 to 58000, in d -dimensional Euclidean space with d between 2 and 64, and 2 to 26
1166 classes. We also generated some synthetic training sets, containing 10^5 uniformly
1167 distributed points, in 2 to 3 dimensions, and 3 classes. All training sets used in
1168 these experimental trials are summarized in Table 4.2. The implementation of the
1169 algorithms, training sets used, and raw results, are publicly available².

1170 We test seven different condensation algorithms, namely CNN, FCNN, SFCNN,
1171 MSS, RSS, VSS, and NET. To compare their results, we consider their runtime and

¹<https://archive.ics.uci.edu/ml/index.php>


²<https://github.com/afloresv/nnc/>

Training set	n	d	c	κ (%)
banana	5300	2	2	811 (15.30%)
cleveland	297	13	5	125 (42.09%)
glass	214	9	6	87 (40.65%)
iris	150	4	3	20 (13.33%)
iris2d	150	2	3	13 (8.67%)
letter	20000	16	26	6100 (30.50%)
magic	19020	10	2	5191 (27.29%)
monk	432	6	2	300 (69.44%)
optdigits	5620	64	10	1245 (22.15%)
pageblocks	5472	10	5	429 (7.84%)
penbased	10992	16	10	1352 (12.30%)
pima	768	8	2	293 (38.15%)
ring	7400	20	2	2369 (32.01%)
satimage	6435	36	6	1167 (18.14%)
segmentation	2100	19	7	398 (18.95%)
shuttle	58000	9	7	920 (1.59%)
thyroid	7200	21	3	779 (10.82%)
twonorm	7400	20	2	1298 (17.54%)
wdbc	569	30	2	123 (21.62%)
wine	178	13	3	37 (20.79%)
wisconsin	683	9	2	35 (5.12%)
v-100000-2-3-15	100000	2	3	1909 (1.90%)
v-100000-2-3-5	100000	2	3	788 (0.78%)
v-100000-3-3-15	100000	3	3	7043 (7.04%)
v-100000-3-3-5	100000	3	3	3738 (3.73%)
v-100000-4-3-15	100000	4	3	13027 (13.02%)
v-100000-4-3-5	100000	4	3	10826 (10.82%)
v-100000-5-3-15	100000	5	3	22255 (22.25%)
v-100000-5-3-5	100000	5	3	17705 (17.70%)


Table 4.2: Training sets used to evaluate the performance of condensation algorithms. Indicates the number of points n , dimensions d , classes c , nearest-enemy points κ (also in percentage *w.r.t.* n).

the size of the selected subset. Clearly, these values might differ greatly on training sets whose size are too distinct. Therefore, before comparing the raw results, these are normalized. The runtime of an algorithm for a given training set is normalized by dividing it by n , the size of the training set. The size of the selected subset is normalized by dividing it by κ , the number of nearest-enemy points in the training set, which characterizes the complexity of the boundaries between classes.

Figures 4.7a and 4.7b summarize the experimental results. Evidently, the performance of SFCNN is equivalent to the original FCNN algorithm, both in terms of runtime and the size of their selected subsets, showing that the proposed modification does not affect the behavior of the algorithm in real-world training sets. Both FCNN and SFCNN outperform other condensation algorithms in terms of runtime, while their subset size is comparable in all cases, with the exception of the NET algorithm.



(a) Running time.



(b) Size of the selected subsets.

Figure 4.7: Evaluating the studied condensation algorithms: CNN, FCNN, SFCNN, NET, MSS, RSS, and VSS.

1185 Chapter 5: ε -Coresets

1186 5.1 Introduction

1187 There are obvious parallels between the problem of nearest-neighbor condensa-
1188 tion and the concept of *coresets* in geometric approximation [57–60]. Intuitively, a
1189 *coreset* is small subset of the original data, that well approximates some statistical
1190 properties of the original set. Coresets have previously been applied to many prob-
1191 lems in machine learning, such as clustering and neural network compression [61–64].
1192 Additionally, coresets have been used towards achieving more efficient classification
1193 techniques, as evidenced by the recent results on coresets for the SVM classifier [65].

1194 While the other chapters of this book deal with training sets in d -dimensional
1195 Euclidean space (*i.e.*, $P \subset \mathbb{R}^d$), this chapter deals with the more generic concept
1196 of metric spaces. Therefore, we assume a training set P consisting of n points in a
1197 metric space $(\mathcal{X}, \mathbf{d})$, with domain \mathcal{X} and distance function $\mathbf{d} : \mathcal{X}^2 \rightarrow \mathbb{R}^+$. Evidently,
1198 \mathbf{d} must hold the properties of *identity*, *symmetry*, and *triangle inequality*. The rest
1199 remains the same. That is, P is partitioned into a finite set of *classes* by associating
1200 each point $p \in P$ with a *label* $l(p)$, indicating the class to which it belongs.

1201 Given an *unlabeled* query point $q \in \mathcal{X}$, the *exact* nearest-neighbor rule predicts
1202 q 's class using the class of its closest point in P according to metric \mathbf{d} . That is, it
1203 assigns q to the class $l(\text{nn}(q))$, where $\text{nn}(q) = \arg \min_{p \in P} \mathbf{d}(q, p)$.

1204 So far, in Chapters 3 and 4 we have assumed that q 's nearest-neighbor is always
1205 computed exactly. However, it is common practice that these queries are instead
1206 computed approximately, leveraging known techniques for efficient nearest-neighbor
1207 search like Approximate Voronoi Diagrams [19, 20], Locality-Sensitive Hashing [21],
1208 and Hierarchical Navigable Small Worlds graphs [22]. In this context, we are given
1209 an approximation parameter $\varepsilon \in [0, 1]$ and an unlabeled query point $q \in \mathcal{X}$, and the
1210 ε -approximate nearest-neighbor rule assigns q to the class of some point p , where
1211 p is any point of P such that $\mathbf{d}(q, p) \leq (1 + \varepsilon) \mathbf{d}_{\text{nn}}(q)$. Therefore, in both this and
1212 the following chapters, we discuss different techniques to achieve boundary-sensitive
1213 approaches (*i.e.*, dependent on κ and k instead of n) for approximate nearest-neighbor
1214 classification.

1215 This chapter presents the first approach proposed to compute coresets for nearest-
1216 neighbor classification, leveraging its resemblance to the problem of nearest-neighbor
1217 condensation. These results have been published in [66].

1218 Preliminaries

1219 Similarly to Chapter 4, we define notation relevant to understanding the results
1220 presented. Given any point $q \in \mathcal{X}$ in the metric space, its nearest-neighbor, denoted
1221 $\text{nn}(q)$, is the closest point of P according to the distance function \mathbf{d} . The distance
1222 from q to its nearest-neighbor is denoted by $\mathbf{d}_{\text{nn}}(q, P)$, or simply $\mathbf{d}_{\text{nn}}(q)$ when P is
1223 clear. Given a point $p \in P$ from the training set its nearest-neighbor in P is p itself.
1224 Additionally, any point of P whose label differs from p 's is called an *enemy* of p . The
1225 closest such point is called p 's *nearest-enemy*, and the distance to this point is called
1226 p 's *nearest-enemy distance*. These are denoted as $\text{ne}(p)$ and $\mathbf{d}_{\text{ne}}(p, P)$, respectively.
1227 Evidently, we can simplify $\mathbf{d}_{\text{ne}}(p, P)$ as $\mathbf{d}_{\text{ne}}(p)$ when it is clear.

1228 Unsurprisingly, the size of a coresnet for nearest-neighbor classification should
1229 depend on the spatial characteristics of the points of different classes in the training
1230 set. For example, it should be much easier to find a small coresnet for two spatially
1231 well separated clusters than for two classes that have a high degree of overlap. Again,
1232 we use the number of nearest-enemy points of P , denoted as κ , to model the intrinsic
1233 complexity of the problem of nearest-neighbor classification.

1234 Other intrinsic characteristics of P will also come handy when describing these
1235 coresnets. Through a suitable uniform scaling, we may assume that the *diameter* of
1236 P (that is, the maximum distance between any two points in the training set) is 1.
1237 Then, the *spread* of P , denoted as Δ , is the ratio between the largest and smallest
1238 distances in P . Similarly, recall the definition of the *margin* of P , denoted γ , as the
1239 smallest nearest-enemy distance in P . Clearly, we can see that $1/\gamma \leq \Delta$.

1240 Additionally, a metric space $(\mathcal{X}, \mathbf{d})$ is said to be *doubling* [5] if there exist some
1241 bounded value λ such that any metric ball of radius r can be covered with at most
1242 λ metric balls of radius $r/2$. Its *doubling dimension* is the base-2 logarithm of λ ,
1243 denoted as $\text{ddim}(\mathcal{X}) = \log \lambda$. Throughout this chapter, we assume that $\text{ddim}(\mathcal{X})$
1244 is a constant, which means that multiplicative factors depending on $\text{ddim}(\mathcal{X})$ may
1245 be hidden in our asymptotic notation. Many natural metric spaces of interest are
1246 doubling, including d -dimensional Euclidean space whose doubling dimension is $\Theta(d)$.
1247 An important property of doubling spaces, that will come useful throughout this
1248 chapter, is that for any subset $R \subseteq \mathcal{X}$ with spread Δ_R , the size of R is upper-bounded
1249 by $|R| \leq \lceil \Delta_R \rceil^{\text{ddim}(\mathcal{X})+1}$.

1250 Contributions

1251 In this chapter, we introduce the concept of a coresnet for classification with the
1252 nearest-neighbor rule, which provides approximate guarantees on correct classification
1253 for all query points. We demonstrate their existence, analyze their size, and discuss
1254 approaches for their efficient computation.

1255 We say that a subset $R \subseteq P$ is an ε -coresnet for the nearest-neighbor rule on P ,
1256 if and only if for every query point $q \in \mathcal{X}$, the class of its exact nearest-neighbor in
1257 R is the same as the class of some ε -approximate nearest-neighbor of q in P (see
1258 Section 5.2 for definitions). Recalling the concepts of κ and γ introduced in the
1259 preliminaries, here is our main result:

1260 **Theorem 5.1.** *Given a training set P in a doubling metric space $(\mathcal{X}, \mathbf{d})$, there exist*
1261 *an ε -coreset for the nearest-neighbor rule of size $\mathcal{O}(\kappa \log \frac{1}{\gamma} (1/\varepsilon)^{\text{ddim}(\mathcal{X})+1})$, and this*
1262 *coreset can be computed in subquadratic worst-case time.*

1263 The following summarizes of the principal results presented in the remaining
1264 sections, which all together are leveraged to prove the main theorem.

- 1265 • We extend the criteria used for nearest-neighbor condensation, and identify
1266 sufficient conditions to guarantee the correct classification of any query point
1267 after condensation. These conditions are the ones that describe our coresets
1268 construction.
- 1269 • We prove that finding minimum-cardinality subsets with this new criteria is
1270 NP-hard. Moreover, we prove it is even hard to approximate within practical
1271 factors.
- 1272 • We provide quadratic-time approximation algorithms with provable upper-
1273 bounds on the sizes of their selected subsets, and we show that the running time
1274 of one such algorithm can be improved to be subquadratic. This subquadratic-
1275 time algorithm is the first with such worst-case runtime for the problem of
1276 nearest-neighbor condensation.

1277 5.2 Coreset Characterization

1278 In practice, many applications usually rely its efficiency on computing nearest-
1279 neighbors not exactly, but rather approximately. Given an approximation parameter
1280 $\varepsilon \geq 0$, an ε -approximate nearest-neighbor or ε -ANN query returns any point whose
1281 distance from the query point is within a factor of $(1 + \varepsilon)$ times the true nearest-
1282 neighbor distance.

1283 Intuitively, a query point should be easier to classify if its nearest-neighbor is
1284 significantly closer than its nearest-enemy. This intuition can be formalized through
1285 the concept of the *chromatic density* [51] of a query point $q \in \mathcal{X}$ with respect to a
1286 set $R \subseteq P$, defined as:

$$\delta(q, R) = \frac{\mathbf{d}_{\text{ne}}(q, R)}{\mathbf{d}_{\text{nn}}(q, R)} - 1. \quad (5.1)$$

1287 Clearly, if $\delta(q, R) > \varepsilon$ then q will be correctly classified¹ by an ε -ANN query
1288 over R , as all possible candidates for the approximate nearest-neighbor belong to
1289 the same class as q 's true nearest-neighbor. However, as evidenced in Figures 5.1a
1290 and 5.1b, one side effect of existing condensation algorithms is a significant reduction
1291 in the chromatic density of query points. Consequently, we propose new criteria
1292 and algorithms that maintain high chromatic densities after condensation, which are
1293 then leveraged to build coresets for the nearest-neighbor rule.

¹By *correct classification*, we mean that the classification is the same as the classification that results from applying the nearest-neighbor rule exactly on the entire training set P .




Figure 5.1: Heatmap of *chromatic density* values of points in \mathbb{R}^2 w.r.t. the subsets computed by different condensation algorithms: FCNN, RSS, and α -RSS (see Figure 1.1). *Yellow* • corresponds to chromatic density values ≥ 0.5 , while *blue* • corresponds to 0. Evidently, α -RSS helps maintaining high chromatic density values when compared to standard condensation algorithms.

1294 5.2.1 Approximation-Sensitive Condensation

1295 The decision boundaries of the nearest-neighbor rule (that is, points q such that
 1296 $d_{\text{ne}}(q, P) = d_{\text{nn}}(q, P)$) are naturally characterized by points that separate clusters of
 1297 points of different classes. As illustrated in Figures 1.1c-1.1g, condensation algorithms
 1298 tend to select such points. However, this behavior implies a significant reduction of
 1299 the chromatic density of query points that are far from such boundaries, as can be
 1300 seen in Figures 5.1a-5.1b with the selection of the FCNN and RSS algorithms.

1301 A natural way to define an approximate notion of consistency is to ensure that
 1302 all points in P are correctly classified by ANN queries over the condensed subset R .
 1303 Given a condensation parameter $\alpha \geq 0$, we define a subset $R \subseteq P$ to be:

1304 α -consistent if $\forall p \in P, d_{\text{nn}}(p, R) < d_{\text{ne}}(p, R)/(1 + \alpha)$.

1305 α -selective if $\forall p \in P, d_{\text{nn}}(p, R) < d_{\text{ne}}(p, P)/(1 + \alpha)$.

1306 It is easy to see that the standard forms arise as special cases when $\alpha = 0$.
 1307 These new condensation criteria imply that $\delta(p, R) > \alpha$ for every $p \in P$, meaning
 1308 that p is correctly classified using an α -ANN query on R . Note that any α -selective
 1309 subset is also α -consistent. Such subsets always exist for any $\alpha \geq 0$ by taking
 1310 $R = P$. Current condensation algorithms cannot guarantee either α -consistency or
 1311 α -selectiveness unless α is equal to zero. Therefore, the central algorithmic challenge
 1312 is how to efficiently compute such sets whose sizes are significantly smaller than P .
 1313 We propose new algorithms to compute such subsets, which showcase how to maintain
 1314 high chromatic density values after condensation, as evidenced in Figures 5.1c and
 1315 5.1d. This empirical evidence is matched with theoretical guarantees for α -consistent
 1316 and α -selective subsets, described in the following section.

₁₃₁₇ 5.2.2 Guarantees on Classification Accuracy

₁₃₁₈ These newly defined criteria for nearest-neighbor condensation enforce lower-
₁₃₁₉ bounds on the chromatic density of any point of P after condensation. However, this
₁₃₂₀ doesn't immediately imply having similar lower-bounds for unlabeled query points of
₁₃₂₁ \mathcal{X} . In this section, we prove useful bounds on the chromatic density of query points,
₁₃₂₂ and characterize sufficient conditions to correctly classify some of these query points
₁₃₂₃ after condensation.

₁₃₂₄ Intuitively, the chromatic density determines how easy it is to correctly classify
₁₃₂₅ a query point $q \in \mathcal{X}$. We show that the “ease” of classification of q after condensation
₁₃₂₆ (*i.e.*, $\delta(q, R)$) depends on both the condensation parameter α , and the chromatic
₁₃₂₇ density of q before condensation (*i.e.*, $\delta(q, P)$). This result is formalized in the
₁₃₂₈ following lemma:

Lemma 5.2. *Let $q \in \mathcal{X}$ be a query point, and R an α -consistent subset of P , for $\alpha \geq 0$. Then, q ’s chromatic density with respect to R is:*

$$\delta(q, R) > \frac{\alpha \delta(q, P) - 2}{\delta(q, P) + \alpha + 3}.$$

₁₃₂₉ *Proof.* The proof follows by analyzing q ’s nearest-enemy distance in R . To this end,
₁₃₃₀ consider the point $p \in P$ that is q ’s nearest-neighbor in P . There are two possible
₁₃₃₁ cases:

₁₃₃₂ Case 1: If $p \in R$, clearly $d_{nn}(q, R) = d_{nn}(q, P)$. Additionally, it is easy to show that af-
₁₃₃₃ ter condensation, q ’s nearest-enemy distance can only increase: *i.e.*, $d_{ne}(q, P) \leq$
₁₃₃₄ $d_{ne}(q, R)$. This implies that $\delta(q, R) \geq \delta(q, P)$.

₁₃₃₅ Case 2: If $p \notin R$, we can upper-bound q ’s nearest-neighbor distance in R as follows:

Since R is an α -consistent subset of P , we know that there exists a point $r \in R$ such that $d(p, r) < d_{ne}(p, R)/(1+\alpha)$. By the triangle inequality and the definition of nearest-enemy, $d_{ne}(p, R) \leq d(p, ne(q, R)) \leq d(p, q) + d_{ne}(q, R)$. Additionally, applying the definition of chromatic density on q and knowing that $d_{ne}(q, P) \leq d_{ne}(q, R)$, we have $d(p, q) = d_{nn}(q, P) \leq d_{nn}(q, R) = d_{ne}(q, R)/(1 + \delta(q, P))$. Therefore:

$$\begin{aligned} d_{nn}(q, R) &\leq d(q, r) \leq d(q, p) + d(p, r) \\ &< d(q, p) + \frac{d(q, p) + d_{ne}(q, R)}{1 + \alpha} \leq \left(\frac{\delta(q, P) + \alpha + 3}{(1 + \alpha)(1 + \delta(q, P))} \right) d_{ne}(q, R). \end{aligned}$$

₁₃₃₆ Finally, from the definition of $\delta(q, R)$, we have:

$$\delta(q, R) = \frac{d_{ne}(q, R)}{d_{nn}(q, R)} - 1 > \frac{(1 + \alpha)(1 + \delta(q, P))}{\delta(q, P) + \alpha + 3} - 1 = \frac{\alpha \delta(q, P) - 2}{\delta(q, P) + \alpha + 3}. \quad \square$$

₁₃₃₈ The above result can be leveraged to define a coresnet, in the sense that an exact
₁₃₃₉ result on the coresnet corresponds to an approximate result on the original set. As
₁₃₄₀ previously defined, we say that a set $R \subseteq P$ is an ε -coresnet for the nearest-neighbor

1341 rule on P , if and only if for every query point $q \in \mathcal{X}$, the class of q 's exact nearest-
1342 neighbor in R is the same as the class of any of its ε -approximate nearest-neighbors
1343 in P .

1344 **Lemma 5.3.** *Any ε -coreset for the nearest-neighbor rule is an α -consistent subset,*
1345 *for $\alpha \geq 0$.*

1346 *Proof.* Consider any ε -coreset $C \subseteq P$ for the nearest-neighbor rule on P . Since the
1347 approximation guarantee holds for any point in \mathcal{X} , it holds for any $p \in P \setminus C$. We
1348 know p 's nearest-neighbor in the original set P is p itself, thus making $d_{nn}(p, P)$
1349 zero. This implies that p must be correctly classified by a nearest-neighbor query
1350 on C , that is, $d_{nn}(p, C) < d_{ne}(p, C)$, which is the definition of α -consistency for any
1351 $\alpha \geq 0$. \square

1352 **Theorem 5.4.** *Any $2/\varepsilon$ -selective subset is an ε -coreset for the nearest-neighbor rule.*

1353 *Proof.* Let R be an α -selective subset of P , where $\alpha = 2/\varepsilon$. Consider any query
1354 point $q \in \mathcal{X}$ in the metric space. It suffices to show that its nearest-neighbor in R is
1355 of the same class as any ε -approximate nearest-neighbor in P . To this end, consider
1356 q 's chromatic density with respect to both P and R , denoted as $\delta(q, P)$ and $\delta(q, R)$,
1357 respectively. We identify two cases:

1358 Case 1 (Correct-Classification guarantee): If $\delta(q, P) \geq \varepsilon$.

1359 Consider the bound derived in Lemma 5.2. Since $\alpha \geq 0$, and by our assumption
1360 that $\delta(q, P) \geq \varepsilon > 0$, setting $\alpha = 2/\varepsilon$ implies that $\delta(q, R) > 0$. This means that
1361 the nearest-neighbor of q in R belongs to the same class as the nearest-neighbor
1362 of q in P . Intuitively, this guarantees that q is correctly classified by the
1363 nearest-neighbor rule in R .

1364 Case 2 (ε -Approximation guarantee): If $\delta(q, P) < \varepsilon$.

1365 Let $p \in P$ be q 's nearest-neighbor in P , thus $d(q, p) = d_{nn}(q, P)$. Since R
1366 is α -selective, there exists a point $r \in R$ such that $d(p, r) = d_{nn}(p, R) <$
1367 $d_{ne}(p, P)/(1 + \alpha)$. Additionally, by the triangle inequality and the definition of
1368 nearest-enemies, we have

$$d_{ne}(p, P) \leq d(p, ne(q, P)) \leq d(p, q) + d(q, ne(q, P)) = d_{nn}(q, P) + d_{ne}(q, P).$$

From the definition of chromatic density, $d_{ne}(q, P) = (1 + \delta(q, P)) d_{nn}(q, P)$. Together, these inequalities imply that $(1 + \alpha) d(p, r) \leq (2 + \delta(q, P)) d_{nn}(q, P)$. Therefore:

$$d_{nn}(q, R) \leq d(q, r) \leq d(q, p) + d(p, r) \leq \left(1 + \frac{2 + \delta(q, P)}{1 + \alpha}\right) d_{nn}(q, P).$$

1369 Now, assuming $\delta(q, P) < \varepsilon$ and setting $\alpha = 2/\varepsilon$, imply that $d_{nn}(q, R) <$
1370 $(1 + \varepsilon) d_{nn}(q, P)$. Therefore, the nearest-neighbor of q in R is an ε -approximate
1371 nearest-neighbor of q in P .

1372 Cases 1 and 2 imply that setting α to $2/\varepsilon$ is sufficient to ensure that the
 1373 nearest-neighbor rule classifies any query point $q \in \mathcal{X}$ with the class of one of its
 1374 valid ε -approximate nearest-neighbors in P . Therefore, R is an ε -coreset for the
 1375 nearest-neighbor rule on P . \square

1376 So far, we have assumed that nearest-neighbor queries over R are computed
 1377 exactly, as this is the standard notion of coresets. However, it is reasonable to
 1378 compute nearest-neighbors approximately even for R . How should the two approxi-
 1379 mations be combined to achieve a desired final degree of accuracy? Consider another
 1380 approximation parameter ξ , where $0 \leq \xi < \varepsilon$. We say that a set $R \subseteq P$ is an
 1381 (ξ, ε) -coreset for the approximate nearest-neighbor rule on P , if and only if for every
 1382 query point $q \in \mathcal{X}$, the class of any of q 's ξ -approximate nearest-neighbor in R
 1383 is the same as the class of any of its ε -approximate nearest-neighbors in P . The
 1384 following result generalizes Theorem 5.4 to accommodate for ξ -ANN queries after
 1385 condensation.

1386 **Theorem 5.5.** *Any α -selective subset is an (ξ, ε) -coreset for the approximate nearest-
 1387 neighbor rule when $\alpha = \Omega(1/(\varepsilon - \xi))$.*

1388 *Proof.* This follows from similar arguments to the ones described in the proof of
 1389 Theorem 5.4. Instead, here we set $\alpha = (\varepsilon\xi + 3\xi + 2)/(\varepsilon - \xi)$. Let R be an α -selective
 1390 subset of P , and $q \in \mathcal{X}$ any query point in the metric space, consider the same two
 1391 cases:

1392 Case 1 (Correct-Classification guarantee): If $\delta(q, P) \geq \varepsilon$.

1393 Consider the bound derived in Lemma 5.2. By our assumption that $\delta(q, P) \geq$
 1394 $\varepsilon > 0$, and since $\alpha \geq 0$, the following inequality holds true:

$$\delta(q, R) > \frac{\alpha \delta(q, P) - 2}{\delta(q, P) + \alpha + 3} \geq \frac{\alpha\varepsilon - 2}{\varepsilon + \alpha + 3}$$

1395 Based on this, it is easy to see that the assignment of $\alpha = (\varepsilon\xi + 3\xi + 2)/(\varepsilon - \xi)$
 1396 implies that $\delta(q, R) > \xi$, meaning that any of q 's ξ -approximate nearest-
 1397 neighbors in R belong to the same class as q 's nearest-neighbor in P . Intuitively,
 1398 this guarantees that q is correctly classified by the ξ -ANN rule in R .

1399 Case 2 (ε -Approximation guarantee): If $\delta(q, P) < \varepsilon$.

1400 The assignment of α implies that $d_{nn}(q, R) < \frac{1+\varepsilon}{1+\xi} d_{nn}(q, P)$. This means that an
 1401 ξ -ANN query for q in R , will return one of q 's ε -approximate nearest-neighbors
 1402 in P .

1403 All together, this implies that R is an (ξ, ε) -coreset for the nearest-neighbor rule on
 1404 P . \square

1405 In contrast with standard condensation criteria, these new results provide
 1406 guarantees on either approximation or the correct classification, of any query point in
 1407 the metric space. This is true even for query points that were “hard” to classify with

1408 the entire training set, formally defined as query points with low chromatic density.
 1409 Consequently, Theorems 5.4 and 5.5 show that α must be set to large values if we
 1410 hope to provide any sort of guarantees for these query points. However, better results
 1411 can be achieved by restricting the set of points that are guaranteed to be correctly
 1412 classified. This relates to the notion of *weak coresets*, which provide approximation
 1413 guarantees only for a subset of the possible queries. Given $\beta \geq 0$, we define \mathcal{Q}_β as
 1414 the set of query points in \mathcal{X} whose chromatic density with respect to P is at least
 1415 β (*i.e.*, $\mathcal{Q}_\beta = \{q \in \mathcal{X} \mid \delta(q, P) \geq \beta\}$). The following result describes the trade-off
 1416 between α and β to guarantee the correct classification of query points in \mathcal{Q}_β after
 1417 condensation.

1418 **Theorem 5.6.** *Any α -consistent subset is a weak ε -coreset for the nearest-neighbor
 1419 rule for queries in \mathcal{Q}_β , for $\beta = 2/\alpha$. Moreover, all query points in \mathcal{Q}_β are correctly
 1420 classified.*

1421 The proof of this theorem is rather simple, and follows the same arguments
 1422 outlined in Case 1 of the proof of Theorem 5.4. Basically, we use Lemma 5.2 to
 1423 show that for any query point $q \in \mathcal{Q}_\beta$, q 's chromatic density after condensation
 1424 is greater than zero if $\alpha\beta \geq 2$. Note that ε plays no role in this result, as the
 1425 guarantee on query points of \mathcal{Q}_β is of correct classification (*i.e.*, the class of its *exact*
 1426 nearest-neighbor in P), rather than an approximation.

1427 The trade-off between α and β is illustrated in Figure 5.2. From an initial
 1428 training set $P \subset \mathbb{R}^2$ (Figure 5.2a), we show the regions of \mathbb{R}^2 that comprise the sets \mathcal{Q}_β
 1429 for $\beta = 2/\alpha$, using $\alpha = \{0.1, 0.2, \sqrt{2}\}$ (Figures 5.2b-5.2d). While evidently, increasing
 1430 α guarantees that more query points will be correctly classified after condensation,
 1431 this example demonstrates a phenomenon commonly observed experimentally: most
 1432 query points lie far from enemy points, and thus have high chromatic density with
 1433 respect to P . Therefore, while Theorem 5.4 states that α must be set to $2/\varepsilon$ to
 1434 provide approximation guarantees on all query points, Theorem 5.6 shows that much
 1435 smaller values of α are sufficient to provide guarantees on some query points, as
 1436 evidenced in the example in Figure 5.2.




Figure 5.2: Depiction of the \mathcal{Q}_β sets for which any α -consistent subset is weak coresset ($\beta = 2/\alpha$). Query points in the *yellow* \bullet areas are inside \mathcal{Q}_β , and thus correctly classified after condensation. Query points in the *blue* \bullet areas are not in \mathcal{Q}_β , and have no guarantee of correct classification.

1437 These results establish a clear connection between the problem of condensation
 1438 and that of finding coresets for the nearest-neighbor rule, and provides a roadmap
 1439 to prove Theorem 5.1. This is the first characterization of sufficient conditions to
 1440 correctly classify any query point in \mathcal{X} after condensation, and not just the points
 1441 in P (as the original consistency criteria implies). In the following section, these
 1442 existential results are matched with algorithms to compute α -selective subsets of P
 1443 of bounded cardinality.

1444 5.3 Coreset Computation

1445 5.3.1 Hardness Results

1446 Define MIN- α -CS to be the problem of computing an α -consistent subset of
 1447 minimum cardinality for a given training set P . Similarly, let MIN- α -SS be the
 1448 corresponding optimization problem for α -selective subsets. Following known results
 1449 from standard condensation [10–12], when α is set to zero, the MIN-0-CS and
 1450 MIN-0-SS problems are both known to be NP-hard. Being special cases of the
 1451 general problems just defined, this implies that both MIN- α -CS and MIN- α -SS are
 1452 NP-hard.

1453 Here we present results related to the hardness of approximation of both
 1454 problems, along with simple algorithmic approaches with tight approximation factors.

1455 **Theorem 5.7.** *The MIN- α -CS problem is NP-hard to approximate in polynomial
 1456 time within a factor of $2^{(\text{ddim}(\mathcal{X}) \log((1+\alpha)/\gamma))^{1-o(1)}}$.*

1457 The full proof is omitted, as it follows from a modification of the hardness
 1458 bounds proof for the MIN-0-CS problem described in [36], which is based on a
 1459 reduction from the *Label Cover* problem. Proving Theorem 5.7 involves a careful
 1460 adjustment of the distances in this reduction, so that all the points in the construction
 1461 have chromatic density at least α . Consequently, this would imply that the minimum
 1462 nearest-enemy distance is reduced by a factor of $1/(1 + \alpha)$, explaining the resulting
 1463 bound for MIN- α -CS.

1464 The NET algorithm [36] can also be generalized to compute α -consistent subsets
 1465 of P as follows. We define α -NET as the algorithm that computes a $\gamma/(1 + \alpha)$ -net
 1466 of P , where γ is the smallest nearest-enemy distance in P . The covering property
 1467 of nets [67] implies that the resulting subset is α -consistent, while the packing
 1468 property suggests that its cardinality is $\mathcal{O}(((1 + \alpha)/\gamma)^{\text{ddim}(\mathcal{X})+1})$, implying a tight
 1469 approximation to the MIN- α -CS problem.

1470 **Theorem 5.8.** *The MIN- α -SS problem is NP-hard to approximate in polynomial
 1471 time within a factor of $(1 - o(1)) \ln n$ unless $\text{NP} \subseteq \text{DTIME}(n^{\log \log n})$.*

1472 *Proof.* The result follows from the hardness of another related covering problem: the
 1473 minimum *dominating set* [68–70]. We describe a simple L-reduction from any instance
 1474 of this problem to an instance of MIN- α -SS, which preserves the approximation
 1475 ratio.

1476 1. Consider any instance of minimum dominating set, consisting of the graph
1477 $G = (V, E)$.

1478 2. Generate a new edge-weighted graph G' as follows:

1479 Create two copies of G , namely $G_r = (V_r, E_r)$ and $G_b = (V_b, E_b)$, of *red* and *blue*
1480 nodes respectively. Set all edge-weights of G_r and G_b to be 1. Finally, connect
1481 each red node v_r to its corresponding blue node v_b by an edge $\{v_r, v_b\}$ of weight
1482 $1 + \alpha + \xi$ for a sufficiently small constant $\xi > 0$. Formally, G' is defined as the
1483 edge-weighted graph $G' = (V', E')$ where the set of nodes is $V' = V_r \cup V_b$, the
1484 set of edges is $E' = E_r \cup E_b \cup \{(v_r, v_b) \mid v \in V\}$, and an edge-weight function
1485 $w : E' \rightarrow \mathbb{R}^+$ where $w(e) = 1$ iff $e \in E_r \cup E_b$, and $w(e) = 1 + \alpha + \xi$ otherwise.

1486 3. A labeling function l where $l(v) = \text{red}$ iff $v \in V_r$, and $l(v) = \text{blue}$ iff $v \in V_b$.

1487 4. Compute the shortest-path metric of G' , denoted as $\mathbf{d}_{G'}$.

1488 5. Solve the MIN- α -SS problem for the set V' , on metric $\mathbf{d}_{G'}$, and the labels
1489 defined by l .

1490 A dominating set of G consists of a subset of nodes $D \subseteq V$, such that every
1491 node $v \in V \setminus D$ is adjacent to a node in D . Given any dominating set $D \subseteq V$ of
1492 G , it is easy to see that the subset $R = \{v_r, v_b \mid v \in D\}$ is an α -selective subset
1493 of V' , where $|R| = 2|D|$. Similarly, given an α -selective subset $R \subseteq V'$, there is
1494 a corresponding dominating set D of G , where $|D| \leq |R|/2$, as D can be either
1495 $R \cap V_r$ or $R \cap V_b$. Therefore, MIN- α -SS is as hard to approximate as the minimum
1496 dominating set problem. \square

1497 There is a clear connection between the MIN- α -SS problem and covering
1498 problems, in particular that of finding an optimal hitting set. Given a set of elements
1499 U and a family C of subsets of U , a *hitting set* of (U, C) is a subset $H \subseteq U$ such
1500 that every set in C contains at least one element of H . Therefore, let $N_{p,\alpha}$ be the
1501 set of points of P whose distance to p is less than $\mathbf{d}_{\text{he}}(p)/(1 + \alpha)$, then any hitting
1502 set of $(P, \{N_{p,\alpha} \mid p \in P\})$ is also an α -selective subset of P , and vice versa. This
1503 simple reduction implies a $\mathcal{O}(n^3)$ worst-case time $\mathcal{O}(\log n)$ -approximation algorithm
1504 for MIN- α -SS, based on the classic greedy algorithm for set cover [71, 72]. Call this
1505 approach α -HSS or α -*Hitting Selective Subset*. It follows from Theorem 5.8 that for
1506 training sets in general metric spaces, this is the best approximation possible under
1507 standard complexity assumptions.

1508 While both α -NET and α -HSS compute tight approximations of their corre-
1509 sponding problems, their performance in practice does not compare to heuristic
1510 approaches for standard condensation (see Section 5.4 for experimental results).
1511 Therefore, in the following subsections, we consider two practical algorithms for this
1512 problem, namely SFCNN and RSS, and extend them to compute subsets with the
1513 newly defined criteria.

1514 5.3.2 An Algorithm for α -Selective Subsets

1515 For standard condensation, we have already analyzed the RSS algorithm (see
1516 Chapter 4) used to compute selective subsets. It runs in quadratic worst-case time
1517 and exhibits good performance in practice. The selection process of this algorithm is
1518 heuristic in nature and can be described as follows: beginning with an empty set, the
1519 points in $p \in P$ are examined in increasing order with respect to their nearest-enemy
1520 distance $d_{ne}(p)$. The point p is added to the subset R if $d_{nn}(p, R) \geq d_{ne}(p)$. It is easy
1521 to see that the resulting subset is selective.

1522 Now, we define a generalization called α -RSS, to compute α -selective subsets
1523 of P . The condition to add a given point $p \in P$ to the selected subset checks if any
1524 previously selected point is closer to p than $d_{ne}(p)/(1 + \alpha)$, instead of just $d_{ne}(p)$.
1525 See Algorithm 10 for a formal description, and Figure 5.3 for an illustration. It is
1526 easy to see that this algorithm computes an α -selective subset, while keeping the
1527 quadratic time complexity of the original RSS algorithm.

Algorithm 10: α -RSS

Input: Initial training set P and parameter $\alpha \geq 0$
Output: α -selective subset $R \subseteq P$

- 1 $R \leftarrow \emptyset$
- 2 Let $\{p_i\}_{i=1}^n$ be the points of P sorted increasingly w.r.t. their nearest-enemy
distance $d_{ne}(p_i)$
- 3 **foreach** $p_i \in P$, where $i = 1 \dots n$ **do**
- 4 **if** $(1 + \alpha) \cdot d_{nn}(p_i, R) \geq d_{ne}(p_i)$ **then**
- 5 $R \leftarrow R \cup \{p_i\}$
- 6 **return** R

1528 Naturally, we want to analyze the number of points this algorithm selects. The
1529 remainder of this section establishes upper-bounds and approximation guarantees of
1530 the α -RSS algorithm for any doubling metric space, with improved results in the
1531 Euclidean space. This proves the result mentioned in Chapter 4 that RSS computes
1532 an approximation of the MIN-0-CS and MIN-0-SS problems.

1533 5.3.2.1 Size in Doubling spaces

1534 First, we consider the case where the underlying metric space (\mathcal{X}, d) of P
1535 is doubling. The following results depend on the doubling dimension $ddim(\mathcal{X})$ of
1536 the metric space (which is assumed to be constant), the margin γ (the smallest
1537 nearest-enemy distance of any point in P), and κ (the number of nearest-enemy
1538 points in P).

1539 **Theorem 5.9.** α -RSS computes a tight approximation for the MIN- α -CS problem.

1540 *Proof.* This follows from a direct comparison to the resulting subset of the α -NET
1541 algorithm from the previous section. For any point p selected by α -NET, let $B_{p,\alpha}$




Figure 5.3: Selection of α -RSS for $\alpha=0.5$. Faded points are not selected, while selected points are drawn along with a ball of radius $d_{ne}(p)$ (dotted outline) and a ball of radius $d_{ne}(p)/(1 + \alpha)$ (solid outline). A point p is selected if no previously selected point is closer to p than $d_{ne}(p)/(1 + \alpha)$.

be the set of points of P “covered” by p , that is, whose distance to p is at most $\gamma/(1 + \alpha)$. By the covering property of ε -nets, this defines a partition on P when considering every point p selected by α -NET.

Let R be the set of points selected by α -RSS, we analyze the size of $B_{p,\alpha} \cap R$, that is, for any given $B_{p,\alpha}$ how many points could have been selected by the α -RSS algorithm. Let $a, b \in B_{p,\alpha} \cap R$ be any two such points, where without loss of generality, $d_{ne}(a) \leq d_{ne}(b)$. By the selection process of the algorithm, we know that $d(a, b) \geq d_{ne}(b)/(1 + \alpha) \geq \gamma/(1 + \alpha)$. A simple packing argument in doubling metrics implies that $|B_{p,\alpha} \cap R| \leq 2^{\text{ddim}(\mathcal{X})+1}$. Altogether, we have that the size of the subset selected by α -RSS is $\mathcal{O}((2(1 + \alpha)/\gamma)^{\text{ddim}(\mathcal{X})+1})$. \square

Theorem 5.10. α -RSS computes an $\mathcal{O}(\log(\min(1 + 2/\alpha, 1/\gamma)))$ -factor approximation for the MIN- α -SS problem. For $\alpha = \Omega(1)$, this is a constant-factor approximation.

Proof. Let OPT_α be the optimum solution to the MIN- α -SS problem, i.e., the minimum cardinality α -selective subset of P . For every point $p \in \text{OPT}_\alpha$ in such solution, define $S_{p,\alpha}$ to be the set of points in P “covered” by p , or simply $S_{p,\alpha} = \{r \in P \mid d(r, p) < d_{ne}(r)/(1 + \alpha)\}$. Additionally, let R be the set of points selected by α -RSS, define $R_{p,\sigma}$ to be the points selected by α -RSS which also belong to $S_{p,\alpha}$ and whose nearest-enemy distance is between σ and 2σ , for $\sigma \in [\gamma, 1]$. That is,

1561 $R_{p,\sigma} = \{r \in R \cap S_{p,\alpha} \mid d_{ne}(r) \in [\sigma, 2\sigma]\}$. Clearly, these subsets define a partitioning
1562 of R for all $p \in \text{OPT}_\alpha$ and values of $\sigma = \gamma 2^i$ for $i = \{0, 1, 2, \dots, \lceil \log \frac{1}{\gamma} \rceil\}$.

1563 However, depending on α , some values of σ would yield empty $R_{p,\sigma}$ sets.
1564 Consider some point $q \in S_{p,\alpha}$, we can bound its nearest-enemy distance with respect
1565 to the nearest-enemy distance of point p . In particular, by leveraging simple triangle-
1566 inequality arguments, it is possible to prove that $\frac{1+\alpha}{2+\alpha} d_{ne}(p) \leq d_{ne}(q) \leq \frac{1+\alpha}{\alpha} d_{ne}(p)$.
1567 Therefore, the values of σ for which $R_{p,\sigma}$ sets are not empty, are $\sigma = 2^j \frac{1+\alpha}{2+\alpha} d_{ne}(p)$
1568 for $j = \{0, \dots, \lceil \log(1 + 2/\alpha) \rceil\}$.

1569 The proof now follows by bounding the size of $R_{p,\sigma}$ which can be achieved by
1570 bounding its spread. Thus, let's consider the smallest and largest pairwise distances
1571 among points in $R_{p,\sigma}$. Take any two points $a, b \in R_{p,\sigma}$ where without loss of
1572 generality, $d_{ne}(a) \leq d_{ne}(b)$. Note that points selected by α -RSS cannot be “too close”
1573 to each other; that is, as a and b were selected by the algorithm, we know that
1574 $(1 + \alpha) d(a, b) \geq d_{ne}(b) \geq \sigma$. Therefore, the smallest pairwise distance in $R_{p,\sigma}$ is at
1575 least $\sigma/(1 + \alpha)$. Additionally, by the triangle inequality, we can bound the maximum
1576 pairwise distance using their distance to p as $d(a, b) \leq d(a, p) + d(p, b) \leq 4\sigma/(1 + \alpha)$.
1577 Then, by the packing properties of doubling spaces, the size of $R_{p,\sigma}$ is at most
1578 $4^{\text{ddim}(\mathcal{X})+1}$.

1579 Altogether, for every $p \in \text{OPT}_\alpha$ there are up to $\lceil \log(\min(1 + 2/\alpha, 1/\gamma)) \rceil$ non-
1580 empty $R_{p,\sigma}$ subsets, each containing at most $4^{\text{ddim}(\mathcal{X})+1}$ points. In doubling spaces
1581 with constant doubling dimension, the size of these subsets is also constant. \square

1582 While these results are meaningful from a theoretical perspective, it is also
1583 useful to establishing bounds in terms of the geometry of the learning space, which
1584 is characterized by the boundaries between points of different classes. Thus, using
1585 similar packing arguments as above, we bound the selection size of the algorithm
1586 with respect to κ .

1587 **Theorem 5.11.** α -RSS selects $\mathcal{O}\left(\kappa \log \frac{1}{\gamma} (1 + \alpha)^{\text{ddim}(\mathcal{X})+1}\right)$ points.

1588 *Proof.* This follows from similar arguments to the ones used to prove Theorem 5.10,
1589 using an alternative charging scheme for each nearest-enemy point in the training set.
1590 Consider one such point $p \in \{\text{ne}(r) \mid r \in P\}$ and a value $\sigma \in [\gamma, 1]$, we define $R'_{p,\sigma}$ to
1591 be the subset of points from α -RSS whose nearest-enemy is p , and their nearest-enemy
1592 distance is between σ and 2σ . That is, $R'_{p,\sigma} = \{r \in R \mid \text{ne}(r) = p \wedge d_{ne}(r) \in [\sigma, 2\sigma]\}$.
1593 These subsets partition R for all nearest-enemy points of P , and values of $\sigma = \gamma 2^i$
1594 for $i = \{0, 1, 2, \dots, \lceil \log \frac{1}{\gamma} \rceil\}$.

1595 For any two points $a, b \in R'_{p,\sigma}$, the selection criteria of α -RSS implies some
1596 separation between selected points, which can be used to prove that $d(a, b) \geq \sigma/(1 + \alpha)$.
1597 Additionally, we know that $d(a, b) \leq d(a, p) + d(p, b) = d_{ne}(a) + d_{ne}(b) \leq 4\sigma$. Using
1598 a simple packing argument, we have that $|R'_{p,\sigma}| \leq [4(1 + \alpha)]^{\text{ddim}(\mathcal{X})+1}$.

1599 Altogether, by counting all sets $R'_{p,\sigma}$ for each nearest-enemy in the training set
1600 and values of σ , the size of R is upper-bounded by $|R| \leq \kappa \lceil \log 1/\gamma \rceil [4(1 + \alpha)]^{\text{ddim}(\mathcal{X})+1}$.
1601 Based on the assumption that $\text{ddim}(\mathcal{X})$ is constant, this completes the proof. \square

1602 As a corollary, this result implies that when $\alpha = 2/\varepsilon$, the α -selective subset
 1603 computed by α -RSS contains $\mathcal{O}(\kappa \log 1/\gamma (1/\varepsilon)^{\text{ddim}(\mathcal{X})+1})$ points. This establishes
 1604 the size bound on the ε -coreset given in Theorem 5.1, which can be computed using
 1605 the α -RSS algorithm.

1606 5.3.2.2 Size in Euclidean space

1607 In the case where $P \subset \mathbb{R}^d$ lies in d -dimensional Euclidean space, the analysis
 1608 of α -RSS can be further improved, leading to a constant-factor approximation of
 1609 MIN- α -SS for values of $\alpha \geq 0$, and reduced dependency on the dimensionality of P .

1610 **Theorem 5.12.** α -RSS computes an $\mathcal{O}(1)$ -approximation for the MIN- α -SS problem
 1611 in \mathbb{R}^d .

1612 *Proof.* Similar to the proof of Theorem 5.10, define $R_p = S_{p,\alpha} \cap R$ as the points selected
 1613 by α -RSS that are “covered” by p in the optimum solution OPT_α . Consider two such
 1614 points $a, b \in R_p$ where without loss of generality, $d_{\text{ne}}(a) \leq d_{\text{ne}}(b)$. By the definition of
 1615 $S_{p,\alpha}$ we know that $d(a, p) < d_{\text{ne}}(a)/(1 + \alpha)$, and similarly with b . Additionally, from
 1616 the selection of the algorithm we know that $d(a, b) \geq d_{\text{ne}}(b)/(1 + \alpha)$. Overall, these
 1617 inequalities imply that the angle $\angle apb \geq \pi/3$. By a simple packing argument, the
 1618 size of R_p is bounded by the *kissing number* in d -dimensional Euclidean space, or
 1619 simply $\mathcal{O}((3/\pi)^{d-1})$. Therefore, we have that $|R| \leq \sum_p |R_p| = |\text{OPT}_\alpha| \mathcal{O}((3/\pi)^{d-1})$.
 1620 Assuming d is constant, this completes the proof. \square




Figure 5.4: Training set where the analysis of the approximation factor of α -RSS in \mathbb{R}^d is tight.

1621 This analysis is tight up to constant factors. In Figure 5.4, we illustrate
 1622 a training set P consisting of *red* and *blue* points in \mathbb{R}^d , where α -RSS selects
 1623 $\Theta(c^{d-2} |\text{OPT}_\alpha|)$ points. Consider two helper points (which do not belong to P)
 1624 $c_r = 0\vec{u}_d$ and $c_b = (1 + \alpha)\vec{u}_d$, where \vec{u}_d is the unit vector parallel to the d -th

coordinate. Add red points r_i on the surface of the $d - 1$ unit ball centered at c_r and perpendicular to \vec{u}_d . Similarly with blue points b_i around c_b . Finally, add two points $r_* = -\xi \vec{u}_d$ and $b_* = (1 + \alpha + \xi) \vec{u}_d$, for a suitable value ξ such that $\|r_* r_i\| < 1$. Clearly, the nearest-enemy distance of all r_i and b_i points is $1 + \alpha$, while the one of r_* and b_* is strictly greater than $1 + \alpha$. Thus, $\text{OPT}_\alpha = \{r_*, b_*\}$ but $\alpha\text{-RSS}$ selects $\Theta(c^{d-2})$ points r_i and b_i at distance greater than 1 from each other.

Furthermore, a similar constant-factor approximation can be achieved for any training set P in ℓ_p space for $p \geq 3$. This follows analogously to the proof of Theorem 5.12, exploiting the bounds between ℓ_p and ℓ_2 metrics, where $1/\sqrt{d} \|v\|_p \leq \|v\|_2 \leq \sqrt{d} \|v\|_p$. This would imply that the angle between any two points in $\alpha\text{-RSS}_p$ is $\Omega(1/d)$. Therefore, it shows that $\alpha\text{-RSS}$ achieves an approximation factor of $\mathcal{O}(d^{d-1})$, or simply $\mathcal{O}(1)$ for constant dimension.

Similarly to the case of doubling spaces, we also establish upper-bounds in terms of κ for the selection size of the algorithm in Euclidean space. The following result improves the exponential dependence on the dimensionality of P (from $\text{ddim}(\mathbb{R}^d) = \Theta(d)$ to $d - 1$), while keeping the dependency on the margin γ , which contrast with the approximation factor results.

Theorem 5.13. *In Euclidean space \mathbb{R}^d , $\alpha\text{-RSS}$ selects $\mathcal{O}\left(\kappa \log \frac{1}{\gamma} (1 + \alpha)^{d-1}\right)$ points.*

Proof. Let p be any nearest-enemy point of P and $\sigma \in [\gamma, 1]$, similarly define $R'_{p,\sigma}$ to be the set of points selected by $\alpha\text{-RSS}$ whose nearest-enemy is p and their nearest-enemy distance is between σ and $b\sigma$, for $b = \frac{(1+\alpha)^2}{\alpha(2+\alpha)}$. Equivalently, these subsets define a partitioning of R for all nearest-enemy points p and values of $\sigma = \gamma b^k$ for $k = \{0, 1, 2, \dots, \lceil \log_b \frac{1}{\gamma} \rceil\}$. Thus, the proof follows from bounding the minimum angle between points in these subsets. For any two such points $p_i, p_j \in R'_{p,\sigma}$, we lower bound the angle $\angle p_i p p_j$. Assume without loss of generality that $\mathbf{d}_{\text{ne}}(p_i) \leq \mathbf{d}_{\text{ne}}(p_j)$. By definition of the partitioning, we also know that $\mathbf{d}_{\text{ne}}(p_j) \leq b\sigma \leq b\mathbf{d}_{\text{ne}}(p_i)$. Therefore, altogether we have that $\mathbf{d}_{\text{ne}}(p_i) \leq \mathbf{d}_{\text{ne}}(p_j) \leq b\mathbf{d}_{\text{ne}}(p_i)$.

First, consider the set of points whose distance to p_i is $(1 + \alpha)$ times their distance to p , which defines a multiplicative weighted bisector [73] between points p and p_i , with weights equal to 1 and $1/(1 + \alpha)$ respectively. This is characterized as a d -dimensional ball (see Figure 5.5a) with center $c_i = (p_i - p)b + p$ and radius $\mathbf{d}_{\text{ne}}(p_i)b/(1 + \alpha)$. Thus p , p_i and c_i are collinear, and the distance between p and c_i is $\mathbf{d}(p, c_i) = b\mathbf{d}_{\text{ne}}(p_i)$. In particular, let's consider the relation between p_j and such bisector. As p_j was selected by the algorithm after p_i , we know that $(1 + \alpha)\mathbf{d}(p_j, p_i) \geq \mathbf{d}_{\text{ne}}(p_j)$ where $\mathbf{d}_{\text{ne}}(p_j) = \mathbf{d}(p_j, p)$. Therefore, clearly p_j lies either outside or in the surface of the weighted bisector between p and p_i (see Figure 5.5b).

For angle $\angle p_i p p_j$, we can frame the analysis to the plane defined by p , p_i and p_j . Let x and y be two points in this plane, such that they are the intersection points between the weighted bisector and the balls centered at p of radii $\mathbf{d}_{\text{ne}}(p_i)$ and $b\mathbf{d}_{\text{ne}}(p_i)$ respectively (see Figure 5.5c). By the convexity of the weighted bisector between p and p_i , we can say that $\angle p_i p p_j \geq \min(\angle x p p_i, \angle y p c_j)$. Now, consider the triangles $\triangle p x p_i$ and $\triangle p y c_j$. By the careful selection of b , these triangles are both isosceles and similar. In particular, for $\triangle p x p_i$ the two sides incident to p have




Figure 5.5: Construction for the analysis of the minimum angle between two points in $R'_{p,\sigma}$ w.r.t. some nearest-enemy point $p \in P$. Let points $p_i, p_j \in R'_{p,\sigma}$, we analyze the angle $\angle p_i p p_j$.

length equal to $d_{\text{ne}}(p_i)$, and the side opposite to p has length equal to $d_{\text{ne}}(p_i)/(1 + \alpha)$.
 For $\triangle pyc_i$, the side lengths are $b d_{\text{ne}}(p_i)$ and $d_{\text{ne}}(p_i) b/(1 + \alpha)$. Therefore, the angle $\angle p_i p p_j \geq \angle xpp_i \geq 1/(1 + \alpha)$.

By a simple packing argument based on this minimum angle, we have that the size of $R'_{p,\sigma}$ is $\mathcal{O}((1 + \alpha)^{d-1})$. All together, following the defined partitioning, we have that:

$$|R| = \sum_p \sum_{k=0}^{\lceil \log_b \frac{1}{\gamma} \rceil} |R'_{p,b^k}| \leq \kappa \left[\log_b \frac{1}{\gamma} \right] \mathcal{O}((1 + \alpha)^{d-1})$$

For constant α and d , the size of α -RSS is $\mathcal{O}(\kappa \log \frac{1}{\gamma})$. Moreover, when α is zero α -RSS selects $\mathcal{O}(\kappa c^{d-1})$, matching the bound presented in Chapter 4 for RSS in Euclidean space. \square

5.3.2.3 Subquadratic Implementation

Here we present a subquadratic implementation for the α -RSS algorithm, which completes the proof of our main result, Theorem 5.1. Prior to this result, among algorithms for nearest-neighbor condensation, FCNN and SFCNN achieve the best worst-case time complexity, running in $\mathcal{O}(nm)$ time, where $m = |R|$ is the size of the selected subset.

The α -RSS algorithm consists of two main stages: computing the nearest-enemy distances of all points in P (and sorting the points based on these), and the selection process itself. The first stage requires a total of n nearest-enemy queries, plus additional $\mathcal{O}(n \log n)$ time for sorting. The second stage performs n nearest-neighbor queries on the current selected subset R , which needs to be updated m times. In both cases, using exact nearest-neighbor search would degenerate into linear search due to the *curse of dimensionality*. Thus, the first and second stage of the algorithm would need $\mathcal{O}(n^2)$ and $\mathcal{O}(nm)$ worst-case time respectively.

1691 These bottlenecks can be overcome by leveraging approximate nearest-neighbor
 1692 techniques. Clearly, the first stage of the algorithm can be improved by computing
 1693 nearest-enemy distances approximately, using as many ANN structures as classes
 1694 there are in P , which is considered to be a small constant. Therefore, by also applying
 1695 a simple brute-force search for nearest-neighbors in the second stage, result (i) of
 1696 the next theorem follows immediately. Moreover, by combining this with standard
 1697 techniques for static-to-dynamic conversions [74], we have result (ii) below. Denote
 1698 this variant of α -RSS as (α, ξ) -RSS, for a parameter $\xi \geq 0$.

1699 **Theorem 5.14.** *Given a data structure for ξ -ANN searching with construction time
 1700 t_c and query time t_q (which potentially depend on n and ξ), the (α, ξ) -RSS variant
 1701 can be implemented with the following worst-case time complexities, where m is the
 1702 size of the selected subset.*

1703 (i) $\mathcal{O}(t_c + n(t_q + m + \log n))$

1704 (ii) $\mathcal{O}((t_c + n t_q) \log n)$

1705 More generally, if we are given an additional data structure for dynamic ξ -
 1706 ANN searching with construction time t'_c , query time t'_q , and insertion time t'_i , the
 1707 overall running time will be $\mathcal{O}(t_c + t'_c + n(t_q + t'_q + \log n) + m t'_i)$. Indeed, this can
 1708 be used to obtain (ii) from the static-to-dynamic conversions [74], which propose an
 1709 approach to convert static search structures into dynamic ones. These results directly
 1710 imply implementations of (α, ξ) -RSS with subquadratic worst-case time complexities,
 1711 based on ANN techniques [43, 75] for low-dimensional Euclidean space, and using
 1712 techniques like LSH [50] that are suitable for ANN in high-dimensional Hamming
 1713 and Euclidean spaces. More generally, subquadratic runtimes can be achieved by
 1714 leveraging techniques [76] for dynamic ANN search in doubling spaces.

Algorithm 11: (α, ε) -RSS

Input: Initial training set P and parameters $\alpha, \varepsilon \geq 0$

Output: α -selective subset $R \subseteq P$

```

1  $R \leftarrow \emptyset$ 
2 Let  $\{p_i\}_{i=1}^n$  be the points of  $P$  sorted increasingly w.r.t. their approximate
   nearest-enemy distance  $d_{ne}(p_i, \varepsilon)$ 
3 foreach  $p_i \in P$ , where  $i = 1 \dots n$  do
4   if  $(1 + \alpha)(1 + \varepsilon) \cdot d_{nn}(p_i, R, \varepsilon) \geq d_{ne}(p_i, \varepsilon)$  then
5      $R \leftarrow R \cup \{p_i\}$ 
6 return  $R$ 

```

1715 It remains unclear how these new implementations of the algorithm would
 1716 deal with ‘‘uncertainty’’. That is, such implementation schemes for α -RSS would
 1717 incur an approximation error (of up to $1 + \xi$) on the computed distances: either
 1718 only during the first stage if (i) is implemented, or during both stages if (ii) or the
 1719 dynamic-structure scheme are implemented. The uncertainty introduced by these

1720 approximate queries, imply that in order to guarantee finding α -selective subsets, we
 1721 must modify the condition for adding point during the second stage of the algorithm.
 1722 Let $d_{ne}(p, \xi)$ denote the ξ -approximate nearest-enemy distance of p computed in the
 1723 first stage, and let $d_{nn}(p, R, \xi)$ denote the ξ -approximate nearest-neighbor distance of
 1724 p over points of the current subset (computed in the second stage). Then, (α, ξ) -RSS
 1725 adds a point p into the subset if $(1 + \xi)(1 + \alpha)d_{nn}(p, R, \xi) \geq d_{ne}(p, \xi)$.

1726 By similar arguments to the ones described in Section 5.3.2, size guarantees
 1727 can be extended to (α, ξ) -RSS. First, the size of the subset selected by (α, ξ) -RSS,
 1728 in terms of the number of nearest-enemy points in the set, would be bounded by the
 1729 size of the subset selected by $\hat{\alpha}$ -RSS with $\hat{\alpha} = (1 + \alpha)(1 + \xi)^2 - 1$. Additionally, the
 1730 approximation factor of (α, ξ) -RSS in both doubling and Euclidean metric spaces
 1731 would increase by a factor of $\mathcal{O}((1 + \xi)^{2(\text{ddim}(\mathcal{X})+1)})$.

1732 This completes the proof of Theorem 5.1.

1733 **Lemma 5.15.** *There exist a data structure for dynamic ξ -ANN queries in sets P
 1734 in d -dimensional Euclidean space, that can be constructed in $t'_c = \mathcal{O}(n \log n)$ time,
 1735 queried in $t'_q = \mathcal{O}(\log n + 1/\xi^{d-1})$ time, and where points of P can be inserted in
 1736 $t'_i = \mathcal{O}(\log n)$ time.*

1737 Together with Theorem 5.14 described above, this lemma implies that there is
 1738 a variant of α -RSS for Euclidean space that runs in $\mathcal{O}(n \log n + n/\xi^{d-1})$ time. Such
 1739 data structure can be build from a standard BBD tree [42, 77] as follows. First,
 1740 construct the tree from the entire set P , thus taking $t'_c = \mathcal{O}(n \log n)$ time. However,
 1741 each node of the tree has some additional data: a boolean flag indicating if the
 1742 subtree rooted at such node contains a point of the “active” subset R . Initially,
 1743 all flags are set to *false*, making the initial active subset being empty. To add a
 1744 point $p \in P$ to the active subset R , all the flags from the root of the tree to the leaf
 1745 node containing p must be set to *true*, thus making the insertion time $t'_i = \mathcal{O}(\log n)$.
 1746 Finally, an ξ -ANN query on such tree would perform as usual, only avoiding to visit
 1747 nodes whose flag is set to *false*, yielding a query time of $t'_q = \mathcal{O}(\log n + 1/\xi^{d-1})$.

1748 5.3.3 An Algorithm for α -Consistent Subsets

1749 Even thought the main result of this chapter relies on the computation of α -
 1750 selective subsets, Theorem 5.6 shows that even α -consistency is enough to guarantee
 1751 the correct classification of certain query points. In practice, FCNN [16] has been
 1752 acknowledged as the most efficient algorithm for computing consistent subsets. From
 1753 the results presented in Chapter 4, we know that while FCNN cannot be upper-
 1754 bounded in terms of k or κ , the simple modification of this algorithm called SFCNN
 1755 can be successfully upper-bounded. Therefore, in this section, we discuss a simple
 1756 extension of this algorithm in order to compute α -consistent subsets.

1757 Recall the workings of both the FCNN and SFCNN algorithms, which select
 1758 points iteratively as follows. First, the subset R is initialized with one point per class
 1759 (*e.g.*, the centroids of each class). During every iteration, the algorithm identifies all
 1760 the points in P that are incorrectly classified with the current R , or simply, those

1761 whose nearest-neighbor in R is of different class. This is formalized as the *voren*
1762 function, defined for every point $p \in R$ as follows:

$$\text{voren}(p, R, P) = \{q \in P \mid \text{nn}(q, R) = p \wedge l(q) \neq l(p)\}$$

1763 This function identifies all the enemies of p whose nearest-neighbor in R is p
1764 itself. The only difference between the original FCNN algorithm and the modified
1765 SFCNN appears next. While FCNN adds one point per each $p \in R$ in a batch²,
1766 potentially doubling the size of R , SFCNN adds only *one* point per iteration. Then,
1767 both algorithms terminate when no other points can be added (*i.e.*, all $\text{voren}(p, R, P)$
1768 are empty), implying that R is consistent.

1769 We can now extend both algorithms to compute α -consistent subsets, namely
1770 α -FCNN and α -SFCNN, by redefining the *voren* function. The idea is simple: to
1771 identify those points whose nearest-neighbor in R is p , such that are either enemies
1772 of p , or whose chromatic density with respect to R is less than α . This is formally
1773 defined as follows:

$$\text{voren}(\alpha, p, R, P) = \{q \in P \mid \text{nn}(q, R) = p \wedge (l(q) = l(p) \Rightarrow \delta(q, R) < \alpha)\}$$

1774 By plugging this function into the algorithms (see Algorithm 12), it is easy
1775 to show that the resulting subsets are α -consistent. Moreover, this can be easily
1776 implemented to run in $\mathcal{O}(nm)$ worst-case time, where m is the final size of R ,
1777 extending the implementation scheme described in the paper where FCNN was
1778 initially proposed [16].

Algorithm 12: α -SFCNN

Input: Initial training set P and parameter $\alpha \geq 0$
Output: α -consistent subset $R \subseteq P$

```

1  $R \leftarrow \emptyset$ 
2  $S \leftarrow \text{centroids}(P)$ 
3 while  $S \neq \emptyset$  do
4    $R \leftarrow R \cup \{\text{Choose one point of } S\}$ 
5    $S \leftarrow \emptyset$ 
6   foreach  $p \in R$  do
7      $S \leftarrow S \cup \{\text{rep}(p, \text{voren}(\alpha, p, R, P))\}$ 
8 return  $R$ 
```

1779 Finally, leveraging the analysis described in Section 4.3.3, together with the
1780 proofs of Theorems 5.9 and 5.11, we upper-bound the selection size of the α -SFCNN
1781 algorithm. The proofs of the next results depend on the following observation. Let
1782 $a, b \in R$ be two points selected by α -SFCNN, where $d_{\text{ne}}(a), d_{\text{ne}}(b) \geq \beta$ for some
1783 $\beta \geq 0$, it is easy to show that $d(a, b) \geq \beta/(1 + \alpha)$. This follows from a fairly simple

2For FCNN, line 4 of Algorithm 12 updates R by adding all the points in S , instead of only one point of S .

1784 argument: to the contrary, suppose that $d(a, b) < \beta/(1 + \alpha)$, which would imply that
 1785 a and b belong to the same class. Without loss of generality, point a was added to R
 1786 before point b . Note that after adding point a to R , the chromatic density of b w.r.t.
 1787 R is $\delta(b, R) > \alpha$, which contradicts the statement that b could be added to R .

1788 **Theorem 5.16.** α -SFCNN computes a tight approximation for the MIN- α -CS
 1789 problem.

1790 This result follows by similar arguments as the proof of Theorem 5.9. By
 1791 considering any two points $a, b \in B_{p,\alpha} \cap R$, we know that $d(a, b) \geq \gamma/(1 + \alpha)$ as γ is
 1792 the smallest nearest-enemy distance in P . This implies α -SFCNN can select up to
 1793 $2^{\text{ddim}(\mathcal{X})+1}$ times more points as the α -NET algorithm, which yields the proof.

1794 **Theorem 5.17.** α -SFCNN selects $\mathcal{O}\left(\kappa \log \frac{1}{\gamma} (1 + \alpha)^{\text{ddim}(\mathcal{X})+1}\right)$ points.

1795 Similarly, this result can be proven using the same arguments outlined to prove
 1796 Theorem 5.11. After partitioning the selection of α -SFCNN into $\mathcal{O}(\kappa \log 1/\gamma)$ subsets,
 1797 consider any two points a, b in one of these subsets, where $d_{\text{ne}}(a), d_{\text{ne}}(b) \in [\sigma, 2\sigma]$, for
 1798 some $\sigma \in [\gamma, 1]$. Therefore, we can show that $d(a, b) \geq \sigma/(1 + \alpha)$, which implies that
 1799 each subset in the partitioning contains at most $\lceil 4(1 + \alpha) \rceil^{\text{ddim}(\mathcal{X})+1}$ points. This
 1800 yields the proof.

1801 5.4 Experimental Comparison

1802 In order to get a clearer impression of the relevance of these results in practice,
 1803 we performed experimental trials on several training sets, both synthetically generated
 1804 and widely used benchmarks. First, we consider 21 training sets from the UCI
 1805 *Machine Learning Repository*³ which are commonly used in the literature to evaluate
 1806 condensation algorithms [18]. These consist of a number of points ranging from 150
 1807 to 58000, in d -dimensional Euclidean space with d between 2 and 64, and 2 to 26
 1808 classes. We also generated some synthetic training sets, containing 10^5 uniformly
 1809 distributed points, in 2 to 3 dimensions, and 3 classes. All training sets used in
 1810 these experimental trials are summarized in Table 4.2. The implementation of the
 1811 algorithms, training sets used, and raw results, are publicly available⁴.

1812 These experimental trials compare the performance of different condensation
 1813 algorithms when applied to vastly different training sets. We use two measures of
 1814 comparison on these algorithms: their runtime in the different training sets, and
 1815 the size of the subset selected. Clearly, these values might differ greatly on training
 1816 sets whose size are too distinct. Therefore, before comparing the raw results, these
 1817 are normalized. The runtime of an algorithm for a given training set is normalized
 1818 by dividing it by n , the size of the training set. The size of the selected subset is
 1819 normalized by dividing it by κ , the number of nearest-enemy points in the training
 1820 set, which characterizes the complexity of the boundaries between classes.

³<https://archive.ics.uci.edu/ml/index.php>

⁴<https://github.com/afloresv/nnc/>

1821 5.4.1 Algorithm Comparison

1822 The first experiment evaluates the performance of the five algorithms discussed
 1823 in this chapter: α -RSS, α -FCNN, α -SFCNN, α -HSS, and α -NET. The evaluation is
 1824 carried out by varying the value of the α parameter from 0 to 1, to understand the
 1825 impact of increasing this parameter. The implementation of α -HSS uses the well-
 1826 known greedy algorithm for set cover [71], and solves the problem using the reduction
 1827 described in Section 5.3.1. In the other hand, recall that the original NET algorithm
 1828 (for $\alpha = 0$) implements an extra pruning technique to further reduce the training set
 1829 after computing the γ -net [36]. For a fair comparison, we implemented the α -NET
 1830 algorithm with a modified version of this pruning technique that guarantees that the
 1831 selected subset is still α -selective.

1832 The results show that α -RSS outperforms the other algorithms in terms of
 1833 running time by a big margin, and irrespective of the value of α (see Figure 5.6a).
 1834 Additionally, the number of points selected by α -RSS, α -FCNN, and α -SFCNN is
 1835 comparable to α -HSS, which guarantees the best possible approximation factor in
 1836 general metrics, while α -NET is significantly outperformed.




Figure 5.6: Comparison α -RSS, α -FCNN, α -SFCNN, α -NET, and α -HSS, for different values of α .

1837 5.4.2 Subquadratic Approach

1838 Using the same experimental framework, we evaluate performance of the
 1839 subquadratic implementation (α, ξ) -RSS described in Section 5.3.2.3. In this case,
 1840 we change the value of parameter ξ to assess its effect on the running time and
 1841 selection size over the algorithm, for two different values of α (see Figure 5.7). The
 1842 results show an expected increase of the number of selected points, while significantly
 1843 improving its running time.




Figure 5.7: Evaluating the effect of increasing the parameter ξ on (α, ξ) -RSS for $\alpha = \{0, 0.2\}$.

1844 Chapter 6: Chromatic Approximate Voronoi Diagram

1845 6.1 Introduction

1846 As evidenced from the previous chapters, a common approach towards dealing
1847 with nearest-neighbor classification, is to reduce this problem to the one of nearest-
1848 neighbor search. That is, taking the initial training set P and using any out-of-the-box
1849 solution for computing nearest-neighbors of P , like AVDs [19, 20], LSH [21], and
1850 HNSW graphs [22]. Assuming this reduction for nearest-neighbor classification, there
1851 are only two ways of having more efficient queries. The first, by preprocessing the
1852 training set via some condensation algorithm or by computing an ε -coresets, as
1853 described in Chapters 3 to 5. The second option is to improve the complexity of
1854 nearest-neighbor search techniques (usually involving approximate solutions), which
1855 is a known and extensive line of research.

1856 In this chapter we explore an alternative approach towards efficient nearest-
1857 neighbor classification. The idea then is to avoid reducing this problem to the one of
1858 nearest-neighbor search, but instead proposing an approach that directly computes
1859 the predicted class. Therefore, our approach would bypass the preprocessing of the
1860 training set, and build a tailor-made data structure for approximate nearest-neighbor
1861 classification. Formally, we are given training set P in a metric space $(\mathcal{X}, \mathbf{d})$ and an
1862 approximation parameter $0 < \varepsilon \leq \frac{1}{2}$, and the goal is to construct a data structure so
1863 that given any query point $q \in \mathcal{X}$, it is possible to efficiently classify q according to
1864 any valid ε -approximate nearest-neighbor of q in P . Throughout, we take domain \mathcal{X}
1865 to be the d -dimensional Euclidean space \mathbb{R}^d , distance function \mathbf{d} to be the L_2 norm,
1866 and we assume that the dimension d is a fixed constant, independent of n and ε .

1867 These results have been published in [78].

1868 Related Work

1869 When working with ε -approximate nearest-neighbor searching (ε -ANN), the
1870 objective is to compute a point whose distance from the query point is within a
1871 factor of $1 + \varepsilon$ of the true nearest-neighbor. This problem, referred to as “standard
1872 ANN” throughout this chapter, has been widely studied. In *chromatic ε -ANN search*
1873 the objective is to return just the class (or more visually, the “color”) of any such
1874 point [51]. We refer to this as ε -classification, and it is the focus of this chapter.

1875 Clearly, chromatic ANN queries can be reduced to standard ANN queries.
1876 Hence, most of the efficiency improvements in nearest-neighbor classification have
1877 arisen from improvements to the standard ANN problem. While standard ANN

1878 has been well studied in high-dimensional spaces (see, e.g., [21, 22, 50]), in constant-
1879 dimensional Euclidean space, the most efficient data structures involve variants of
1880 the *Approximate Voronoi Diagram* (or AVD) (see [19, 20, 43, 75]). Mount *et al.* [51]
1881 proposed a data structure specifically tailored for ε -classification. Unfortunately,
1882 this work was based on older technology, and its results are not competitive when
1883 compared to the most recent advances on standard ANN search via AVDs.

1884 All previous results have query and space complexities that depend on n , the
1885 total size of the training set P . In many cases, a much smaller portion of the training
1886 set may suffice to correctly ε -classify queries. In the exact setting, these smaller set
1887 of points would be the set of border points of P , of size k , which are the ones that
1888 define the boundaries between classes. Moreover, the concept of border points can be
1889 generalized in the context of ε -classification (see Section 6.2 for a formal definition).
1890 Thus, denote k_ε as the number of ε -border points, where $k \leq k_\varepsilon \leq n$. Ideally, we
1891 would like the query and space complexities of answering chromatic ε -ANN queries
1892 to depend on k_ε instead of n .

1893 There are different approaches to achieve this goal via training set reduction,
1894 as described in Chapters 3 to 5. In general, the idea in such case would be to select
1895 a subset $R \subseteq P$, which is then used to build a standard AVD to answer ε -ANN
1896 queries over R . However, depending on the approach, one obtains different subset
1897 sizes and classification guarantees. From the condensation heuristics described in
1898 Chapter 4, we know that it is possible to compute subsets R of size $\mathcal{O}(k)$ in $\mathcal{O}(n^2)$
1899 time. However, when AVDs are built from these subsets, the resulting data structures
1900 are likely to introduce classification errors, especially for query points that should be
1901 easily ε -classified (as described in Chapter 5). Thus, while often used in practice,
1902 these approaches do not guarantee that chromatic ε -ANN queries are answered
1903 correctly. The second approach would involve computing a coresnet for ε -classification,
1904 as defined in Chapter 5. Recall that a coresnet R guarantees that every query point
1905 will be correctly classified when assigning the class of the point of R returned by the
1906 AVD. That is, for any query $q \in \mathbb{R}^d$, the point of R returned by the AVD belongs to
1907 the same class as one of q 's ε -approximate nearest-neighbors in P . Unfortunately,
1908 the size of the described coresnet can be as large as $\mathcal{O}((k \log \Delta)/\varepsilon^{d-1})$, where Δ is
1909 the spread of P .

1910 Contributions

1911 From the previous section, we have seen that existing approaches for ε -
1912 classification achieve only one of the following goals:

- 1913 • The size of the resulting data structure is dependent only on ε , k_ε (the number
1914 of ε -border points) and d , while being independent from n and Δ .
- 1915 • It guarantees correct ε -classification for any query point.

1916 The main result of this chapter is an approach that achieves both goals. We
1917 propose a new data structure built specifically to answer chromatic ε -ANN queries
1918 over the training set P , which we call a *Chromatic AVD*. Given any query point

1919 $q \in \mathbb{R}^d$, this data structure returns the class to be assigned to q , which matches the
 1920 class of at least one of q 's ε -approximate nearest-neighbors in P . More generally,
 1921 our data structure returns a set of classes such that there is an ε -approximate
 1922 nearest-neighbor of q from each of these classes.

1923 Therefore, the Chromatic AVD can be used to correctly ε -classify any query
 1924 point. The main result of this work is summarized in the following theorem, expressed
 1925 in the form of a space-time tradeoff based on a parameter γ .

1926 **Theorem 6.1.** *Given a training set P of n labeled points in \mathbb{R}^d , an error parameter
 1927 $0 < \varepsilon \leq \frac{1}{2}$, and a separation parameter $2 \leq \gamma \leq \frac{1}{\varepsilon}$. Let k_ε be the number of ε -border
 1928 points of P . There exists a data structure for ε -classification, called Chromatic AVD,
 1929 with:*

$$\text{Query time: } \mathcal{O}\left(\log(k_\varepsilon \gamma) + \frac{1}{(\varepsilon \gamma)^{\frac{d-1}{2}}}\right) \quad \text{Space: } \mathcal{O}\left(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon}\right).$$

1930 Which can be constructed in time $\tilde{\mathcal{O}}\left(\left(n + k_\varepsilon / (\varepsilon \gamma)^{\frac{3}{2}(d-1)}\right) \gamma^d \log \frac{1}{\varepsilon}\right)$.




Figure 6.1: Examples of the space partitioning achieved by any standard AVD, compared to the Chromatic AVD data structure proposed in this chapter. Our approach subdivides the space around the boundaries defined by the ε -border points, while ignoring other boundaries.

1931 By setting γ to either of its extreme values, we obtain the following query times
 1932 and space complexities.

Corollary 6.2. *The separation parameter γ describes the tradeoffs between the query time and space complexity of the Chromatic AVD. This yields the following results:*

$$\begin{aligned} \text{If } \gamma = 2 &\rightarrow \quad \text{Query time: } \mathcal{O}\left(\log k_\varepsilon + \frac{1}{\varepsilon^{\frac{d-1}{2}}}\right) \quad \text{Space: } \mathcal{O}\left(k_\varepsilon \log \frac{1}{\varepsilon}\right). \\ \text{If } \gamma = \frac{1}{\varepsilon} &\rightarrow \quad \text{Query time: } \mathcal{O}\left(\log \frac{k_\varepsilon}{\varepsilon}\right) \quad \text{Space: } \mathcal{O}\left(\frac{k_\varepsilon}{\varepsilon^d}\right). \end{aligned}$$

1933 The approach towards constructing this data structure is hybrid, combining
 1934 a quadtree-induced partitioning of space (leveraging similar techniques to the ones
 1935 used for standard AVDs), with the construction of coresets for only some cells of this
 1936 partition. All other cells can be discarded, and a new quadtree can be built with
 1937 only the remaining cells. The final size of the tree is bounded in terms of k_ε . This
 1938 technique allows us to maintain coresets in the most critical regions of space, and
 1939 thus, avoiding the dependency on the spread of P .

1940 6.2 Preliminary Ideas and Intuition

1941 First, we need to introduce some preliminary definitions and notations that
 1942 are relevant to the results presented in the remaining of the chapter. Given any
 1943 point $q \in \mathbb{R}^d$, denote its nearest-neighbor as $\text{nn}(q)$, and the distance between them
 1944 by $\mathbf{d}_{\text{nn}}(q) = \mathbf{d}(q, \text{nn}(q))$.

1945 Additionally, let's introduce a few concepts and related properties that will
 1946 prove useful in the construction of the Chromatic AVD. These are Well-Separated
 1947 Pair Decompositions [79] (WSPDs), Quadtrees [48, 80], and Approximate Voronoi
 1948 Diagrams [19, 20, 43, 75] (AVDs).

1949 Well-Separated Pair Decompositions: Given the point set P , and a separation factor
 1950 $\sigma > 2$, we say that two sets $X, Y \subseteq P$ are *well separated* if they can be enclosed
 1951 within two disjoint balls of radius r , such that the distance between the centers
 1952 of these balls is at least σr . We say that X and Y form a *dumbbell*, where both
 1953 sets are the *heads* of this dumbbell. Consider the line segment that connects
 1954 the centers of both balls, and let z and ℓ be the center and length of this line
 1955 segment, respectively (*i.e.*, the center and the length of the dumbbell). The
 1956 following properties hold when $\sigma > 4$, for $x \in X$ and $y \in Y$:

$$1957 \quad \mathbf{d}(x, z) < \ell \quad \ell < 2\mathbf{d}(x, y) \quad \ell > \mathbf{d}(x, y)/2.$$

1958 Furthermore, a well-separated pair decomposition of P is defined as a set
 1959 $\mathcal{D} = \{(X_i, Y_i)\}_i$ where every X_i and Y_i are well separated, and for every two
 1960 distinct points $p_1, p_2 \in P$ there exists a unique pair $\mathcal{P} = (X, Y) \in \mathcal{D}$ such that
 1961 $p_1 \in X$ and $p_2 \in Y$, or vice-versa. It is known how to construct a WSPD of P
 1962 with $\mathcal{O}(\sigma^d n)$ pairs in $\mathcal{O}(n \log n + \sigma^d n)$ time.

1963 Quadtrees: These are tree data structures that provide a hierarchical partition of
 1964 space. Each node in this tree consists of a d -dimensional hypercube, where
 1965 non-leaf nodes partition its corresponding hypercube into 2^d equal parts. The
 1966 root of this tree corresponds to the $[0, 1]^d$ hypercube. We will use a variant of
 1967 this structure called a *balanced box-decomposition* tree (BBD tree) [42]. Such
 1968 data structure satisfies the following properties:

- 1969 1. Given a point set P , such a tree can be built in $\mathcal{O}(n \log n)$ time, having
 1970 space $\mathcal{O}(n)$ such that each leaf node contains at most one point of P .

- 1971 2. Given a collection \mathcal{U} of n quadtree boxes in $[0, 1]^d$, such a tree can be built
 1972 in $\mathcal{O}(n \log n)$ time, having $\mathcal{O}(n)$ nodes such that the subdivision induced
 1973 by its leaf cells is a refinement of the subdivision induced by the Quadtree
 1974 boxes in \mathcal{U} .
 1975 3. Given the trees from 1 or 2, it is possible to determine the leaf cell
 1976 containing any arbitrary query point q in $\mathcal{O}(\log n)$ time.

1977 Approximate Voronoi Diagram: Generally, AVDs are quadtree-based data structures
 1978 that can be used to efficiently answer ANN queries. The partitioning of
 1979 space induced by this data structure is often generated from a WSPD of
 1980 P . Additionally, every leaf cell w of this quadtree has an associated set of
 1981 ε -representatives R_w that has the following property: for any query point $q \in w$,
 1982 at least one point in R_w is one of q 's ε -approximate nearest-neighbors in P .

1983 New Ideas and Intuitions. Consider the space partitioning induced by a
 1984 standard AVD, as previously described. By construction, any leaf cell w of this
 1985 partition has an associated set of ε -representatives R_w . Evidently, for the purposes
 1986 of ε -classification, the most important information related to this leaf cell comes
 1987 from the classes of the points in R_w , and not necessarily the points themselves.

1988 This leads to an initial approach to simplify an AVD. We distinguish between
 1989 two types of leaf cells, based on the points inside their corresponding ε -representative
 1990 sets. Any leaf cell w is said to be:

- 1991 • **Resolved:** If every point in R_w belongs to the same class.
- 1992 • **Ambiguous:** Otherwise, if at least two points in R_w belong to different classes.

1993 Clearly, there is no need to store the set of ε -representatives of any resolved
 1994 leaf cell, as instead, we can simply mark the leaf cell w as *resolved with the class*
 1995 that is shared by all the points in R_w . This effectively reduces the space needed for
 1996 such cells to be constant.

1997 Furthermore, it seems that the bulk of the “work” needed to decide the class
 1998 of a given query point can be carried out by the ambiguous leaf cells, along with
 1999 some groupings of resolved leaf cells. The data structure presented in this chapter,
 2000 called Chromatic AVD, builds upon this hypothesis.

2001 Additionally, we formally define the set of ε -border points of the training set
 2002 P . This set, denoted as K_ε , contains any point $p \in P$ for which there exist some
 2003 $q \in \mathbb{R}^d$ and $\bar{p} \in P$, such that p and \bar{p} are ε -approximate nearest-neighbors of q , and
 2004 both belong to different classes. Denote $k_\varepsilon = |K_\varepsilon|$ as the number of ε -border points
 2005 of the training set P . Note that $K_\varepsilon \subseteq K_{\varepsilon'}$ if and only if $\varepsilon \leq \varepsilon'$. Additionally, note
 2006 that K_0 defines the set of (exact) border points of P , where $k = k_0$.

2007 This generalization of the definition of border points seems better suited to
 2008 analyze the problem of ε -classification, as illustrated in Figure 6.2. Figure 6.2b shows
 2009 the ε -approximate bisectors between the two closest and two farthest points (the first
 2010 two belong to K_0 , while the others belong to K_ε but not K_0). A hypothetical leaf
 2011 cell w is sufficiently separated from the only two exact border points, but intersects




Figure 6.2: Intuition to assume that K_ε (and not K_0) is needed to ε -classify some query points.

the ε -approximate bisectors between the two farthest points. This implies that inside the cell w lie query points that *can only* be ε -classified with one class, and others with the other class, forcing this cell to be ambiguous. This suggests that K_0 is insufficient to account for the necessary complexity of ε -classification.

6.3 Chromatic AVD Construction

In this section, we describe our method for constructing the proposed Chromatic AVD. The following overview outlines the necessary steps followed to construct this data structure.

- *The Build step* (Section 6.3.1): Consists of building an initial quadtree-based subdivision of space, designed specifically to achieve the properties described in Lemma 6.3.
- *The Reduce step* (Section 6.3.2): Seeks to identify the leaf cells of the initial subdivision that are relevant for ε -classification, as well as those that can be ignored or simplified. This process consists of the following substeps.
 - Computing the sets of ε -representatives for every leaf cell of the initial quadtree.
 - Based on these sets, marking the leaf cells as either ambiguous or resolved.
 - Selecting those leaf cells which are relevant for ε -classification.
 - Building a new quadtree-based subdivision using the previously selected leaf cells.

6.3.1 The Build Step

We begin by constructing the tree T_{init} using similar methods as the ones used to construct a standard AVD. Thus, the first step is to compute a well-separated

2035 pair decomposition \mathcal{D} of P using a constant separation factor of $\sigma > 4$. While the
 2036 standard construction would use all pairs in this decomposition, for the purpose of
 2037 the Chromatic AVD, we filter \mathcal{D} to only keep bichromatic pairs. Denote $\mathcal{D}' \subseteq \mathcal{D}$ to
 2038 be the set of bichromatic pairs in \mathcal{D} , where a pair $\mathcal{P} \in \mathcal{D}$ is said to be bichromatic if
 2039 and only if the dumbbell heads separate points of different classes. Note that \mathcal{D}' can
 2040 be computed similarly to \mathcal{D} , using a simple modification of the well-known algorithm
 2041 for computing WSVDs [79] (the details are left to the reader).

2042 Next, we compute an initial set of quadtree boxes $\mathcal{U}(\mathcal{P})$ for every pair in \mathcal{D}'
 2043 as follows. This construction depends on two constants c_1 and c_2 whose assignment
 2044 will be described later in this section. For $0 \leq i \leq \lceil \log(c_1 1/\varepsilon) \rceil$, we define $b_i(\mathcal{P})$ as
 2045 the ball centered at z of radius $r_i = 2^i \ell$. Thus, this set of balls involves radius values
 2046 ranging from ℓ to $\Theta(\ell/\varepsilon)$. For each such ball $b_i(\mathcal{P})$, let $\mathcal{U}_i(\mathcal{P})$ be the set of quadtree
 2047 boxes of size $r_i/(c_2 \gamma)$ that overlap the ball. Let $\mathcal{U}(\mathcal{P})$ denote the union of all these
 2048 boxes over all the $\mathcal{O}(\log 1/\varepsilon)$ values of i .

2049 After performing this process on every pair of the filtered decomposition \mathcal{D}' ,
 2050 take the union of all these boxes denoted as $\mathcal{U} = \bigcup_{\mathcal{P} \in \mathcal{D}'} \mathcal{U}(\mathcal{P})$. Finally, build the tree
 2051 T_{init} from the set of quadtree boxes \mathcal{U} , leveraging property 2 of quadtrees described
 2052 in Section 6.2. Additionally, for each class i in the training set, build an auxiliary
 2053 tree T_{aux}^i from the point set P_i (*i.e.*, the points of P of that are labeled with class i),
 2054 using property 1 of quadtrees. These auxiliary trees will be used together with T_{init}
 2055 in order to build our final tree T , the Chromatic AVD.

2056 While the standard AVD construction satisfies that all resulting leaf cells of
 2057 the tree have certain separation properties from the points of set P , the same is not
 2058 true for tree T_{init} . However, the following result describes a relaxed notion of the
 2059 separation properties, now based on the classes of the points, which are achieved by
 2060 T_{init} .

2061 **Lemma 6.3** (Chromatic Separation Properties). *Given two separation parameters
 2062 $\gamma > 2$ and $\varphi > 3$, every leaf cell w of the tree T_{init} satisfies at least one of the
 2063 following separation properties, where b_w is the minimum enclosing ball of w :*

- 2064 (i) *$P \cap \gamma b_w$ is empty (see Figure 6.3a), and hence b_w is concentrically γ -separated
 2065 from P .*
- 2066 (ii) *The cell w can be resolved with the classes present inside $P \cap \varphi b_w$ (see Fig-
 2067 ure 6.3b).*

2068 *Proof.* Let w be any leaf cell of T_{init} , with center c_w and side length s_w , where its
 2069 (minimum) enclosing ball b_w has radius $r_w = \sqrt{d}/2 s_w$ and shares the center c_w .
 2070 Additionally, let $x_i \in P_i$ be a 1-approximate nearest-neighbor of c_w among the points
 2071 of P of class i . In other words, for each class of points we use the auxiliary trees T_{aux}^i
 2072 to compute a 1-approximate nearest-neighbor of c_w . A few cases unfold from here:

2073 The first case is rather simple. If $4\gamma\varphi b_w \cap \{x_i\}_i = \emptyset$, knowing that the points x_i
 2074 are 1-approximate nearest-neighbors of c_w , this implies that the ball $2\gamma\varphi b_w$ is empty
 2075 (*i.e.*, we know that $2\gamma\varphi b_w \cap P = \emptyset$). Clearly, this means the the first separation
 2076 property holds for w .




Figure 6.3: Basic separation properties achieved by during the construction of the Chromatic AVD.

2077 Consider the case when $|4\gamma\varphi b_w \cap \{x_i\}_i| = 1$, and let i be the class of the point
 2078 that lies inside $4\gamma\varphi b_w$. Following similar arguments to the previous case, this implies
 2079 that only points of class i could potentially lie inside of $2\gamma\varphi b_w$. Then, check if x_i
 2080 lies inside the smaller ball expansion $2\gamma b_w$. If not, we know that γb_w is empty (*i.e.*,
 2081 $\gamma b_w \cap P = \emptyset$), making the first separation property hold for w . Otherwise, we know
 2082 that $2\gamma b_w$ contains at least one point (*i.e.*, x_i), and additionally we know that $2\gamma b_w$
 2083 is φ -separated from points of all other classes but i (as $2\gamma\varphi b_w$ only contains points
 2084 of class i). Given that $\varphi > 3$, the nearest neighbor of every query point inside $2\gamma b_w$
 2085 has class i . Therefore, w can be resolved with class i (namely, $C_w = \{i\}$), satisfying
 2086 the second separation property.

2087 Lastly, it is possible that $|4\gamma\varphi b_w \cap \{x_i\}_i| \geq 2$. However, it is possible to show
 2088 that if this is the case, it immediately implies that every point inside $4\gamma\varphi b_w$ actually
 2089 lies inside of some ball b'_w which is β -separated from w (see Figure 6.4a), where
 2090 $\beta = 1/\varepsilon$. Let $x, y \in 4\gamma\varphi b_w$ be the two points of different classes inside the ball
 2091 $4\gamma\varphi b_w$ with largest pairwise distance. Thus, it is easy to show that all the points
 2092 inside $4\gamma\varphi b_w$ lie inside the two balls centered at x and y with radii equal to $d(x, y)$,
 2093 as shown in Figure 6.4b. By definition of the (bichromatic) well-separated pair
 2094 decomposition \mathcal{D}' , there exists a pair $\mathcal{P} \in \mathcal{D}'$ that contains x and y each in one of
 2095 its dumbbells, with length ℓ and center z . Now, we define the ball b'_w with center at
 2096 z and radius $r'_w = \max(d(z, x), d(z, y)) + d(x, y)$. By definition of \mathcal{P} , we know that
 2097 $d(z, x), d(z, y) \leq \ell$ and $d(x, y) \leq 2\ell$, thus making $r'_w \leq 3\ell$. Let L be the distance
 2098 from w to z , we distinguish two cases based on the relationship between L and ℓ :

- 2099 • $L > c_1\beta\ell$. Consider the distance that separates the ball b'_w from the cell w .

$$d(w, b'_w) = L - r'_w > c_1\beta\ell - r'_w \geq (c_1\beta/3 - 1)r'_w$$




Figure 6.4: It is possible that points of ≥ 2 classes lie inside of $\gamma\varphi b_w$. However, this case can be reduced to the two separation properties illustrated in Figure 6.3.

Since $\beta = 1/\varepsilon$, for all sufficiently large constants $c_1 \geq 3(1 + \varepsilon)$, the distance $d(w, b'_w)$ exceeds $\beta r'_w$ which implies that b'_w is concentrically β -separated from w .

- $L \leq c_1 \beta \ell$. We will show that this case cannot occur, since otherwise the dumbbell \mathcal{P} would have caused w to be split, contradicting the assumption that it is a leaf cell of T_{init} . Since x, y , and w are all contained in the ball $4\gamma\varphi b_w$ whose center lies within w , we have both that $d(x, w) \leq 4\gamma\varphi r_w$, and $\ell < 2d(x, y) \leq 2(8\gamma\varphi r_w) = 16\gamma\varphi r_w$. Thus, by the triangle inequality, we have:

$$L \leq d(x, y) + d(x, w) < \ell + 4\gamma\varphi r_w < 16\gamma\varphi r_w + 4\gamma\varphi r_w = 20\gamma\varphi r_w$$

Because $L \leq c_1 \beta \ell$, it follows from our construction that there is at least one ball $b_i(\mathcal{P})$ (with $0 \leq i \leq \lceil \log c_1 \beta \rceil$) that overlaps w . Let b denote the smallest such ball, and let r denote its radius. By the construction, we have that $r \leq \max(\ell, 2L)$. Since our construction generates all quadtree boxes of size $r/(c_2 \gamma)$ that overlap b , it follows that $s_w \leq r/(c_2 \gamma)$. Thus, we have:

$$r_w = s_w \frac{\sqrt{d}}{2} \leq \frac{r \sqrt{d}}{2c_2 \gamma} \leq \frac{\max(\ell, 2L) \sqrt{d}}{2c_2 \gamma} < \frac{20\gamma\varphi r_w \sqrt{d}}{c_2 \gamma} = \frac{20\varphi r_w \sqrt{d}}{c_2}$$

Choosing $c_2 \geq 20\varphi\sqrt{d}$ yields the desired contradiction.

Finally, if $|4\gamma\varphi b_w \cap \{x_i\}_i| \geq 2$, we know all points inside $4\gamma\varphi b_w$ are β -separated from w . We can now proceed similarly to the previous case, by checking if one of the computed 1-approximate nearest-neighbors lies inside the ball $2\gamma b_w$. If

2117 $2\gamma b_w \cap \{x_i\}_i = \emptyset$ we know that γb_w is empty (*i.e.*, $\gamma b_w \cap P = \emptyset$), making the first
 2118 separation property hold for w . Otherwise, note that b'_w is completely contained
 2119 inside $2\gamma(1 + \varepsilon)b_w$. Given that $\varphi > 3$, it is possible to show that for any query point
 2120 in w , all points in b'_w are valid ε -approximate nearest neighbors. This implies that we
 2121 can resolve w with the class of any of the points inside of b'_w , thus satisfying the second
 2122 separation property. In particular, we mark w as resolved with every class present in
 2123 the inner cluster b'_w , namely, $C_w = \{l(p) \mid \forall p \in b'_w \cap P\} = \{i \mid x_i \in 4\gamma\varphi b_w\}$. \square

2124 6.3.2 The Reduce Step

2125 From the initial partitioning as described in Lemma 6.3, every leaf cell w of
 2126 T_{init} is either concentrically γ -separated from P (*i.e.*, $\gamma b_w \cap P = \emptyset$), or it is already
 2127 marked as resolved. For every leaf cell w in the first case, we will compute a set of
 2128 ε -representatives by leveraging the *concentric ball lemma* (see Lemma 5.1 in [75]). It
 2129 states that there exists a set R_w of ε -representatives for w of size $\mathcal{O}(1/(\varepsilon\gamma)^{\frac{d-1}{2}})$, and
 2130 provides a way to compute such set.

2131 Instead of directly applying this result, we use it to compute a set of $\varepsilon/3$ -
 2132 representatives for any leaf cell w that is yet unresolved. Essentially, this leads to the
 2133 same asymptotic upper-bound on the size of R_w , meaning that $|R_w| = \mathcal{O}(1/(\varepsilon\gamma)^{\frac{d-1}{2}})$.
 2134 Once R_w is computed, we can proceed to mark w as either resolved or ambiguous as
 2135 follows.

2136 *Procedure to Mark Leaf Cells:* For every leaf cell w , this procedure marks w as either
 2137 resolved or ambiguous, following a few defined cases that unfold from the contents
 2138 of the set R_w of points selected as representatives for w . Let $r_w^- = \varepsilon(1 - \gamma)r_w/3$.

2139 1. If all the points in R_w belong to the same class.

2140 For every point $p \in R_w$ and class $i \in C$, compute a 1-approximate nearest-
 2141 neighbor of p among the points of P_i , denoted as the point $x_{p,i}$. If $d(p, x_{p,i}) < r_w^-$,
 2142 then add $x_{p,i}$ to R_w . It is easy to show that $x_{p,i}$ would also be an ε -representative
 2143 for w . Repeat this for every point originally in R_w , and every class in the
 2144 training set.

2145 (a) If any point $x_{p,i}$ was added to R_w , proceed with *Case 2*.

2146 (b) Otherwise, mark w as resolved with the class of the points in R_w . Namely,
 2147 let i be the class of every point in R_w , then $C_w = \{i\}$.

2148 2. If R_w contains points of more than one class.

2149 Before proceeding, we will do some basic pruning of the set R_w . For every class
 2150 i , compute a *net* among the points of R_w of class i , using a radius of r_w^- to
 2151 compute the net, and replace the points of class i in R_w with the computed net.
 2152 It is easy to see that the remaining points of R_w are a set of ε -representatives
 2153 of w , and that every two points in R_w of the same class are at distance at least
 2154 r_w^- .

2155 (a) If the diameter of R_w is less than r_w^- , it is easy to prove that all the points
 2156 in R_w are ε -representatives of any point inside b_w . Therefore, w can be
 2157 labeled as resolved with the class of all of the points in R_w . That is,
 2158 $C_w = \{l(p) \mid \forall p \in R_w\}$.

2159 (b) If the diameter is greater than or equal to r_w^- , w is marked as ambiguous.

2160 Let \mathcal{A} and \mathcal{R} be the sets of ambiguous and resolved leaf cells of T_{init} , respectively.
 2161 We will use some of these cells to build the Chromatic AVD, while ignoring the
 2162 remaining cells.

2163 Consider the set of resolved leaf cells \mathcal{R} , we partition this set into two subsets
 2164 \mathcal{R}_b and \mathcal{R}_i (named *boundary* and *interior* resolved leaf cells, respectively). We say
 2165 a resolved leaf cell w_1 belongs to \mathcal{R}_b , if and only if there exists another resolved
 2166 leaf cell w_2 adjacent to w_1 , such that $C_{w_2} \setminus C_{w_1} \neq \emptyset$. Every other resolved leaf cell
 2167 belongs to \mathcal{R}_i (*i.e.*, $\mathcal{R}_i = \mathcal{R} \setminus \mathcal{R}_b$). Note that both sets \mathcal{R}_b and \mathcal{R}_i can be identified
 2168 by a simple traversal over the leaf cells of T_{init} , using linear time in the size of the
 2169 tree¹.

2170 Finally, we build a new tree T from the set of ambiguous and boundary resolved
 2171 leaf cells $\mathcal{A} \cup \mathcal{R}_b$. By well-known construction methods of quadtrees, as described in
 2172 Section 6.2, the leaf cells of T either belong to $\mathcal{A} \cup \mathcal{R}_b$, or are “Steiner” leaf cells
 2173 added during the construction of T that cover the remainder of the space that is
 2174 uncovered by $\mathcal{A} \cup \mathcal{R}_b$.

2175 **Lemma 6.4.** *For any leaf cell w in the tree T such that $w \notin \mathcal{A} \cup \mathcal{R}_b$, w must cover a
 2176 space that is also covered by a collection of leaf cells of T_{init} , all of which are resolved
 2177 with the same set of classes C_w .*

2178 *Proof.* This becomes apparent from the construction of T . In the new tree T ,
 2179 consider any leaf cell w of T that is not part of $\mathcal{A} \cup \mathcal{R}_b$ (*i.e.*, a “Steiner” leaf cell
 2180 added during the construction of the tree). Now, recall that the leaf cells of both
 2181 T and T_{init} are a partitioning of (the same) space, which means that we can define
 2182 $\mathcal{W}_w = \{w' \in T_{\text{init}} \mid w \cap w' \neq \emptyset\}$ as the collection of leaf cells of T_{init} that cover the
 2183 same space covered by w .

2184 Now, for any fixed w of T , it is easy to see that any two leaf cells $w_1, w_2 \in \mathcal{W}_w$
 2185 must be resolved with the same set of classes. Otherwise, at least one of these two
 2186 would be part of the set \mathcal{R}_b , which would be a contradiction to the fact that w is
 2187 a “Steiner” leaf cell of T . Therefore, any query inside w can be answered with the
 2188 classes $C_w = C_{w_1} = C_{w_2}$, and this can be determined during preprocessing by a
 2189 single query on T_{init} (*e.g.*, finding the leaf cell of T_{init} that contains the center c_w of
 2190 w is sufficient to know the contents of C_w). \square

2191 This implies that after building tree T , and with some extra preprocessing to
 2192 resolve the “Steiner” leaf cells of the tree, we can use the resulting data structure to
 2193 correctly answer chromatic ε -approximate nearest-neighbor queries over the training

¹Two leaf cells are adjacent if and only if a vertex of one of the cells “touches” the other cell. This implies that the number of adjacency relations (*i.e.*, edges in the implicit graph where the leaf cells are the nodes) is $\mathcal{O}(2^d m)$, where m is the number of leaf cells of the tree T_{init} .

2194 set P . In other words, T can be used to answer ε -classification queries over P . We
2195 call this data structure T the Chromatic AVD.

2196 **Lemma 6.5.** *The construction of T takes $\tilde{\mathcal{O}}(n\gamma^d \log \frac{1}{\varepsilon})$ time.*

2197 *Proof.* Let's analyze the total time needed to build our Chromatic AVD, namely the
2198 tree T , by analyzing the time required to perform each step of the construction.

- 2199 • Building T_{init} has essentially the same complexity of building any standard
2200 AVD [19, 20, 43, 75]. This means that constructing T_{init} takes $\mathcal{O}(m \log m)$ time,
2201 where $m = n\gamma^d \log \frac{1}{\varepsilon}$. Note that during the construction, while computing the
2202 set of ε -representatives of each leaf cell, each leaf cell can already be marked
2203 as either ambiguous or resolved.
- 2204 • Building the auxiliary trees T_{aux}^i for every class i , takes $\mathcal{O}(n \log n)$ time, as the
2205 number of classes of P is considered to be a constant. Recall that because
2206 these trees are only used to for 1-ANN queries, they only need to be standard
2207 Quadtrees, and not AVDs.
- 2208 • Identifying the set \mathcal{R}_b requires a traversal over the leaf-level partitioning of
2209 the space, which is linear in terms of the number of cells. Therefore, this step
2210 requires $\mathcal{O}(m)$ time.
- 2211 • Once the sets of ambiguous and boundary resolved leaf cells are identified,
2212 namely, the sets \mathcal{A} and \mathcal{R}_b , the final tree T can be built. Roughly, this step
2213 takes $\mathcal{O}(m \log m)$ time.
- 2214 • Finally, we must resolve each “Steiner” leaf cell of T , which can be done by
2215 a single query over T_{init} , each taking $\mathcal{O}(\log m)$ time. Thus, this step takes
2216 $\mathcal{O}(m \log m)$ total time.

2217 All together, the total construction time is dominated by the first step. There-
2218 fore, the time required to construct T is $\mathcal{O}(m \log m) = \tilde{\mathcal{O}}(m) = \tilde{\mathcal{O}}(n\gamma^d \log \frac{1}{\varepsilon})$. \square

2219 6.4 Tree-size Analysis

2220 6.4.1 Initial Size Bounds

2221 Define the set of important leaf cells \mathcal{I} of the tree T_{init} as those leaf cells w for
2222 which there exists two ε -border points inside $\rho\gamma b_w$ for some constant ρ , such that
2223 the distance between these points is lower-bounded by $\Omega(\varepsilon\gamma r_w)$. Formally, we define
2224 this set as $\mathcal{I} = \{w \in T_{\text{init}} \mid \exists p_1, p_2 \in \rho\gamma b_w \cap K_\varepsilon, d(p_1, p_2) = \Omega(\varepsilon\gamma r_w)\}$.

2225 **Lemma 6.6.** *The number of important leaf cells of T_{init} is $|\mathcal{I}| = \mathcal{O}(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon})$.*

2226 *Proof.* This proof follows from a charging argument on the set K_ε of ε -border points
2227 of P . More specifically, consider a well-separated pair decomposition \mathcal{D}'' of K_ε with
2228 constant separation factor of $\sigma > 4$, the charging scheme assigns every important

leaf cell $w \in \mathcal{I}$ to a pair of \mathcal{D}'' . Recall that \mathcal{D}'' can generally be considered to have $\mathcal{O}(k_\varepsilon)$ pairs, where $k_\varepsilon = |K_\varepsilon|$. It is important to note that both K_ε and \mathcal{D}'' need not be computed.

Given that $w \in \mathcal{I}$, we know there exist two points $p_1, p_2 \in \rho\gamma b_w \cap K_\varepsilon$. Let $\mathcal{P} \in \mathcal{D}''$ be the pair of \mathcal{D}'' that contains both p_1 and p_2 , each in one of its dumbbell heads. We then charge w to the pair \mathcal{P} . Denote z and ℓ to be the center and length of \mathcal{P} , respectively, we know the following. First, note that the distance from c_w (the center of w) to z is $d(c_w, z) \leq \rho\gamma r_w + \ell$. Additionally, we know that $\ell \geq d(p_1, p_2)/2 = \Omega(\varepsilon\gamma r_w)$ by the properties of WSPDs described in Section 6.2. Therefore, this implies that the ratio $d(c_w, z)/\ell = \mathcal{O}(1/\varepsilon)$.

Finally, fix some pair $\mathcal{P} \in \mathcal{D}''$ with center z and length ℓ , and consider all important leaf cells according to their distance to z . For any value of $i \in [0, 1, \dots, \mathcal{O}(\log 1/\varepsilon)]$, consider all leaf cells that can charge \mathcal{P} whose distance to z is between $2^i\ell$ and $2^{i+1}\ell$. From our previous analysis, $r_w \geq d(c_w, z)/\rho\gamma \geq 2^i\ell/\rho\gamma$. By a simple packing argument, the number of such leaf cells is at most $\mathcal{O}(\gamma^d)$. Thus, a total of $\mathcal{O}(\gamma^d \log 1/\varepsilon)$ cells can charge \mathcal{P} . Note that no leaf cell whose distance to z is $\Omega(\ell/\varepsilon)$ can charge \mathcal{P} , as it would contradict the fact that both p_1 and p_2 are separated by a distance of $\Omega(\varepsilon\gamma r_w)$. Finally, the proof follows by knowing that there are at most $\mathcal{O}(k_\varepsilon)$ pairs in \mathcal{D}'' . \square

Lemma 6.7. *Every ambiguous leaf cell of T_{init} is important, namely $\mathcal{A} \subseteq \mathcal{I}$.*

Proof. Consider any ambiguous leaf cell $w \in \mathcal{A}$ of the tree T_{init} . Knowing that w is ambiguous implies that there must exist some point $q \in \frac{\gamma}{2}b_w$ for which two of the ε -representatives of w are valid ε -approximate nearest neighbors for q , both points belong to different classes, and the distance between them is $\Omega(\varepsilon\gamma r_w)$. Formally, denote these points as $p_1, p_2 \in P$ such that $l(p_1) \neq l(p_2)$, $d(p_1, p_2) \geq \varepsilon(1-\gamma)r_w/4$, and $d(q, p_1), d(q, p_2) \leq (1+\varepsilon)d_{\text{nn}}(q)$.

We will see now how to bound the distance from c_w to any of these points as a constant factor of r_w (recall that $r_w = \sqrt{d}/2 s_w$). From the proof of Lemma 6.3 in [75], we know that the ball $c_3\gamma b_w \cap P \neq \emptyset$, for some constant $c_3 \geq 1 + 2c_2/\sqrt{d}$. In other words, $d_{\text{nn}}(c_w) \leq c_3\gamma r_w$. From this, we can say that $d_{\text{nn}}(q) \leq (\frac{1}{2} + c_3)\gamma r_w$. Applying the triangle inequality yields that $d(c_w, p_1) \leq d(c_w, q) + d(q, p_1) \leq (\frac{1}{2} + (1+\varepsilon)(\frac{1}{2} + c_3))\gamma r_w$. Similarly, we can achieve the same bound for $d(c_w, p_2)$.

Therefore, both $p_1, p_2 \in \rho\gamma b_w$ for sufficiently large constant ρ (i.e., $\rho \geq \varepsilon(\frac{1}{2} + c_3) + c_3 + 1$). Knowing also that $d(p_1, p_2) = \Omega(r_w^-) = \Omega(\varepsilon\gamma r_w)$ yields that the leaf cell $w \in \mathcal{I}$. \square

Lemma 6.8. *Every boundary resolved leaf cell of T_{init} is important, namely $\mathcal{R}_b \subseteq \mathcal{I}$.*

Proof. Let $w_1 \in \mathcal{R}_b$ be any boundary resolved leaf cell of the tree T_{init} , we know there exists another leaf cell $w_2 \in \mathcal{R}_b$ adjacent to w_1 , such that there exists some class $i \in C_{w_2} \setminus C_{w_1}$. Let b_{w_1} and b_{w_2} be the corresponding bounding balls of w_1 and w_2 . By definition, any point q on the boundary shared by w_1 and w_2 has at least one ε -representative from each cell, namely some points $p_1 \in R_{w_1}$ and $p_2 \in R_{w_2}$, where

2271 $l(p_1) \neq i$ and $l(p_2) = i$. Additionally, by similar arguments to the ones described in
 2272 Lemma 6.7, we know that both $p_1, p_2 \in \rho\gamma b_w$ for sufficiently large constant ρ .

2273 Now, we proceed to prove that $d(p_1, p_2) \geq r_w^-/2$. First, note that if w_1 was
 2274 resolved by the initial marking of leaf cells as described in Lemma 6.3, then p_2
 2275 must lie outside of γb_w . In such cases, clearly $d(p_1, p_2) \geq r_w^-/2$. The remaining
 2276 possibility is that w_1 was resolved after computing the set of representatives. From
 2277 the described procedure, in Case 1, we know that if $d(p_1, p_2) < r_w^-/2$, the point $x_{p_1,i}$
 2278 (which is a 1-approximate nearest-neighbor of p_1 among points in P_i) would hold that
 2279 $d(p_1, x_{p_1,i}) < r_w^-$. Hence, $x_{p_1,i}$ should have been added to the set of representatives of
 2280 w_1 , contradicting the assumption that w_1 is resolved, or that C_{w_1} does not contain i .
 2281 All together, we have that $d(p_1, p_2) = \Omega(r_w^-) = \Omega(\varepsilon\gamma r_w)$. From the definition of the
 2282 set of important leaf cells, $w \in \mathcal{I}$. \square

2283 Lemmas 6.7 and 6.8 imply that all the leaf cells used to build T belong to
 2284 the set of important leaf cells (*i.e.*, $\mathcal{A} \cup \mathcal{R}_b \subseteq \mathcal{I}$), whose size is upper-bounded by
 2285 Lemma 6.6. All together, and leveraging construction methods of quadtrees (see
 2286 Section 6.2), the size of T is proportional to the total number of leaf cells used to
 2287 build it, which we now know is $\mathcal{O}(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon})$. However, we also need to account for
 2288 the set of ε -representatives stored for each ambiguous leaf cell, leading to a worst-case
 2289 upper-bound of $\mathcal{O}(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon} \cdot 1/(\varepsilon\gamma)^{\frac{d-1}{2}})$ total space to store T .

2290 6.4.2 Spatial Amortization

2291 The previous result can be improved by applying a technique called *spatial
 2292 amortization*, described by Arya *et al.* [75]. That is, we can remove the extra
 2293 $\mathcal{O}(1/(\varepsilon\gamma)^{\frac{d-1}{2}})$ factor from the analysis of the space requirements for T .

2294 This will be twofold process, as in order to successfully apply spatial amortiza-
 2295 tion to the analysis of the data structure, we first need to further reduce the set of
 2296 ε -representatives of every ambiguous leaf cell in the tree. Actually, the new set will
 2297 no longer be a set of ε -representatives, but it will just be a *weak* ε -coreset for query
 2298 points inside of each leaf cell.

2299 **Lemma 6.9.** *The total space required to store the ambiguous leaf cells of T is
 2300 $\mathcal{O}(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon})$.*

2301 Consider any ambiguous leaf cell w of T , and in particular, consider the set
 2302 R_w of ε -representatives of w . By construction, R_w has the property that every point
 2303 $q \in b_w$ has at least one ε -approximate nearest-neighbor in the set R_w . However, note
 2304 that the opposite is not necessarily true, as not every $p \in R_w$ is an ε -approximate
 2305 nearest-neighbor of some point in b_w . Even worst, while the fact the w is ambiguous
 2306 indicates that at least two points in R_w belong to K_ε , the remaining points of R_w
 2307 might not, which in turn prevents the application of a spatial amortization analysis.
 2308 Overall, this suggests some of the points of R_w might not be necessary to distinguish
 2309 between the classes that change the classification of points inside b_w (see Figure 6.5a).

2310 This can be resolved as follows. Suppose we have access to the Voronoi diagram
 2311 of the set of points R_w , and consider the boundaries between adjacent cells of this

diagram. Any boundary that separates two points of R_w of different classes, and that intersects b_w , is relevant to the classification any query point inside b_w . Formally, we define the set $R'_w \subseteq R_w$ of border points of R_w as (see Figure 6.5b):

$$R'_w = \{p \in R_w \mid \exists q \in b_w, p' \in R_w \text{ such that } l(p) \neq l(p') \wedge d(q, p) = d(q, p')\}$$

This new set R'_w has some useful properties. Note that for any query point $q \in b_w$, its (exact) nearest-neighbor in R'_w belongs to the same class as its (exact) nearest-neighbor in R_w , which itself is an ε -approximate nearest-neighbor of q among the points of P . In other words, R'_w is an ε -coreset for any query point in b_w . This implies that we can replace the set of ε -representatives of w with the set R'_w . Moreover, this means that by the definition of ε -border points, $R'_w \subseteq K_\varepsilon$. Note that we can use the algorithm described in Chapter 3 to compute R'_w in $\mathcal{O}(|R_w| \cdot |R'_w|^2)$ time, increases the construction time described in Lemma 6.5 by a factor of k_ε^2 .




Figure 6.5: The set R_w of ε -representatives of w can be reduced to the set R'_w . This later set is a subset of K_ε , and can be charged to a proportional number of relevant pairs of \mathcal{D}'' .

Now, let's proceed with the charging argument over the pairs of the same WSPD \mathcal{D}'' used in Lemma 6.6. Instead of only charging w to a single pair (as described in Lemma 6.7), we charge every point stored in R'_w to a pair of \mathcal{D}'' . Thus, consider the following procedure to find a sufficient number of pairs to charge all the points in R'_w , which is derived from a similar procedure proposed in [75]. See Figure 6.5c for an illustrative example.

1. Compute a net of R'_w using radius r_w^- , and denote this subset R''_w . Given that all the points of R'_w that belong to the same class are already separated by a distance of at least r_w^- , we know that $|R''_w| = \Theta(|R'_w|)$, hiding constants² that depend exponentially on d .
2. Find the two of points of $p_1, p_2 \in R''_w$ with smallest pairwise distance, and consider the pair of $\mathcal{P} \in \mathcal{D}''$ that contains both points p_1 and p_2 , each in one

²More specifically, we know that $|R''_w| \geq |R'_w|/\phi^{d-1}c$, where ϕ^{d-1} is the kissing number in d -dimensional Euclidean space, and c is the number of classes in P .

2335 of its dumbbell heads. Note that by having computed a net in the previous
2336 step, $\mathbf{d}(p_1, p_2) \geq r_w^-$.

- 2337 3. Delete one of the two points from R''_w (without lost of generality, delete p_1).
2338 4. Charge every point of R'_w that is covered by p_1 (*i.e.*, whose distance to p_1 is
2339 $\leq r_w^-$) to the pair \mathcal{P} . By the arguments described in step 1 on the size of R''_w ,
2340 we know that \mathcal{P} receives a charge from $\mathcal{O}(1)$ points of R'_w .
2341 5. Repeat steps 2-4 with the remaining points of R''_w until the set is empty.

2342 Evidently, the number of pairs found (and charged) equals $|R''_w| - 1$. All
2343 together, we have that the total space required to store all the ambiguous leaf cells
2344 is proportional to the sum of charges to every pair of \mathcal{D}'' . Using the same arguments
2345 as Lemma 6.6, this implies that the total storage is $\mathcal{O}(k_\varepsilon \gamma^d \log \frac{1}{\varepsilon})$. This completes
2346 the proof of Theorem 6.1.

2347 Chapter 7: Conclusions

2348 In this dissertation, we have presented different techniques that successfully
2349 reduce the query time and space complexities of the nearest-neighbor rule. While
2350 usually nearest-neighbor classification would be highly dependent on n , the size of
2351 the training set P , our results provide a clear way to reduce its dependency to k .
2352 This parameter is defined as the number of border points of P , and characterizes
2353 the intrinsic complexity of the class boundaries of the training set.

2354 The results presented here open a series of potential directions for future re-
2355 search to further improve the efficiency of the nearest-neighbor rule. While mostly
2356 theoretical, these results can help expand the scope of applicability of this classifica-
2357 tion method for larger real-world use cases. Furthermore, our results on training set
2358 reduction algorithms imply their potential application beyond the scope of nearest-
2359 neighbor classification, by reducing the data used to train other classification models.
2360 While many of the classification guarantees provided here would not hold –with most
2361 likeliness– for other classification methods, the upper-bounds on subset sizes would.

2362 In the following sections, we retrace most of the results presented throughout
2363 this dissertation, and discuss some of the remaining open problems and limitations.

2364 7.1 Training Set Reduction

2365 Most of the results presented in this dissertation can be described as training
2366 set reduction approaches. In this case, the goal is to find a subset $R \subseteq P$, whose
2367 induced class boundaries resemble the original class boundaries of P . This subset is
2368 then used by the nearest-neighbor rule to answer any incoming queries, instead of
2369 using the full training set. As seen throughout this book, different approaches yield
2370 different subsets, and depending on the approach used, the classification guarantees
2371 of the nearest-neighbor rule after training set reduction can vary.

2372 7.1.1 Boundary Preservation

2373 The goal of boundary preservation algorithms is to compute the set of border
2374 points of P . Given that by definition, this subset of points fully characterize the class
2375 boundaries of P , this reduction of the training set does not affect the classification
2376 accuracy of the nearest-neighbor rule. That is, the class boundaries remain the same.
2377 Moreover, these algorithms achieve the expected dependency reduction from n to k ,
2378 as k is defined as the size of their selected subset.

2379 In Chapter 3, we present an improvement over Eppstein’s recent algorithm [8]

2380 to compute the set of border points of any training set $P \subset \mathbb{R}^d$, with constant d .
2381 This improved algorithm reduces the complexity of computing such subset to $\mathcal{O}(nk^2)$
2382 worst-case time, while Eppstein’s original algorithm has $\mathcal{O}(n^2 + nk^2)$ runtime. A
2383 paper describing this result is under submission, and can be found here [52].

2384 7.1.2 Condensation Algorithms

2385 Alternatively, condensation algorithms aim to compute subsets $R \subseteq P$ whose
2386 induced class boundaries are “similar” to the original class boundaries, albeit not the
2387 same as with boundary preservation algorithms. As formally described in Chapter 2,
2388 condensation algorithms select either consistent or selective subsets of P , which can
2389 only guarantee the correct classification of points in P . This is a rather popular line
2390 of research [9, 13–17, 34, 38], with many heuristic approaches proposed to compute
2391 such subsets. Until our work, theoretical results on these heuristics were scarce.

2392 In Chapter 4, we present the first theoretical results on upper-bounding the
2393 subset sizes of condensation algorithms. In particular, we show that FCNN and
2394 MSS, which were considered state-of-the-art for the condensation problem, can not
2395 be bounded in terms of k . Additionally, we propose new quadratic-time algorithms
2396 called SFCNN, RSS, and VSS, and prove upper-bounds on their subset sizes as a
2397 function of k . These upper-bounds are tight up-to constant factors, and supports
2398 the good empirical behavior observed for these condensation algorithms. Most of
2399 these results were published in a series of papers [54–56].

2400 Despite our efforts, some questions remain open. First, it is unclear whether
2401 our upper-bound for the RSS algorithm in terms of k can be improved to have linear
2402 dependencies on d , as our current result has a term with exponential dependency
2403 on d . So far, we have been unable to find a matching lower-bound for RSS in
2404 terms of k , keeping the hope that this improvement is indeed possible. The other
2405 important open question revolves around SFCNN. While this algorithm performs
2406 really well in practice, we were only able to prove a $\mathcal{O}(k \log(1/\gamma))$ upper-bound on
2407 its selected subset (compared with the $\mathcal{O}(k)$ upper-bound on RSS). Thus, it remains
2408 open whether the $\log(1/\gamma)$ factor in the upper-bound of SFCNN can be dropped.

2409 7.1.3 ε -Coresets

2410 Both *boundary preservation* and *condensation* algorithms make the assumption
2411 that the nearest-neighbor rule computes nearest-neighbors exactly. However, efficient
2412 data structures for nearest-neighbor search compute nearest-neighbor approximately
2413 rather than exactly. That is, given an approximation parameter $\varepsilon \in [0, 1]$, a query
2414 point $q \in \mathcal{X}$ can be assigned the class of any point of P whose distance to q is
2415 at most $1 + \varepsilon$ times the distance from q to its true nearest-neighbor. This relaxed
2416 assumption immediately breaks the classification guarantees provided by both types
2417 of training set reduction techniques.

2418 Chapter 5 presents a new framework for training set reduction that is sensitive
2419 to these approximations, as well as a characterization of ε -coresets for the nearest-
2420 neighbor rule. First, we define such an ε -coreset as a subset $R \subseteq P$ where for every

2421 query point $q \in \mathcal{X}$ the class of its exact nearest-neighbor in R is the same as the
 2422 class of a valid ε -approximate nearest-neighbor of q in P . In order to compute
 2423 such subsets, we extended the criteria used for condensation to be approximation-
 2424 sensitive, and modified the condensation algorithms from Chapter 4 to compute
 2425 subsets under these new criteria. Finally, this gives us a way to compute ε -coresets
 2426 of size $\mathcal{O}(k \log \frac{1}{\gamma} (1/\varepsilon)^{\text{ddim}(\mathcal{X})+1})$. These results have been published in the following
 2427 paper [66].

2428 The biggest shortcoming of this result is the size of the ε -coresets, having high
 2429 dependencies on γ and ε . The lower-bound presented in Lemma 5.3 indicate that
 2430 our coreset construction might not be optimal, and that the terms on γ and ε could
 2431 be potentially dropped. However, despite our efforts, we were unable to achieve this.

2432 7.2 Chromatic Nearest-Neighbor Search

2433 Unsurprisingly, the standard approach to answer nearest-neighbor classification
 2434 queries is to reduce them to nearest-neighbor search queries. However, as described
 2435 in Chapter 6, there exists an alternative approach towards achieving more efficient
 2436 nearest-neighbor classification. This consists of having algorithms and data structures
 2437 to directly compute the class of the nearest-neighbor of any given query point; *i.e.*,
 2438 without computing the nearest-neighbor itself.

2439 To this end, we proposed a tailor-made data structure for approximate nearest-
 2440 neighbor classification called *Chromatic AVD*. Given a training set P in d -dimensional
 2441 Euclidean space (assuming constant d and the ℓ_2 metric) and an approximation
 2442 parameter $\varepsilon \in [0, \frac{1}{2}]$, we construct a quadtree-based partitioning of space to answer
 2443 any classification query approximately. That is, for any query point $q \in \mathbb{R}^d$ this data
 2444 structure returns the class of any of q 's valid ε -approximate nearest-neighbors in P .
 2445 This data structure is designed as a simplification of state-of-the-art AVDs [20] for
 2446 approximate nearest-neighbor search, and completely reduces its dependency from
 2447 n to k_ε , which is defined as the number of ε -border points of P , and describes the
 2448 intrinsic complexity of the ε -approximate class boundaries. These results have been
 2449 published in the following paper [78].

2450 There are a few possible improvements and extensions that would be worth
 2451 exploring. First, the query time and space dependencies of the Chromatic AVD are
 2452 expressed in terms of k_ε and not k , where $k_\varepsilon \geq k$. However, it is unclear if k_ε can
 2453 be expressed in terms of k , or if the analysis and construction of the Chromatic
 2454 AVD can be improved to reduce the expressed dependencies to be in terms of k .
 2455 Another interesting direction is trying to obtain query times that are query-sensitive,
 2456 similar to the result by Mount *et al.* [51] in 1997, where query points with high
 2457 chromatic density are easier to answer than query points with low chromatic density.
 2458 Finally, yet another direction would be on extending this data structure to answer
 2459 ε -approximate k-NN classification queries, for $k \geq 1$. While there exists some work
 2460 on this problem [51, 81], all existing results have dependencies on n .

2461

Bibliography

2462

- [1] E. Fix and J. L. Hodges. Discriminatory analysis, nonparametric discrimination: Consistency properties. *US Air Force School of Aviation Medicine*, Technical Report 4(3):477+, January 1951.
- [2] Charles J. Stone. Consistent nonparametric regression. *The annals of statistics*, pages 595–620, 1977.
- [3] T. Cover and P. Hart. Nearest neighbor pattern classification. *IEEE Trans. Inf. Theor.*, 13(1):21–27, January 1967.
- [4] Luc Devroye. On the inequality of cover and hart in nearest neighbor discrimination. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, pages 75–78, 1981.
- [5] Juha Heinonen. *Lectures on analysis on metric spaces*. Springer Science & Business Media, 2012.
- [6] David Bremner, Erik Demaine, Jeff Erickson, John Iacono, Stefan Langerman, Pat Morin, and Godfried Toussaint. Output-sensitive algorithms for computing nearest-neighbour decision boundaries. In Frank Dehne, Jörg-Rüdiger Sack, and Michiel Smid, editors, *Algorithms and Data Structures: 8th International Workshop, WADS 2003, Ottawa, Ontario, Canada, July 30 - August 1, 2003. Proceedings*, pages 451–461, Berlin, Heidelberg, 2003. Springer Berlin Heidelberg.
- [7] Kenneth L Clarkson. More output-sensitive geometric algorithms. In *Proceedings 35th Annual Symposium on Foundations of Computer Science*, pages 695–702. IEEE, 1994.
- [8] David Eppstein. Finding relevant points for nearest-neighbor classification. In *Symposium on Simplicity in Algorithms (SOSA)*, pages 68–78. SIAM, 2022.
- [9] P. Hart. The condensed nearest neighbor rule (corresp.). *IEEE Trans. Inf. Theor.*, 14(3):515–516, September 1968.

- 2488 [10] Gordon Wilfong. Nearest neighbor problems. In *Proceedings of the Seventh*
 2489 *Annual Symposium on Computational Geometry*, SCG '91, pages 224–233, New
 2490 York, NY, USA, 1991. ACM.
- 2491 [11] A. V. Zukhba. NP-completeness of the problem of prototype selection in the
 2492 nearest neighbor method. *Pattern Recog. Image Anal.*, 20(4):484–494, December
 2493 2010.
- 2494 [12] Kamyar Khodamoradi, Ramesh Krishnamurti, and Bodhayan Roy. Consistent
 2495 subset problem with two labels. In *Conference on Algorithms and Discrete*
 2496 *Applied Mathematics*, pages 131–142. Springer, 2018.
- 2497 [13] Godfried Toussaint. Open problems in geometric methods for instance-based
 2498 learning. In Jin Akiyama and Mikio Kano, editors, *JCDCG*, volume 2866 of
 2499 *Lecture Notes in Computer Science*, pages 273–283. Springer, 2002.
- 2500 [14] Godfried Toussaint. Proximity graphs for nearest neighbor decision rules: Recent
 2501 progress. In *Progress”, Proceedings of the 34th Symposium on the INTERFACE*,
 2502 pages 17–20, 2002.
- 2503 [15] Norbert Jankowski and Marek Grochowski. Comparison of instances selection
 2504 algorithms I. Algorithms survey. In *Artificial Intelligence and Soft Computing-
 2505 ICAISC 2004*, pages 598–603. Springer, 2004.
- 2506 [16] Fabrizio Angiulli. Fast nearest neighbor condensation for large data sets classifica-
 2507 tion. *IEEE Transactions on Knowledge and Data Engineering*, 19(11):1450–1464,
 2508 2007.
- 2509 [17] Ricardo Barandela, Francesc J Ferri, and J Salvador Sánchez. Decision boundary
 2510 preserving prototype selection for nearest neighbor classification. *International*
 2511 *Journal of Pattern Recognition and Artificial Intelligence*, 19(06):787–806, 2005.
- 2512 [18] Salvador Garcia, Joaquin Derrac, Jose Cano, and Francisco Herrera. Prototype
 2513 selection for nearest neighbor classification: Taxonomy and empirical study.
 2514 *IEEE Trans. Pattern Anal. Mach. Intell.*, 34(3):417–435, March 2012.
- 2515 [19] Sariel Har-Peled. A replacement for Voronoi diagrams of near linear size. In
 2516 *Proc. 42nd Annu. IEEE Sympos. Found. Comput. Sci.*, pages 94–103, 2001.
- 2517 [20] Sunil Arya, Guilherme D. da Fonseca, and David M. Mount. Optimal ap-
 2518 proximate polytope membership. In *Proceedings of the Twenty-Eighth Annual*
 2519 *ACM-SIAM Symposium on Discrete Algorithms*, pages 270–288. SIAM, 2017.
- 2520 [21] Piotr Indyk, Rajeev Motwani, Prabhakar Raghavan, and Santosh Vempala.
 2521 Locality-preserving hashing in multidimensional spaces. In *Proceedings of the*
 2522 *Twenty-Ninth Annual ACM Symposium on Theory of Computing*, STOC '97,
 2523 page 618–625, New York, NY, USA, 1997. Association for Computing Machinery.

- 2524 [22] Yu A. Malkov and D. A. Yashunin. Efficient and robust approximate nearest
2525 neighbor search using hierarchical navigable small world graphs. 42(4):824–836,
2526 apr 2020.
- 2527 [23] Corinna Cortes and Vladimir Vapnik. Support-vector networks. In *Machine*
2528 *Learning*, pages 273–297, 1995.
- 2529 [24] Jürgen Schmidhuber. Deep learning in neural networks: An overview. *CoRR*,
2530 abs/1404.7828, 2014.
- 2531 [25] Oren Boiman, Eli Shechtman, and Michal Irani. In defense of nearest-neighbor
2532 based image classification. In *2008 IEEE Conference on Computer Vision and*
2533 *Pattern Recognition*, pages 1–8. IEEE, 2008.
- 2534 [26] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search
2535 with GPUs. *arXiv preprint arXiv:1702.08734*, 2017.
- 2536 [27] Ruiqi Guo, Philip Sun, Erik Lindgren, Quan Geng, David Simcha, Felix Chern,
2537 and Sanjiv Kumar. Accelerating large-scale inference with anisotropic vector
2538 quantization. In *International Conference on Machine Learning*, 2020.
- 2539 [28] Marc Khouri and Dylan Hadfield-Menell. Adversarial training with Voronoi
2540 constraints. *CoRR*, abs/1905.01019, 2019.
- 2541 [29] Neehar Peri, Neal Gupta, W. Ronny Huang, Liam Fowl, Chen Zhu, Soheil Feizi,
2542 Tom Goldstein, and John P. Dickerson. Deep k -nn defense against clean-label
2543 data poisoning attacks. In *European Conference on Computer Vision*, pages
2544 55–70. Springer, 2020.
- 2545 [30] Chawin Sitawarin and David Wagner. On the robustness of deep k -nearest
2546 neighbors. In *2019 IEEE Security and Privacy Workshops (SPW)*, pages 1–7.
2547 IEEE, 2019.
- 2548 [31] Nicolas Papernot and Patrick McDaniel. Deep k -nearest neighbors: Towards
2549 confident, interpretable and robust deep learning. *arXiv preprint arXiv:1803.04765*,
2550 2018.
- 2551 [32] Belur V. Dasarathy. Nearest unlike neighbor (nun): an aid to decision confidence
2552 estimation. *Optical Engineering*, 34(9):2785–2792, 1995.
- 2553 [33] D. Randall Wilson and Tony R. Martinez. Instance pruning techniques. In
2554 *Proceedings of the Fourteenth International Conference on Machine Learning*,
2555 ICML ’97, pages 403–411, San Francisco, CA, USA, 1997. Morgan Kaufmann
2556 Publishers Inc.
- 2557 [34] G. L. Ritter, H. B. Woodruff, S. R. Lowry, and T. L. Isenhour. An algorithm for
2558 a selective nearest neighbor decision rule. *IEEE Transactions on Information*
2559 *Theory*, 21(6):665–669, 1975.

- 2560 [35] Richard Nock and Marc Sebban. Sharper bounds for the hardness of prototype
 2561 and feature selection. In *Proceedings of the 11th International Conference on*
 2562 *Algorithmic Learning Theory*, ALT '00, page 224–237, Berlin, Heidelberg, 2000.
 2563 Springer-Verlag.
- 2564 [36] Lee-Ad Gottlieb, Aryeh Kontorovich, and Pinhas Nisnevitch. Near-optimal
 2565 sample compression for nearest neighbors. In *Advances in Neural Information*
 2566 *Processing Systems*, pages 370–378, 2014.
- 2567 [37] Rajesh Chitnis. Refined lower bounds for nearest neighbor condensation. In
 2568 Sanjoy Dasgupta and Nika Haghtalab, editors, *Proceedings of The 33rd Interna-*
 2569 *tional Conference on Algorithmic Learning Theory*, volume 167 of *Proceedings*
 2570 *of Machine Learning Research*, pages 262–281. PMLR, 29 Mar–01 Apr 2022.
- 2571 [38] Geoffrey W. Gates. The reduced nearest neighbor rule (corresp.). *IEEE*
 2572 *Transactions on Information Theory*, 18(3):431–433, 1972.
- 2573 [39] Ahmad Biniaz, Sergio Cabello, Paz Carmi, Jean-Lou De Carufel, Anil Mahesh-
 2574 wari, Saeed Mehrabi, and Michiel Smid. On the minimum consistent subset
 2575 problem. In *Workshop on Algorithms and Data Structures*, pages 155–167.
 2576 Springer, 2019.
- 2577 [40] Kenneth L Clarkson. A randomized algorithm for closest-point queries. *SIAM*
 2578 *Journal on Computing*, 17(4):830–847, 1988.
- 2579 [41] Andrew C Yao and F Frances Yao. A general approach to d-dimensional
 2580 geometric queries. In *Proceedings of the seventeenth annual ACM symposium*
 2581 *on Theory of computing*, pages 163–168, 1985.
- 2582 [42] Sunil Arya, David M. Mount, Nathan S. Netanyahu, Ruth Silverman, and
 2583 Angela Y. Wu. An optimal algorithm for approximate nearest neighbor searching
 2584 in fixed dimensions. *J. ACM*, 45(6):891–923, November 1998.
- 2585 [43] Sunil Arya, Guilherme D. Da Fonseca, and David M. Mount. Approximate
 2586 polytope membership queries. *SIAM Journal on Computing*, 47(1):1–51, 2018.
- 2587 [44] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search
 2588 with gpus. *IEEE Transactions on Big Data*, 7(3):535–547, 2019.
- 2589 [45] Ruiqi Guo, Philip Sun, Erik Lindgren, Quan Geng, David Simcha, Felix Chern,
 2590 and Sanjiv Kumar. Accelerating large-scale inference with anisotropic vector
 2591 quantization. In Hal Daumé III and Aarti Singh, editors, *Proceedings of the*
 2592 *37th International Conference on Machine Learning*, volume 119 of *Proceedings*
 2593 *of Machine Learning Research*, pages 3887–3896. PMLR, 13–18 Jul 2020.
- 2594 [46] Jon Louis Bentley. Multidimensional binary search trees used for associative
 2595 searching. *Communications of the ACM*, 18(9):509–517, 1975.

- 2596 [47] Hanan Samet. The quadtree and related hierarchical data structures. *ACM*
 2597 *Comput. Surv.*, 16(2):187–260, jun 1984.
- 2598 [48] Hanan Samet. *The design and analysis of spatial data structures*, volume 85.
 2599 Addison-Wesley Reading, MA, 1990.
- 2600 [49] Hanan Samet. *Foundations of multidimensional and metric data structures*.
 2601 Morgan Kaufmann, 2006.
- 2602 [50] Alexandr Andoni, Piotr Indyk, and Ilya Razenshteyn. Approximate nearest
 2603 neighbor search in high dimensions. *arXiv preprint arXiv:1806.09823*, 2018.
- 2604 [51] David M. Mount, Nathan S. Netanyahu, Ruth Silverman, and Angela Y. Wu.
 2605 Chromatic nearest neighbor searching: A query sensitive approach. *Computa-*
 2606 *tional Geometry*, 17(3):97 – 119, 2000.
- 2607 [52] Alejandro Flores-Velazco. Improved search of relevant points for nearest-neighbor
 2608 classification. *CoRR*, abs/2203.03567, 2022.
- 2609 [53] Timothy M Chan. Output-sensitive results on convex hulls, extreme points, and
 2610 related problems. *Discrete & Computational Geometry*, 16(4):369–387, 1996.
- 2611 [54] Alejandro Flores-Velazco and David M. Mount. Guarantees on nearest-neighbor
 2612 condensation heuristics. In *Proceedings of the 31st Canadian Conference on*
 2613 *Computational Geometry, CCCG 2019, August 8-10, 2019, University of Alberta,*
 2614 *Edmonton, Alberta, Canada*, 2019.
- 2615 [55] Alejandro Flores-Velazco. Social distancing is good for points too! In *Pro-*
 2616
 2617 *2020, August 5-7, 2020, University of Saskatchewan, Saskatoon, Saskatchewan,*
 2618 *Canada*, 2020.
- 2619 [56] Alejandro Flores-Velazco and David M. Mount. Guarantees on nearest-neighbor
 2620 condensation heuristics. *Computational Geometry*, 95:101732, 2021.
- 2621 [57] Pankaj K. Agarwal, Sariel Har-Peled, and Kasturi R. Varadarajan. Geometric
 2622 approximation via coresets. *Combinatorial and computational geometry*, 52:1–30,
 2623 2005.
- 2624 [58] Jeff M. Phillips. Coresets and sketches, 2016.
- 2625 [59] Dan Feldman. Core-sets: Updated survey. In *Sampling Techniques for Supervised*
 2626 *or Unsupervised Tasks*, pages 23–44. Springer, 2020.
- 2627 [60] Sariel Har-Peled and Soham Mazumdar. On coresets for k-means and k-median
 2628 clustering. In *Proceedings of the thirty-sixth annual ACM symposium on Theory*
 2629 *of computing*, pages 291–300, 2004.

- 2630 [61] Cenk Baykal, Lucas Liebenwein, Igor Gilitschenski, Dan Feldman, and Daniela
 2631 Rus. Data-dependent coresets for compressing neural networks with applications
 2632 to generalization bounds. *arXiv preprint arXiv:1804.05345*, 2018.
- 2633 [62] Vladimir Braverman, Dan Feldman, and Harry Lang. New frameworks for
 2634 offline and streaming coreset constructions, 2016.
- 2635 [63] Dan Feldman and Michael Langberg. A unified framework for approximating
 2636 and clustering data. In *Proceedings of the forty-third annual ACM symposium*
 2637 *on Theory of computing*, pages 569–578, 2011.
- 2638 [64] Lucas Liebenwein, Cenk Baykal, Harry Lang, Dan Feldman, and Daniela
 2639 Rus. Provable filter pruning for efficient neural networks. *arXiv preprint*
 2640 *arXiv:1911.07412*, 2019.
- 2641 [65] Murad Tukan, Cenk Baykal, Dan Feldman, and Daniela Rus. On coresets for
 2642 support vector machines, 2020.
- 2643 [66] Alejandro Flores-Velazco and David M. Mount. Coresets for the nearest-neighbor
 2644 rule, 2020.
- 2645 [67] Sariel Har-Peled and Manor Mendel. Fast construction of nets in low-dimensional
 2646 metrics and their applications. *SIAM Journal on Computing*, 35(5):1148–1184,
 2647 2006.
- 2648 [68] Uriel Feige. A threshold of $\ln n$ for approximating set cover. *Journal of the*
 2649 *ACM (JACM)*, 45(4):634–652, 1998.
- 2650 [69] Azaria Paz and Shlomo Moran. Non deterministic polynomial optimization
 2651 problems and their approximations. *Theoretical Computer Science*, 15(3):251–
 2652 277, 1981.
- 2653 [70] Carsten Lund and Mihalis Yannakakis. On the hardness of approximating
 2654 minimization problems. *Journal of the ACM (JACM)*, 41(5):960–981, 1994.
- 2655 [71] V. Chvatal. A greedy heuristic for the set-covering problem. *Mathematics of*
 2656 *Operations Research*, 4(3):233–235, 1979.
- 2657 [72] Petr Slavik. A tight analysis of the greedy algorithm for set cover. In *Proceedings*
 2658 *of the Twenty-eighth Annual ACM Symposium on Theory of Computing*, STOC
 2659 ’96, pages 435–441, New York, NY, USA, 1996. ACM.
- 2660 [73] F. Aurenhammer and H. Edelsbrunner. An optimal algorithm for constructing
 2661 the weighted Voronoi diagram in the plane. *Pattern Recognition*, 17(2):251 –
 2662 257, 1984.
- 2663 [74] Jon Louis Bentley and James B. Saxe. Decomposable searching problems i.
 2664 static-to-dynamic transformation. *Journal of Algorithms*, 1(4):301–358, 1980.

- 2665 [75] Sunil Arya, Theocharis Malamatos, and David M. Mount. Space-time tradeoffs
2666 for approximate nearest neighbor searching. *Journal of the ACM (JACM)*,
2667 57(1):1, 2009.
- 2668 [76] Richard Cole and Lee-Ad Gottlieb. Searching dynamic point sets in spaces with
2669 bounded doubling dimension. In *Proceedings of the thirty-eighth annual ACM*
2670 *symposium on Theory of computing*, pages 574–583, 2006.
- 2671 [77] Timothy M. Chan. A minimalist’s implementation of an approxi-
2672 mate nearest neighbor algorithm in fixed dimensions (2006). URL:
2673 <http://tmc.web.engr.illinois.edu/ssss.ps>.
- 2674 [78] Alejandro Flores-Velazco and David M. Mount. Boundary-sensitive approach
2675 for approximate nearest-neighbor classification. In Petra Mutzel, Rasmus Pagh,
2676 and Grzegorz Herman, editors, *29th Annual European Symposium on Algo-*
2677 *rithms, ESA 2021, September 6-8, 2021, Lisbon, Portugal (Virtual Conference)*,
2678 volume 204 of *LIPICS*, pages 44:1–44:15. Schloss Dagstuhl - Leibniz-Zentrum für
2679 Informatik, 2021.
- 2680 [79] Paul B. Callahan and S. Rao Kosaraju. A decomposition of multidimensional
2681 point sets with applications to k -nearest-neighbors and n -body potential fields.
2682 *Journal of the ACM (JACM)*, 42(1):67–90, 1995.
- 2683 [80] Raphael A. Finkel and Jon Louis Bentley. Quad trees a data structure for
2684 retrieval on composite keys. *Acta informatica*, 4(1):1–9, 1974.
- 2685 [81] TWJ van der Horst. Chromatic k-nearest neighbors queries using (near-) linear
2686 space. Master’s thesis, 2021.